

# AWS SDK for Java Developer Guide

March 09, 2017



# Contents

<b>AWS SDK for Java</b>	<b>1</b>
AWS SDK for Java Developer Guide	2
Additional Documentation and Resources	2
Eclipse IDE Support	2
Developing AWS Applications for Android	2
Viewing the SDK's Revision History	2
Building Java Reference Documentation for Earlier SDK versions	2
Getting Started	4
Sign Up for AWS and Create an IAM User	4
Set up the AWS SDK for Java	5
Prerequisites	5
Including the SDK in your project	5
Downloading and extracting the SDK	5
Installing previous versions of the SDK	6
Installing a Java Development Environment	7
Choosing a JVM	7
Set up AWS Credentials and Region for Development	7
Setting AWS Credentials	7
Setting the AWS Region	8
Using the SDK with Apache Maven	9
Create a new Maven package	9
Configure the SDK as a Maven dependency	10
Specifying individual SDK modules	10
Importing all SDK modules	11
Build your project	11
Build the SDK with Maven	11
Using the SDK with Gradle	11
Using the AWS SDK for Java	13
Best Practices for AWS Development with the AWS SDK for Java	13
Amazon S3	13
Avoid ResetExceptions	13
Creating Service Clients	13
Obtaining a Client Builder	13

Creating Async Clients	15
Using DefaultClient	15
Client Lifecycle	15
Working with AWS Credentials	16
Using the Default Credential Provider Chain	16
Setting Credentials	17
Setting an Alternate Credentials Profile	17
Setting an Alternate Credentials File Location	17
AWS Credentials File Format	18
Loading Credentials	18
Specifying a Credential Provider or Provider Chain	19
Explicitly Specifying Credentials	19
More Info	20
AWS Region Selection	20
Checking for Service Availability in an AWS Region	20
Choosing a Region	20
Choosing a Specific Endpoint	21
Automatically Determine the AWS Region from the Environment	21
Default Region Provider Chain	22
Exception Handling	22
Why Unchecked Exceptions?	22
AmazonServiceException (and Subclasses)	23
AmazonClientException	23
Asynchronous Programming	23
Java Futures	24
Asynchronous Callbacks	25
Best Practices	26
Callback Execution	26
Thread Pool Configuration	26
Amazon S3 Asynchronous Access	27
Logging AWS SDK for Java Calls	27
Download the Log4J JAR	27
Setting the Classpath	27
Service-Specific Errors and Warnings	28
Request/Response Summary Logging	28

Verbose Wire Logging	29
Client Networking Configuration	29
Proxy Configuration	29
HTTP Transport Configuration	30
Local Address	30
Maximum Connections	30
Proxy Options	30
Timeouts and Error Handling	30
TCP Socket Buffer Size Hints	30
Access Control Policies	31
Amazon S3 Example	31
Amazon SQS Example	32
Amazon SNS Example	32
Setting the JVM TTL for DNS Name Lookups	33
How to Set the JVM TTL	33
Programming Examples	34
SDK Code Samples	34
How to Get the Samples	34
Building and Running the Samples Using the Command Line	34
Prerequisites	34
Running the Samples	34
Building and Running the Samples Using the Eclipse IDE	35
Prerequisites	35
Running the Samples	35
DynamoDB Examples	36
Working with Tables in DynamoDB	36
Create a Table	37
Create a Table with a Simple Primary Key	37
Create a Table with a Composite Primary Key	38
List Tables	38
Describe (Get Information about) a Table	39
Modify (Update) a Table	41
Delete a Table	41
More Info	42
Working with Items in DynamoDB	42

Retrieve (Get) an Item from a Table	42
Add a New Item to a Table	43
Update an Existing Item in a Table	44
More Info	45
Managing Tomcat Session State with DynamoDB	45
Download the Session Manager	46
Configure the Session-State Provider	46
Configure a Tomcat Server to Use DynamoDB as the Session-State Server	46
Configure Your AWS Security Credentials	46
Configure with Elastic Beanstalk	46
Manage Tomcat Session State with DynamoDB	47
Options Specified in context.xml	47
Troubleshooting	48
Limitations	48
Amazon EC2 Examples	48
Tutorial: Starting an EC2 Instance	48
Prerequisites	48
Create an Amazon EC2 Security Group	48
Create a Key Pair	50
Run an Amazon EC2 Instance	51
Using IAM Roles to Grant Access to AWS Resources on Amazon EC2	52
The default provider chain and EC2 instance profiles	52
Walkthrough: Using IAM roles for EC2 instances	53
Create an IAM Role	53
Launch an EC2 Instance and Specify Your IAM Role	53
Create your Application	54
Transfer the Compiled Program to Your EC2 Instance	56
Run the Sample Program on the EC2 Instance	57
Tutorial: Amazon EC2 Spot Instances	57
Overview	57
Prerequisites	58
Step 1: Setting Up Your Credentials	58
Step 2: Setting Up a Security Group	58
Step 3: Submitting Your Spot Request	60

Step 4: Determining the State of Your Spot Request	62
Step 5: Cleaning Up Your Spot Requests and Instances	64
Bringing It All Together	66
Next Steps	66
Tutorial: Advanced Amazon EC2 Spot Request Management	66
Prerequisites	66
Setting up your credentials	66
Setting up a security group	67
Detailed spot instance request creation options	68
Persistent vs. one-time requests	69
Limiting the duration of a request	70
Grouping your Amazon EC2 spot instance requests	71
How to persist a root partition after interruption or termination	75
How to tag your spot requests and instances	76
Tagging requests	76
Tagging instances	77
Canceling spot requests and terminating instances	79
Canceling a spot request	79
Terminating spot instances	79
Bringing it all together	80
Amazon S3 Examples	80
Creating, Listing, and Deleting Amazon S3 Buckets	80
Create a Bucket	81
List Buckets	81
Delete a Bucket	82
Remove Objects from an Unversioned Bucket Before Deleting It	82
Remove Objects from a Versioned Bucket Before Deleting It	83
Delete an Empty Bucket	84
Performing Operations on Amazon S3 Objects	85
Upload an Object	85
List Objects	85
Download an Object	86
Copy, Move, or Rename Objects	87
Delete an Object	88
Delete Multiple Objects at Once	88

Managing Access to Amazon S3 Buckets Using Bucket Policies	89
Set a Bucket Policy	89
Use the Policy Class to Generate or Validate a Policy	89
Get a Bucket Policy	90
Delete a Bucket Policy	91
More Info	91
Using TransferManager for Amazon S3 Operations	91
Upload Files and Directories	92
Upload a Single File	92
Upload a List of Files	93
Upload a Directory	94
Download Files or Directories	95
Download a Single File	95
Download a Directory	96
Copy Objects	96
Wait for a Transfer to Complete	97
Get Transfer Status and Progress	97
Poll the Current Progress of a Transfer	98
Get Transfer Progress with a ProgressListener	98
Get the Progress of Subtransfers	99
More Info	99
Amazon SQS Examples	99
Working with Amazon SQS Message Queues	100
Create a Queue	100
Listing Queues	101
Get the URL for a Queue	101
Delete a Queue	102
More Info	102
Sending, Receiving, and Deleting Amazon SQS Messages	102
Send a Message	102
Send Multiple Messages at Once	103
Receive Messages	103
Delete Messages after Receipt	104
More Info	104
Getting Temporary Credentials with AWS STS	104

(Optional) Activate and use an AWS STS region	105
Retrieve temporary security credentials from AWS STS	105
Use the temporary credentials to access AWS resources	106
For more information	107
Amazon SWF	107
Amazon SWF Basics	107
Dependencies	107
Imports	108
Using the SWF client class	108
Building a Simple Amazon SWF Application	109
About the example	109
Prerequisites	109
Development environment	109
AWS access	110
Create a SWF project	110
Code the project	112
Common steps for all source files	112
Register a domain, workflow and activity types	112
Implement the activity worker	115
Implement the workflow worker	117
Implement the workflow starter	121
Build the example	122
Run the example	122
Setting the Java classpath	123
Register the domain, workflow and activity types	123
Start the activity and workflow workers	123
Start the workflow execution	123
Complete source for this example	124
For more information	124
Lambda Tasks	124
Set up a cross-service IAM role to run your Lambda function	124
Create a Lambda function	124
Register a workflow for use with Lambda	126
Schedule a Lambda task	126
Handle Lambda function events in your decider	127

Receive output from your Lambda function	128
Complete source for this example	128
Shutting Down Activity and Workflow Workers Gracefully	128
Registering Domains	130
Listing Domains	131
Document History	132
About Amazon Web Services	134



# AWS SDK for Java

# AWS SDK for Java Developer Guide

The [AWS SDK for Java](#) provides a Java API for Amazon Web Services. Using the SDK, you can easily build Java applications that work with Amazon S3, Amazon EC2, Amazon SimpleDB, and more. We regularly add support for new services to the AWS SDK for Java. For a list of the supported services and their API versions that are included with each release of the SDK, view the [release notes](#) for the version that you're working with.

## Additional Documentation and Resources

In addition to this guide, the following are valuable online resources for AWS SDK for Java developers:

- [AWS SDK for Java Reference](#)
- [Java developer blog](#)
- [Java developer forums](#)
- GitHub:
  - [Documentation source](#)
  - [Documentation issues](#)
  - [SDK source](#)
  - [SDK issues](#)
  - [SDK samples](#)
  - [Gitter channel](#)
- [@awsforjava \(Twitter\)](#)
- [release notes](#)

## Eclipse IDE Support

If you develop code using the Eclipse IDE, you can use the [AWS Toolkit for Eclipse](#) to add the AWS SDK for Java to an existing Eclipse project or to create a new AWS SDK for Java project. The toolkit also supports creating and uploading Lambda functions, launching and monitoring Amazon EC2 instances, managing IAM users and security groups, a CloudFormation template editor, and more.

See the [AWS Toolkit for Eclipse User Guide](#) for full documentation.

## Developing AWS Applications for Android

If you're an Android developer, Amazon Web Services publishes an SDK made specifically for Android development: the [AWS Mobile SDK for Android](#). See the [AWS Mobile SDK for Android Developer Guide](#) for full documentation.

## Viewing the SDK's Revision History

To view the release history of the AWS SDK for Java, including changes and supported services per SDK version, see the SDK's [release notes](#).

## Building Java Reference Documentation for Earlier SDK versions

The *AWS SDK for Java Reference* represents the most recent version of the SDK. If you're using an earlier SDK version, you might want to access the SDK reference documentation that matches the version you're using.

The easiest way to build the documentation is using Apache's [Maven](#) build tool. *Download and install Maven first if you don't already have it on your system*, then use the following instructions to build the reference documentation.

### To build reference documentation for an earlier SDK version

1. Locate and select the SDK version that you're using on the [releases](#) page of the SDK repository on GitHub.
2. Choose either the `zip` (most platforms, including Windows) or `tar.gz` (Linux, macOS, or Unix) link to download the SDK to your computer.
3. Unpack the archive to a local directory.
4. On the command line, navigate to the directory where you unpacked the archive, and type the following.

```
mvn javadoc:javadoc
```

5. After building is complete, you'll find the generated HTML documentation in the `aws-java-sdk/target/site/apidocs/` directory.

# Getting Started

This section provides information about how to install, set up, and use the AWS SDK for Java.

## Sign Up for AWS and Create an IAM User

To use the AWS SDK for Java to access Amazon Web Services (AWS), you will need an AWS account and AWS credentials. To increase the security of your AWS account, we recommend that you use an *IAM user* to provide access credentials instead of using your root account credentials.

### Tip

For an overview of IAM users and why they are important for the security of your account, see [Overview of Identity Management: Users](#) in the *IAM User Guide*.

### To sign up for AWS

1. Open <https://aws.amazon.com/> and click **Sign Up**.
2. Follow the on-screen instructions. Part of the sign-up procedure involves receiving a phone call and entering a PIN using your phone keypad.

Next, create an IAM user and download (or copy) its secret access key.

### To create an IAM user

1. Go to the [IAM console](#) (you may need to sign in to AWS first).
2. Click **Users** in the sidebar to view your IAM users.
3. If you don't have any IAM users set up, click **Create New Users** to create one.
4. Select the IAM user in the list that you'll use to access AWS.
5. Open the **Security Credentials** tab, and click **Create Access Key**.

### Note

You can have a maximum of two active access keys for any given IAM user. If your IAM user has two access keys already, then you'll need to delete one of them before creating a new key.

6. On the resulting dialog, click the **Download Credentials** button to download the credential file to your computer, or click **Show User Security Credentials** to view the IAM user's access key ID and secret access key (which you can copy and paste).

## Important

There is no way to obtain the secret access key once you close the dialog. You can, however, delete its associated access key ID and create a new one.

Next, you should *set your credentials* in the AWS shared credentials file or in the environment.

## Tip

If you use the Eclipse IDE, you should consider installing the [AWS Toolkit for Eclipse](#) and providing your credentials as described in [Set up AWS Credentials](#) in the [AWS Toolkit for Eclipse User Guide](#).

## Set up the AWS SDK for Java

Describes how to use the AWS SDK for Java in your project.

### Prerequisites

To use the AWS SDK for Java, you must have:

- a suitable Java Development Environment.
- An AWS account and access keys. For instructions, see [Sign Up for AWS and Create an IAM User](#).
- AWS credentials (access keys) set in your environment or using the shared (by the AWS CLI and other SDKs) credentials file. For more information, see [Set up AWS Credentials and Region for Development](#).

### Including the SDK in your project

To include the SDK your project, use one of the following methods depending on your build system or IDE:

- **Apache Maven** – If you use [Apache Maven](#), you can specify the entire SDK (or specific SDK components) as dependencies in your project. See [Using the SDK with Apache Maven](#) for details about how to set up the SDK when using Maven.
- **Gradle** – If you use [Gradle](#), you can import the Maven Bill of Materials (BOM) in your Gradle project to automatically manage SDK dependencies. See [Using the SDK with Gradle](#) for more infomation.
- **Eclipse IDE** – If you use the Eclipse IDE, you may want to install and use the [AWS Toolkit for Eclipse](#), which will automatically download, install and update the Java SDK for you. For more information and setup instructions, see the [AWS Toolkit for Eclipse User Guide](#).

If you intend to build your projects using a different IDE, with Apache Ant or by any other means, then download and extract the SDK as shown in the next section.

### Downloading and extracting the SDK

We recommend that you use the most recent pre-built version of the SDK for new projects, which provides you with the latest support for all AWS services.

## Note

For information about how to download and build previous versions of the SDK, see [Installing previous versions of the SDK](#).

### To download and extract the latest version of the SDK

1. Download the SDK from <https://sdk-for-java.amazonwebservices.com/latest/aws-java-sdk.zip>.
2. After downloading the SDK, extract the contents into a local directory.

The SDK contains the following directories:

- documentation – contains the API documentation (also available on the web: [AWS SDK for Java Reference](#)).
- lib – contains the SDK .jar files.
- samples – contains working sample code that demonstrates how to use the SDK.
- third-party – contains third-party libraries that are used by the SDK, such as Apache commons logging, AspectJ and the Spring framework.

To use the SDK, add the full path to the lib and third-party directories to the dependencies in your build file, and add them to your java CLASSPATH to run your code.

### Installing previous versions of the SDK

Only the latest version of the SDK is provided in pre-built form. However, you can build a previous version of the SDK using Apache Maven (open source). Maven will download all necessary dependencies, build and install the SDK in one step. Visit <http://maven.apache.org/> for installation instructions and more information.

### To install a previous version of the SDK

1. Go to the SDK's GitHub page at: [AWS SDK for Java \(GitHub\)](#).
2. Choose the tag corresponding to the version number of the SDK that you want. For example, 1.6.10.
3. Click the **Download ZIP** button to download the version of the SDK you selected.
4. Unzip the file to a directory on your development system. On many systems, you can use your graphical file manager to do this, or use the unzip utility in a terminal window.
5. In a terminal window, navigate to the directory where you unzipped the SDK source.
6. Build and install the SDK with the following command ([Maven](#) required):

```
mvn clean install
```

The resulting .jar file is built into the target directory.

7. (Optional) Build the API Reference documentation using the following command:

## Getting Started

```
mvn javadoc:javadoc
```

The documentation is built into the `target/site/apidocs/` directory.

## Installing a Java Development Environment

The AWS SDK for Java requires J2SE Development Kit *6.0 or later*. You can download the latest Java software from <http://www.oracle.com/technetwork/java/javase/downloads/>.

### Important

Java version 1.6 (JS2E 6.0) did not have built-in support for SHA256-signed SSL certificates, which are required for all HTTPS connections with AWS after September 30, 2015.

Java versions 1.7 or newer are packaged with updated certificates and are unaffected by this issue.

## ***Choosing a JVM***

For the best performance of your server-based applications with the AWS SDK for Java, we recommend that you use the *64-bit version* of the Java Virtual Machine (JVM). This JVM runs only in server mode, even if you specify the `-Client` option at run time.

Using the 32-bit version of the JVM with the `-Server` option at run time should provide comparable performance to the 64-bit JVM.

## Set up AWS Credentials and Region for Development

To connect to any of the supported services with the AWS SDK for Java, you must provide AWS credentials. The AWS SDKs and CLIs use *provider chains* to look for AWS credentials in a number of different places, including system/user environment variables and local AWS configuration files.

This topic provides basic information about setting up your AWS credentials for local application development using the AWS SDK for Java. If you need to set up credentials for use within an EC2 instance or if you're using the Eclipse IDE for development, refer to the following topics instead:

- When using an EC2 instance, create an IAM role and then give your EC2 instance access to that role as shown in *Using IAM Roles to Grant Access to AWS Resources on Amazon EC2*.
- Set up AWS credentials within Eclipse using the [AWS Toolkit for Eclipse](#). See [Set up AWS Credentials](#) in the *AWS Toolkit for Eclipse User Guide* for more information.

## Setting AWS Credentials

Setting your credentials for use by the AWS SDK for Java can be done in a number of ways, but here are the recommended approaches:

- Set credentials in the AWS credentials profile file on your local system, located at:
  - `~/.aws/credentials` on Linux, macOS, or Unix

## Getting Started

- C:\Users\USERNAME\.aws\credentials on Windows

This file should contain lines in the following format:

```
[default]
aws_access_key_id = your_access_key_id
aws_secret_access_key = your_secret_access_key
```

Substitute your own AWS credentials values for the values *your\_access\_key\_id* and *your\_secret\_access\_key*.

- Set the AWS\_ACCESS\_KEY\_ID and AWS\_SECRET\_ACCESS\_KEY environment variables.

To set these variables on Linux, macOS, or Unix, use **export**:

```
export AWS_ACCESS_KEY_ID=your_access_key_id
export AWS_SECRET_ACCESS_KEY=your_secret_access_key
```

To set these variables on Windows, use **set**:

```
set AWS_ACCESS_KEY_ID=your_access_key_id
set AWS_SECRET_ACCESS_KEY=your_secret_access_key
```

- For an EC2 instance, specify an IAM role and then give your EC2 instance access to that role. See [IAM Roles for Amazon EC2](#) in the *Amazon EC2 User Guide for Linux Instances* for a detailed discussion about how this works.

Once you have set your AWS credentials using one of these methods, they will be loaded automatically by the AWS SDK for Java by using the default credential provider chain. For further information about working with AWS credentials in your Java applications, see *Working with AWS Credentials*.

## Setting the AWS Region

You should set a default AWS Region that will be used for accessing AWS services with the AWS SDK for Java. For the best network performance, you should choose a region that's geographically close to you (or to your customers).

### Note

If you *don't* select a region, then us-east-1 will be used by default.

You can use similar techniques to setting credentials to set your default AWS region:

- Set the AWS region in the AWS config file on your local system, located at:

- `~/.aws/config` on Linux, macOS, or Unix
- `C:\Users\USERNAME\.aws\config` on Windows

## Getting Started

This file should contain lines in the following format:

```
[default]
region = your_aws_region
```

Substitute your desired AWS region (for example, "us-west-2") for *your\_aws\_region*.

- Set the AWS\_REGION environment variable.

On Linux, macOS, or Unix, use **export**:

```
export AWS_REGION=your_aws_region
```

On Windows, use **set**:

```
set AWS_REGION=your_aws_region
```

Where *your\_aws\_region* is the desired AWS region name.

## Using the SDK with Apache Maven

You can use [Apache Maven](#) to configure and build AWS SDK for Java projects, or to build the SDK itself.

### Note

You must have Maven installed to use the guidance in this topic. If it isn't already installed, visit <http://maven.apache.org/> to download and install it.

## Create a new Maven package

To create a basic Maven package, open a terminal (command-line) window and run:

```
mvn -B archetype:generate \
-DarchetypeGroupId=org.apache.maven.archetypes \
-DgroupId=org.example.basicapp \
-DartifactId=myapp
```

Replace *org.example.basicapp* with the full package namespace of your application, and *myapp* with the name of your project (this will become the name of the directory for your project).

By default, Maven creates a project template for you using the [quickstart](#) archetype, which is a good starting place for many projects. There are more archetypes available; visit the [Maven archetypes](#) page for a list of archetypes packaged with Maven. You can choose a particular archetype to use by adding the *-DarchetypeArtifactId* argument to the *archetype:generate* command. For example:

## Getting Started

```
mvn archetype:generate \
-DarchetypeGroupId=org.apache.maven.archetypes \
-DarchetypeArtifactId=maven-archetype-webapp \
-DgroupId=org.example.webapp \
-DartifactId=mywebapp
```

### Tip

Much more information about creating and configuring Maven projects is provided in the [Maven Getting Started Guide](#).

## Configure the SDK as a Maven dependency

To use the AWS SDK for Java in your project, you'll need to declare it as a dependency in your project's pom.xml file. Beginning with version 1.9.0, you can import individual components or the entire SDK.

### ***Specifying individual SDK modules***

To select individual SDK modules, use the AWS SDK for Java bill of materials (BOM) for Maven, which will ensure that the modules you specify use the same version of the SDK and that they're compatible with each other.

To use the BOM, add a <dependencyManagement> section to your application's pom.xml file, adding aws-java-sdk-bom as a dependency and specifying the version of the SDK you want to use:

```
<dependencyManagement>
  <dependencies>
    <dependency>
      <groupId>com.amazonaws</groupId>
      <artifactId>aws-java-sdk-bom</artifactId>
      <version>1.11.22</version>
      <type>pom</type>
      <scope>import</scope>
    </dependency>
  </dependencies>
</dependencyManagement>
```

You can now select individual modules from the SDK that you use in your application. Because you already declared the SDK version in the BOM, you don't need to specify the version number for each component.

```
<dependencies>
  <dependency>
    <groupId>com.amazonaws</groupId>
    <artifactId>aws-java-sdk-s3</artifactId>
  </dependency>
```

## Getting Started

```
<dependency>
    <groupId>com.amazonaws</groupId>
    <artifactId>aws-java-sdk-dynamodb</artifactId>
</dependency>
</dependencies>
```

### **Importing all SDK modules**

If you would like to pull in the *entire* SDK as a dependency, don't use the BOM method, but simply declare it in your pom.xml like this:

```
<dependencies>
    <dependency>
        <groupId>com.amazonaws</groupId>
        <artifactId>aws-java-sdk</artifactId>
        <version>1.11.22</version>
    </dependency>
</dependencies>
```

## Build your project

Once you have your project set up, you can build it using Maven's package command:

```
mvn package
```

This will create your .jar file in the target directory.

## Build the SDK with Maven

You can use Apache Maven to build the SDK from source. To do so, download the SDK code from GitHub, unpack it locally, and then execute the following Maven command:

```
mvn clean install
```

## Using the SDK with Gradle

To use the AWS SDK for Java in your [Gradle](#) project, use Spring's [dependency management plugin](#) for Gradle, which can be used to import the SDK's Maven Bill of Materials (BOM) to manage SDK dependencies for your project.

### To configure the SDK for Gradle

1. Add the dependency management plugin to your build.gradle file

```
buildscript {
    repositories {
        mavenCentral()
    }
}
```

## Getting Started

```
dependencies {  
    classpath "io.spring.gradle:dependency-management-plugin:1.0.0.RC2"  
}  
  
apply plugin: "io.spring.dependency-management"
```

2. Add the BOM to the *dependencyManagement* section of the file

```
dependencyManagement {  
    imports {  
        mavenBom 'com.amazonaws:aws-java-sdk-bom:1.10.77'  
    }  
}
```

3. Specify the SDK modules that you'll be using in the *dependencies* section

```
dependencies {  
    compile 'com.amazonaws:aws-java-sdk-s3'  
    testCompile group: 'junit', name: 'junit', version: '4.11'  
}
```

Gradle will automatically resolve the correct version of your SDK dependencies using the information from the BOM.

### Note

For more detail about specifying SDK dependencies using the BOM, see *Using the SDK with Apache Maven*.

# Using the AWS SDK for Java

This section provides important general information about programming with the AWS SDK for Java that applies to all services you might use with the SDK.

For service-specific programming information and examples (for Amazon EC2, Amazon S3, SWF, etc.), see *Programming Examples*.

## Best Practices for AWS Development with the AWS SDK for Java

The following best practices can help you avoid issues or trouble as you develop AWS applications with the AWS SDK for Java. We've organized best practices by service.

### Amazon S3

#### **Avoid ResetExceptions**

When you upload objects to Amazon S3 by using streams (either through an [AmazonS3](#) client or [TransferManager](#)), you might encounter network connectivity or timeout issues. By default, the AWS SDK for Java attempts to retry failed transfers by marking the input stream before the start of a transfer and then resetting it before retrying.

If the stream doesn't support mark and reset, the SDK throws a [ResetException](#) when there are transient failures and retries are enabled.

#### **Best Practice**

We recommend that you use streams that support mark and reset operations.

The most reliable way to avoid a [ResetException](#) is to provide data by using a [File](#) or [FileInputStream](#), which the AWS SDK for Java can handle without being constrained by mark and reset limits.

If the stream isn't a [FileInputStream](#) but does support mark and reset, you can set the mark limit by using the `setReadLimit` method of [RequestClientOptions](#). Its default value is 128 KB. Setting the read limit value to *one byte greater than the size of stream* will reliably avoid a [ResetException](#).

For example, if the maximum expected size of a stream is 100,000 bytes, set the read limit to 100,001 (100,000 + 1) bytes. The mark and reset will always work for 100,000 bytes or less. Be aware that this might cause some streams to buffer that number of bytes into memory.

## Creating Service Clients

To make requests to Amazon Web Services, you first create a service client object. The recommended way is to use the service client builder.

Each AWS service has a service interface with methods for each action in the service API. For example, the service interface for Amazon DynamoDB is named [AmazonDynamoDB](#). Each service interface has a corresponding client builder you can use to construct an implementation of the service interface. The client builder class for DynamoDB is named [AmazonDynamoDBClientBuilder](#).

### Obtaining a Client Builder

To obtain an instance of the client builder, use the static factory method `standard`, as shown in the following example.

```
AmazonDynamoDBClientBuilder builder = AmazonDynamoDBClientBuilder.standard();
```

Once you have a builder, you can customize the client's properties by using many fluent setters in the builder API. For example, you can set a custom region and a custom credentials provider, as follows.

```
AmazonDynamoDB ddb = AmazonDynamoDBClientBuilder.standard()
    .withRegion(Regions.US_WEST_2)
    .withCredentials(new ProfileCredentialsProvider("myProfile"))
    .build();
```

## Note

The fluent `withXXX` methods return the builder object so that you can chain the method calls for convenience and for more readable code. After you configure the properties you want, you can call the `build` method to create the client. Once a client is created, it's immutable and any calls to `setRegion` or `setEndpoint` will fail.

A builder can create multiple clients with the same configuration. When you're writing your application, be aware that the builder is mutable and not thread-safe.

The following code uses the builder as a factory for client instances.

```
public class DynamoDBClientFactory {
    private final AmazonDynamoDBClientBuilder builder =
        AmazonDynamoDBClientBuilder.standard()
            .withRegion(Regions.US_WEST_2)
            .withCredentials(new ProfileCredentialsProvider("myProfile"));

    public AmazonDynamoDB createClient() {
        return builder.build();
    }
}
```

The builder also exposes fluent setters for `ClientConfiguration` and `RequestMetricCollector`, and a custom list of `RequestHandler2`.

The following is a complete example that overrides all configurable properties.

```
AmazonDynamoDB ddb = AmazonDynamoDBClientBuilder.standard()
    .withRegion(Regions.US_WEST_2)
    .withCredentials(new ProfileCredentialsProvider("myProfile"))
    .withClientConfiguration(new ClientConfiguration().withRequestTimeout(5000))
    .withMetricsCollector(new MyCustomMetricsCollector())
    .withRequestHandlers(new MyCustomRequestHandler(), new MyOtherCustomRequestHandler())
    .build();
```

## Creating Async Clients

The AWS SDK for Java has asynchronous (or `async`) clients for every service (except for Amazon S3), and a corresponding `async` client builder for every service.

### To create an `async` DynamoDB client with the default `ExecutorService`

```
AmazonDynamoDBAsync ddbAsync = AmazonDynamoDBAsyncClientBuilder.standard()
    .withRegion(Regions.US_WEST_2)
    .withCredentials(new ProfileCredentialsProvider("myProfile"))
    .build();
```

In addition to the configuration options that the synchronous (or `sync`) client builder supports, the `async` client enables you to set a custom `ExecutorFactory` to change the `ExecutorService` that the `async` client uses. `ExecutorFactory` is a functional interface, so it interoperates with Java 8 lambda expressions and method references.

### To create an `async` client with a custom executor

```
AmazonDynamoDBAsync ddbAsync = AmazonDynamoDBAsyncClientBuilder.standard()
    .withExecutorFactory(() -> Executors.newFixedThreadPool(10))
    .build();
```

## Using DefaultClient

Both the sync and `async` client builders have another factory method named `defaultClient`. This method creates a service client with the default configuration, using the default provider chain to load credentials and the AWS Region. If credentials or the region can't be determined from the environment that the application is running in, the call to `defaultClient` fails. See *Working with AWS Credentials* and *AWS Region Selection* for more information about how credentials and region are determined.

### To create a default service client

```
AmazonDynamoDB ddb = AmazonDynamoDBClientBuilder.defaultClient();
```

## Client Lifecycle

Service clients in the SDK are thread-safe and, for best performance, you should treat them as long-lived objects. Each client has its own connection pool resource that is shut down when the client is garbage collected. To explicitly shut down a client, call the `shutdown` method. After calling `shutdown`, all client resources are released and the client is unusable.

### To shut down a client

```
AmazonDynamoDB ddb = AmazonDynamoDBClientBuilder.defaultClient();
ddb.shutdown();
// Client is now unusable
```

## Working with AWS Credentials

To make requests to Amazon Web Services, you must supply AWS credentials to the AWS SDK for Java. You can do this in the following ways:

- Use the default credential provider chain (*recommended*).
- Use a specific credential provider or provider chain (or create your own).
- Supply the credentials yourself. These can be root account credentials, IAM credentials, or temporary credentials retrieved from AWS STS.

### Important

For security, we *strongly recommend* that you *use IAM users instead of the root account for AWS access*. For more information, see [IAM Best Practices](#) in the *IAM User Guide*.

## Using the Default Credential Provider Chain

When you initialize a new service client without supplying any arguments, the AWS SDK for Java attempts to find AWS credentials by using the *default credential provider chain* implemented by the [DefaultAWSCredentialsProviderChain](#) class. The default credential provider chain looks for credentials in this order:

1. **Environment variables** – AWS\_ACCESS\_KEY\_ID and AWS\_SECRET\_ACCESS\_KEY. The AWS SDK for Java uses the [EnvironmentVariableCredentialsProvider](#) class to load these credentials.
2. **Java system properties** – aws.accessKeyId and aws.secretKey. The AWS SDK for Java uses the [SystemPropertiesCredentialsProvider](#) to load these credentials.
3. **The default credential profiles file** – typically located at `~/.aws/credentials` (location can vary per platform), and shared by many of the AWS SDKs and by the AWS CLI. The AWS SDK for Java uses the [ProfileCredentialsProvider](#) to load these credentials.  
You can create a credentials file by using the `aws configure` command provided by the AWS CLI, or you can create it by editing the file with a text editor. For information about the credentials file format, see [AWS Credentials File Format](#).
4. **Amazon ECS container credentials** – loaded from the Amazon ECS if the environment variable AWS\_CONTAINER\_CREDENTIALS\_RELATIVE\_URI is set. The AWS SDK for Java uses the [ContainerCredentialsProvider](#) to load these credentials.
5. **Instance profile credentials** – used on EC2 instances, and delivered through the Amazon EC2 metadata service. The AWS SDK for Java uses the [InstanceProfileCredentialsProvider](#) to load these credentials.

## Note

Instance profile credentials are used only if `AWS_CONTAINER_CREDENTIALS_RELATIVE_URI` is not set. See [EC2ContainerCredentialsProviderWrapper](#) for more information.

## ***Setting Credentials***

To be able to use AWS credentials, they must be set in *at least one* of the preceding locations. For information about setting credentials, see the following topics:

- To specify credentials in the *environment* or in the default *credential profiles file*, see [Set up AWS Credentials and Region for Development](#).
- To set Java *system properties*, see the [System Properties](#) tutorial on the official *Java Tutorials* website.
- To set up and use *instance profile credentials* with your EC2 instances, see [Using IAM Roles to Grant Access to AWS Resources on Amazon EC2](#).

## ***Setting an Alternate Credentials Profile***

The AWS SDK for Java uses the *default* profile by default, but there are ways to customize which profile is sourced from the credentials file.

You can use the AWS Profile environment variable to change the profile loaded by the SDK.

For example, on Linux, macOS, or Unix you would run the following command to change the profile to *myProfile*.

```
export AWS_PROFILE="myProfile"
```

On Windows you would use the following.

```
set AWS_PROFILE="myProfile"
```

Setting the `AWS_PROFILE` environment variable affects credential loading for all officially supported AWS SDKs and Tools (including the AWS CLI and the AWS CLI for PowerShell). To change only the profile for a Java application, you can use the system property `aws.profile` instead.

## Note

The environment variable takes precedence over the system property.

## ***Setting an Alternate Credentials File Location***

## Using the AWS SDK for Java

The AWS SDK for Java loads AWS credentials automatically from the default credentials file location. However, you can also specify the location by setting the `AWS_CREDENTIAL_PROFILES_FILE` environment variable with the full path to the credentials file.

You can use this feature to temporarily change the location where the AWS SDK for Java looks for your credentials file (for example, by setting this variable with the command line). Or you can set the environment variable in your user or system environment to change it for the user or systemwide.

### To override the default credentials file location

- Set the `AWS_CREDENTIAL_PROFILES_FILE` environment variable to the location of your AWS credentials file.
  - On Linux, macOS, or Unix, use `export`:

```
export AWS_CREDENTIAL_PROFILES_FILE=path/to/credentials_file
```

- On Windows, use `set`:

```
set AWS_CREDENTIAL_PROFILES_FILE=path/to/credentials_file
```

### AWS Credentials File Format

When you use the `aws configure` command to create an AWS credentials file, the command creates a file with the following format.

```
[default]
aws_access_key_id={YOUR_ACCESS_KEY_ID}
aws_secret_access_key={YOUR_SECRET_ACCESS_KEY}

[profile2]
aws_access_key_id={YOUR_ACCESS_KEY_ID}
aws_secret_access_key={YOUR_SECRET_ACCESS_KEY}
```

The profile name is specified in square brackets (for example, `[default]`), followed by the configurable fields in that profile as key-value pairs. You can have multiple profiles in your credentials file, which can be added or edited using `aws configure --profile PROFILE_NAME` to select the profile to configure.

You can specify additional fields, such as `aws_session_token`, `metadata_service_timeout`, and `metadata_service_num_attempts`. These are not configurable with the CLI—you must edit the file by hand if you want to use them. For more information about the configuration file and its available fields, see [Configuring the AWS Command Line Interface](#) in the [AWS CLI User Guide](#).

### Loading Credentials

After you set credentials, you can load them by using the default credential provider chain.

To do this, you instantiate an AWS Service client without explicitly providing credentials to the builder, as follows.

```
AmazonS3 s3Client = AmazonS3ClientBuilder.standard()
    .withRegion(Regions.US_WEST_2)
    .build();
```

## Specifying a Credential Provider or Provider Chain

You can specify a credential provider that is different from the *default* credential provider chain by using the client builder.

You provide an instance of a credentials provider or provider chain to a client builder that takes an [AWSCredentialsProvider](#) interface as input. The following example shows how to use *environment* credentials specifically.

```
AmazonS3 s3Client = AmazonS3ClientBuilder.standard()
    .withCredentials(new EnvironmentVariableCredentialsProvider())
    .build();
```

For the full list of AWS SDK for Java-supplied credential providers and provider chains, see [All Known Implementing Classes](#) in [AWSCredentialsProvider](#).

### Tip

You can use this technique to supply credential providers or provider chains that you create by using your own credential provider that implements the [AWSCredentialsProvider](#) interface, or by subclassing the [AWSCredentialsProviderChain](#) class.

## Explicitly Specifying Credentials

If the default credential chain or a specific or custom provider or provider chain doesn't work for your code, you can set credentials that you supply explicitly. If you've retrieved temporary credentials using AWS STS, use this method to specify the credentials for AWS access.

### To explicitly supply credentials to an AWS client

1. Instantiate a class that provides the [AWSCredentials](#) interface, such as [BasicAWSCredentials](#), and supply it with the AWS access key and secret key you will use for the connection.
2. Create an [AWSStaticCredentialsProvider](#) with the [AWSCredentials](#) object.
3. Configure the client builder with the [AWSStaticCredentialsProvider](#) and build the client.

The following is an example.

```
BasicAWSCredentials awsCreds = new BasicAWSCredentials("access_key_id", "secret_key_id");
AmazonS3 s3Client = AmazonS3ClientBuilder.standard()
    .withCredentials(new AWSStaticCredentialsProvider(awsCreds))
    .build();
```

## Using the AWS SDK for Java

When using *temporary credentials obtained from STS*, create a `BasicSessionCredentials` object, passing it the STS-supplied credentials and session token.

```
BasicSessionCredentials sessionCredentials = new BasicSessionCredentials(  
    session_creds.getAccessKeyId(),  
    session_creds.getSecretAccessKey(),  
    session_creds.getSessionToken());  
  
AmazonS3 s3 = AmazonS3ClientBuilder.standard()  
    .withCredentials(new AWSStaticCredentialsProvider(sessionCredentials))  
    .build();
```

## More Info

- *Sign Up for AWS and Create an IAM User*
- *Set up AWS Credentials and Region for Development*
- *Using IAM Roles to Grant Access to AWS Resources on Amazon EC2*

## AWS Region Selection

Regions enable you to access AWS services that physically reside in a specific geographic area. This can be useful both for redundancy and to keep your data and applications running close to where you and your users will access them.

### Checking for Service Availability in an AWS Region

To see if a particular AWS service is available in a region, use the `isServiceSupported` method on the region that you'd like to use.

```
Region.getRegion(Regions.US_WEST_2)  
    .isServiceSupported(AmazonDynamoDB.ENDPOINT_PREFIX);
```

See the `Regions` class documentation for the regions you can specify, and use the endpoint prefix of the service to query. Each service's endpoint prefix is defined in the service interface. For example, the DynamoDB endpoint prefix is defined in `AmazonDynamoDB`.

### Choosing a Region

Beginning with version 1.4 of the AWS SDK for Java, you can specify a region name and the SDK will automatically choose an appropriate endpoint for you. To choose the endpoint yourself, see Choosing a Specific Endpoint.

To explicitly set a region, we recommend that you use the `Regions` enum. This is an enumeration of all publicly available regions. To create a client with a region from the enum, use the following code.

```
AmazonEC2 ec2 = AmazonEC2ClientBuilder.standard()  
    .withRegion(Regions.US_WEST_2)  
    .build();
```

## Using the AWS SDK for Java

If the region you are attempting to use isn't in the `Regions` enum, you can set the region using a *string* that represents the name of the region.

```
AmazonEC2 ec2 = AmazonEC2ClientBuilder.standard()
    .withRegion("us-west-2")
    .build();
```

### Note

After you build a client with the builder, it's *immutable* and the region *cannot be changed*. If you are working with multiple AWS Regions for the same service, you should create multiple clients—one per region.

## Choosing a Specific Endpoint

Each AWS client can be configured to use a *specific endpoint* within a region by calling the `setEndpoint` method.

For example, to configure the Amazon EC2 client to use the EU (Ireland) Region, use the following code.

```
AmazonEC2 ec2 = new AmazonEC2(myCredentials);
ec2.setEndpoint("https://ec2.eu-west-1.amazonaws.com");
```

See [Regions and Endpoints](#) for the current list of regions and their corresponding endpoints for all AWS services.

## Automatically Determine the AWS Region from the Environment

### Important

This section applies only when using a *client builder* to access AWS services. AWS clients created by using the client constructor will not automatically determine region from the environment and will, instead, use the *default* SDK region (USEast1).

When running on Amazon EC2 or Lambda, you might want to configure clients to use the same region that your code is running on. This decouples your code from the environment it's running in and makes it easier to deploy your application to multiple regions for lower latency or redundancy.

*You must use client builders to have the SDK automatically detect the region your code is running in.*

To use the default credential/region provider chain to determine the region from the environment, use the client builder's `defaultClient` method.

```
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();
```

This is the same as using `standard` followed by `build`.

```
AmazonEC2 ec2 = AmazonEC2ClientBuilder.standard()  
    .build();
```

If you don't explicitly set a region using the `withRegion` methods, the SDK consults the default region provider chain to try and determine the region to use.

### **Default Region Provider Chain**

The following is the region lookup process:

1. Any explicit region set by using `withRegion` or `setRegion` on the builder itself takes precedence over anything else.
2. The `AWS_REGION` environment variable is checked. If it's set, that region is used to configure the client.

#### **Note**

This environment variable is set by the Lambda container.

3. The SDK checks the AWS shared configuration file (usually located at `~/.aws/config`). If the `region` property is present, the SDK uses it.
  - The `AWS_CONFIG_FILE` environment variable can be used to customize the location of the shared config file.
  - The `AWS_PROFILE` environment variable or the `aws.profile` system property can be used to customize the profile that is loaded by the SDK.
4. The SDK attempts to use the Amazon EC2 instance metadata service to determine the region of the currently running Amazon EC2 instance.
5. If the SDK still hasn't found a region by this point, client creation fails with an exception.

When developing AWS applications, a common approach is to use the *shared configuration file* (described in Using the Default Credential Provider Chain) to set the region for local development, and rely on the default region provider chain to determine the region when running on AWS infrastructure. This greatly simplifies client creation and keeps your application portable.

## **Exception Handling**

Understanding how and when the AWS SDK for Java throws exceptions is important to building high-quality applications using the SDK. The following sections describe the different cases of exceptions that are thrown by the SDK and how to handle them appropriately.

### **Why Unchecked Exceptions?**

## Using the AWS SDK for Java

The AWS SDK for Java uses runtime (or unchecked) exceptions instead of checked exceptions for these reasons:

- To allow developers fine-grained control over the errors they want to handle without forcing them to handle exceptional cases they aren't concerned about (and making their code overly verbose)
- To prevent scalability issues inherent with checked exceptions in large applications

In general, checked exceptions work well on small scales, but can become troublesome as applications grow and become more complex.

For more information about the use of checked and unchecked exceptions, see:

- [Unchecked Exceptions—The Controversy](#)
- [The Trouble with Checked Exceptions](#)
- [Java's checked exceptions were a mistake \(and here's what I would like to do about it\)](#)

## AmazonServiceException (and Subclasses)

[AmazonServiceException](#) is the most common exception that you'll experience when using the AWS SDK for Java. This exception represents an error response from an AWS service. For example, if you try to terminate an Amazon EC2 instance that doesn't exist, EC2 will return an error response and all the details of that error response will be included in the [AmazonServiceException](#) that's thrown. For some cases, a subclass of [AmazonServiceException](#) is thrown to allow developers fine-grained control over handling error cases through catch blocks.

When you encounter an [AmazonServiceException](#), you know that your request was successfully sent to the AWS service but couldn't be successfully processed. This can be because of errors in the request's parameters or because of issues on the service side.

[AmazonServiceException](#) provides you with information such as:

- Returned HTTP status code
- Returned AWS error code
- Detailed error message from the service
- AWS request ID for the failed request

[AmazonServiceException](#) also includes information about whether the failed request was the caller's fault (a request with illegal values) or the AWS service's fault (an internal service error).

## AmazonClientException

[AmazonClientException](#) indicates that a problem occurred inside the Java client code, either while trying to send a request to AWS or while trying to parse a response from AWS. An [AmazonClientException](#) is generally more severe than an [AmazonServiceException](#), and indicates a major problem that is preventing the client from making service calls to AWS services. For example, the AWS SDK for Java throws an [AmazonClientException](#) if no network connection is available when you try to call an operation on one of the clients.

## Asynchronous Programming

## Using the AWS SDK for Java

You can use either *synchronous* or *asynchronous* methods to call operations on AWS services. Synchronous methods block your thread's execution until the client receives a response from the service. Asynchronous methods return immediately, giving control back to the calling thread without waiting for a response.

Because an asynchronous method returns before a response is available, you need a way to get the response when it's ready. The AWS SDK for Java provides two ways: *Future objects* and *callback methods*.

### Java Futures

Asynchronous methods in the AWS SDK for Java return a `Future` object that contains the results of the asynchronous operation *in the future*.

Call the `Future isDone()` method to see if the service has provided a response object yet. When the response is ready, you can get the response object by calling the `Future get()` method. You can use this mechanism to periodically poll for the asynchronous operation's results while your application continues to work on other things.

Here is an example of an asynchronous operation that calls a Lambda function, receiving a `Future` that can hold an `InvokeResult` object. The `InvokeResult` object is retrieved only after `isDone()` is true.

```
import com.amazonaws.services.lambda.AWSLambdaAsyncClient;
import com.amazonaws.services.lambda.model.InvokeRequest;
import com.amazonaws.services.lambda.model.InvokeResult;
import java.nio.ByteBuffer;
import java.util.concurrent.Future;
import java.util.concurrent.ExecutionException;

public class InvokeLambdaFunctionAsync
{
    public static void main(String[] args)
    {
        String function_name = "HelloFunction";
        String function_input = "{\"who\":\"AWS SDK for Java\"}";

        AWSLambdaAsync lambda = AWSLambdaAsyncClientBuilder.defaultClient();
        InvokeRequest req = new InvokeRequest()
            .withFunctionName(function_name)
            .withPayload(ByteBuffer.wrap(function_input.getBytes()));

        Future<InvokeResult> future_res = lambda.invokeAsync(req);

        System.out.print("Waiting for future");
        while (future_res.isDone() == false) {
            System.out.print(".");
            try {
                Thread.sleep(1000);
            }
            catch (InterruptedException e) {
                System.err.println("\nThread.sleep() was interrupted!");
                System.exit(1);
            }
        }
    }
}
```

```
try {
    InvokeResult res = future_res.get();
    if (res.getStatusCode() == 200) {
        System.out.println("\nLambda function returned:");
        ByteBuffer response_payload = res.getPayload();
        System.out.println(new String(response_payload.array()));
    }
    else {
        System.out.format("Received a non-OK response from AWS: %d\n",
                          res.getStatusCode());
    }
}
catch (InterruptedException | ExecutionException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}

System.exit(0);
}
```

## Asynchronous Callbacks

In addition to using the Java `Future` object to monitor the status of asynchronous requests, the SDK also enables you to implement a class that uses the `AsyncHandler` interface. `AsyncHandler` provides two methods that are called depending on how the request completed: `onSuccess` and `onError`.

The major advantage of the callback interface approach is that it frees you from having to poll the `Future` object to find out when the request has completed. Instead, your code can immediately start its next activity, and rely on the SDK to call your handler at the right time.

```
import com.amazonaws.services.lambda.AWSLambdaAsync;
import com.amazonaws.services.lambda.AWSLambdaAsyncClientBuilder;
import com.amazonaws.services.lambda.model.InvokeRequest;
import com.amazonaws.services.lambda.model.InvokeResult;
import com.amazonaws.handlers.AsyncHandler;
import java.nio.ByteBuffer;
import java.util.concurrent.Future;

public class InvokeLambdaFunctionCallback
{
    private class AsyncLambdaHandler implements AsyncHandler<InvokeRequest, InvokeResult>
    {
        public void onSuccess(InvokeRequest req, InvokeResult res) {
            System.out.println("\nLambda function returned:");
            ByteBuffer response_payload = res.getPayload();
            System.out.println(new String(response_payload.array()));
            System.exit(0);
        }
    }
}
```

```
public void onError(Exception e) {
    System.out.println(e.getMessage());
    System.exit(1);
}

public static void main(String[] args)
{
    String function_name = "HelloFunction";
    String function_input = "{\"who\":\"AWS SDK for Java\"}";

    AWSLambdaAsync lambda = AWSLambdaAsyncClientBuilder.defaultClient();
    InvokeRequest req = new InvokeRequest()
        .withFunctionName(function_name)
        .withPayload(ByteBuffer.wrap(function_input.getBytes()));

    Future<InvokeResult> future_res = lambda.invokeAsync(req, new AsyncLambdaHandler()

    System.out.print("Waiting for async callback");
    while (!future_res.isDone() && !future_res.isCancelled()) {
        // perform some other tasks...
        try {
            Thread.sleep(1000);
        }
        catch (InterruptedException e) {
            System.err.println("Thread.sleep() was interrupted!");
            System.exit(0);
        }
        System.out.print(".");
    }
}
}
```

## Best Practices

### **Callback Execution**

Your implementation of `AsyncHandler` is executed inside the thread pool owned by the asynchronous client. Short, quickly executed code is most appropriate inside your `AsyncHandler` implementation. Long-running or blocking code inside your handler methods can cause contention for the thread pool used by the asynchronous client, and can prevent the client from executing requests. If you have a long-running task that needs to begin from a callback, have the callback run its task in a new thread or in a thread pool managed by your application.

### **Thread Pool Configuration**

The asynchronous clients in the AWS SDK for Java provide a default thread pool that should work for most applications. You can implement a custom `ExecutorService` and pass it to AWS SDK for Java asynchronous clients for more control over how the thread pools are managed.

## Using the AWS SDK for Java

For example, you could provide an `ExecutorService` implementation that uses a custom `ThreadFactory` to control how threads in the pool are named, or to log additional information about thread usage.

### Amazon S3 Asynchronous Access

The `TransferManager` class in the SDK offers asynchronous support for working with the Amazon S3. `TransferManager` manages asynchronous uploads and downloads, provides detailed progress reporting on transfers, and supports callbacks into different events.

## Logging AWS SDK for Java Calls

The AWS SDK for Java is instrumented with [Apache Commons Logging](#), which is an abstraction layer that enables the use of any one of several logging systems at runtime.

Supported logging systems include the Java Logging Framework and Apache Log4j, among others. This topic shows you how to use Log4j. You can use the SDK's logging functionality without making any changes to your application code.

To learn more about Log4j, see the [Apache website](#).

### Note

This topic focuses on Log4j 1.x. Log4j2 doesn't directly support Apache Commons Logging, but provides an adapter that directs logging calls automatically to Log4j2 using the Apache Commons Logging interface. For more information, see [Commons Logging Bridge](#) in the Log4j2 documentation.

## Download the Log4J JAR

To use Log4j with the SDK, you need to download the Log4j JAR from the Apache website. The SDK doesn't include the JAR. Copy the JAR file to a location that is on your classpath.

Log4j uses a configuration file, `log4j.properties`. Example configuration files are shown below. Copy this configuration file to a directory on your classpath. The Log4j JAR and the `log4j.properties` file don't have to be in the same directory.

The `log4j.properties` configuration file specifies properties such as [logging level](#), where logging output is sent (for example, [to a file](#) or [to the console](#)), and the [format of the output](#). The logging level is the granularity of output that the logger generates. Log4j supports the concept of multiple logging *hierarchies*. The logging level is set independently for each hierarchy. The following two logging hierarchies are available in the AWS SDK for Java:

- `log4j.logger.com.amazonaws`
- `log4j.logger.org.apache.http.wire`

## Setting the Classpath

Both the Log4j JAR and the `log4j.properties` file must be located on your classpath. If you're using [Apache Ant](#), set the classpath in the `path` element in your Ant file. The following example shows a `path` element from the Ant file for the Amazon S3 example included with the SDK.

```
<path id="aws.java.sdk.classpath">
  <fileset dir="..../third-party" includes="**/*.jar"/>
  <fileset dir="..../lib" includes="**/*.jar"/>
  <pathelement location=". "/>
</path>
```

If you're using the Eclipse IDE, you can set the classpath by opening the menu and navigating to **Project | Properties | Java Build Path**.

## Service-Specific Errors and Warnings

We recommend that you always leave the "com.amazonaws" logger hierarchy set to "WARN" to catch any important messages from the client libraries. For example, if the Amazon S3 client detects that your application hasn't properly closed an `InputStream` and could be leaking resources, the S3 client reports it through a warning message to the logs. This also ensures that messages are logged if the client has any problems handling requests or responses.

The following `log4j.properties` file sets the `rootLogger` to `WARN`, which causes warning and error messages from all loggers in the "com.amazonaws" hierarchy to be included. Alternatively, you can explicitly set the `com.amazonaws` logger to `WARN`.

```
log4j.rootLogger=WARN, A1
log4j.appenders.A1=org.apache.log4j.ConsoleAppender
log4j.appenders.A1.layout=org.apache.log4j.PatternLayout
log4j.appenders.A1.layout.ConversionPattern=%d [%t] %-5p %c - %m%n
# Or you can explicitly enable WARN and ERROR messages for the AWS Java clients
log4j.logger.com.amazonaws=WARN
```

## Request/Response Summary Logging

Every request to an AWS service generates a unique AWS request ID that is useful if you run into an issue with how an AWS service is handling a request. AWS request IDs are accessible programmatically through `Exception` objects in the SDK for any failed service call, and can also be reported through the `DEBUG` log level in the "com.amazonaws.request" logger.

The following `log4j.properties` file enables a summary of requests and responses, including AWS request IDs.

```
log4j.rootLogger=WARN, A1
log4j.appenders.A1=org.apache.log4j.ConsoleAppender
log4j.appenders.A1.layout=org.apache.log4j.PatternLayout
log4j.appenders.A1.layout.ConversionPattern=%d [%t] %-5p %c - %m%n
# Turn on DEBUG logging in com.amazonaws.request to log
# a summary of requests/responses with AWS request IDs
log4j.logger.com.amazonaws.request=DEBUG
```

Here is an example of the log output.

```
2009-12-17 09:53:04,269 [main] DEBUG com.amazonaws.request - Sending
Request: POST https://rds.amazonaws.com / Parameters: (MaxRecords: 20,
Action: DescribeEngineDefaultParameters, SignatureMethod: HmacSHA256,
```

```
AWSAccessKeyId: ACCESSKEYID, Version: 2009-10-16, SignatureVersion: 2,
Engine: mysql5.1, Timestamp: 2009-12-17T17:53:04.267Z, Signature:
q963XH63Lcovl5Rr71APlzlye99rmWwT9DfuQaNznkD, ) 2009-12-17 09:53:04,464
[main] DEBUG com.amazonaws.request - Received successful response: 200, AWS
Request ID: 694d1242-cee0-c85e-f31f-5dab1ea18bc6 2009-12-17 09:53:04,469
[main] DEBUG com.amazonaws.request - Sending Request: POST
https://rds.amazonaws.com / Parameters: (ResetAllParameters: true, Action:
ResetDBParameterGroup, SignatureMethod: HmacSHA256, DBParameterGroupName:
java-integ-test-param-group-0000000000000, AWSAccessKeyId: ACCESSKEYID,
Version: 2009-10-16, SignatureVersion: 2, Timestamp:
2009-12-17T17:53:04.467Z, Signature:
9WcfgfPwTobvLVcpvhbrdN7P7l3uH0oviYQ4yZ+TQjsQ=, )

2009-12-17 09:53:04,646 [main] DEBUG com.amazonaws.request - Received
successful response: 200, AWS Request ID:
694d1242-cee0-c85e-f31f-5dab1ea18bc6
```

## Verbose Wire Logging

In some cases, it can be useful to see the exact requests and responses that the AWS SDK for Java sends and receives. You shouldn't enable this logging in production systems because writing out large requests (e.g., a file being uploaded to Amazon S3) or responses can significantly slow down an application. If you really need access to this information, you can temporarily enable it through the Apache HttpClient 4 logger. Enabling the DEBUG level on the `apache.http.wire` logger enables logging for all request and response data.

The following log4j.properties file turns on full wire logging in Apache HttpClient 4 and should only be turned on temporarily because it can have a significant performance impact on your application.

```
log4j.rootLogger=WARN, A1
log4j.appenders.A1=org.apache.log4j.ConsoleAppender
log4j.appenders.A1.layout=org.apache.log4j.PatternLayout
log4j.appenders.A1.layout.ConversionPattern=%d [%t] %-5p %c - %m%n
# Log all HTTP content (headers, parameters, content, etc) for
# all requests and responses. Use caution with this since it can
# be very expensive to log such verbose data!
log4j.logger.org.apache.http.wire=DEBUG
```

## Client Networking Configuration

The AWS SDK for Java enables you to change the default client configuration, which is helpful when you want to:

- Connect to the Internet through proxy
- Change HTTP transport settings, such as connection timeout and request retries
- Specify TCP socket buffer size hints

### Proxy Configuration

When constructing a client object, you can pass in an optional [ClientConfiguration](#) object to customize the client's configuration.

If you connect to the Internet through a proxy server, you'll need to configure your proxy server settings (proxy host, port, and username/password) through the [ClientConfiguration](#) object.

## HTTP Transport Configuration

You can configure several HTTP transport options by using the [ClientConfiguration](#) object. New options are occasionally added; to see the full list of options you can retrieve or set, see the *AWS SDK for Java Reference*.

Each of the configurable values has a default value defined by a constant. For a list of the constant values for [ClientConfiguration](#), see [Constant Field Values](#) in the *AWS SDK for Java Reference*.

### **Local Address**

To set the local address that the HTTP client will bind to, use [ClientConfiguration.setLocalAddress](#).

### **Maximum Connections**

You can set the maximum allowed number of open HTTP connections by using the [ClientConfiguration.setMaxConnections](#) method.

### **Proxy Options**

If you use a proxy with your HTTP connections, you might need to set certain options related to HTTP proxies.

### **Timeouts and Error Handling**

You can set options related to timeouts and handling errors with HTTP connections.

#### **• Connection Timeout**

The connection timeout is the amount of time (in milliseconds) that the HTTP connection will wait to establish a connection before giving up. The default is 50,000 ms.

To set this value yourself, use the [ClientConfiguration.setConnectionTimeout](#) method.

#### **• Connection Time to Live (TTL)**

By default, the SDK will attempt to reuse HTTP connections as long as possible. In failure situations where a connection is established to a server that has been brought out of service, having a finite TTL can help with application recovery. For example, setting a 15 minute TTL will ensure that even if you have a connection established to a server that is experiencing issues, you'll reestablish a connection to a new server within 15 minutes.

To set the HTTP connection TTL, use the [ClientConfiguration.setConnectionTTL](#) method.

#### **• Maximum Error Retries**

You can set the maximum retry count for retriable errors by using the [ClientConfiguration.setMaxErrorRetry](#) method.

## TCP Socket Buffer Size Hints

Advanced users who want to tune low-level TCP parameters can additionally set TCP buffer size hints through the [ClientConfiguration](#) object. The majority of users will never need to tweak these values, but they are provided for advanced users.

Optimal TCP buffer sizes for an application are highly dependent on network and operating system configuration and capabilities. For example, most modern operating systems provide auto-tuning logic for TCP buffer sizes. This can have a big impact on performance for TCP connections that are held open long enough for the auto-tuning to optimize buffer sizes.

Large buffer sizes (e.g., 2 MB) allow the operating system to buffer more data in memory without requiring the remote server to acknowledge receipt of that information, and so can be particularly useful when the network has high latency.

This is only a *hint*, and the operating system might not honor it. When using this option, users should always check the operating system's configured limits and defaults. Most operating systems have a maximum TCP buffer size limit configured, and won't let you go beyond that limit unless you explicitly raise the maximum TCP buffer size limit.

Many resources are available to help with configuring TCP buffer sizes and operating system-specific TCP settings, including the following:

- [TCP Tuning and Network Troubleshooting](#)
- [Host Tuning](#)

## Access Control Policies

AWS *access control policies* enable you to specify fine-grained access controls on your AWS resources. An access control policy consists of a collection of *statements*, which take the form:

*Account A* has permission to perform *action B* on *resource C* where *condition D* applies.

Where:

- *A* is the *principal* – The AWS account that is making a request to access or modify one of your AWS resources.
- *B* is the *action* – The way in which your AWS resource is being accessed or modified, such as sending a message to an Amazon SQS queue, or storing an object in an Amazon S3 bucket.
- *C* is the *resource* – The AWS entity that the principal wants to access, such as an Amazon SQS queue, or an object stored in Amazon S3.
- *D* is a *set of conditions* – The optional constraints that specify when to allow or deny access for the principal to access your resource. Many expressive conditions are available, some specific to each service. For example, you can use date conditions to allow access to your resources only after or before a specific time.

## Amazon S3 Example

The following example demonstrates a policy that allows anyone access to read all the objects in a bucket, but restricts access to uploading objects to that bucket to two specific AWS accounts (in addition to the bucket owner's account).

```
Statement allowPublicReadStatement = new Statement(Effect.Allow)
    .withPrincipals(Principal.AllUsers)
```

## Using the AWS SDK for Java

```
.withActions(S3Actions.GetObject)
.withResources(new S3ObjectResource(myBucketName, "*"));
Statement allowRestrictedWriteStatement = new Statement(Effect.Allow)
    .withPrincipals(new Principal("123456789"), new Principal("876543210"))
    .withActions(S3Actions.PutObject)
    .withResources(new S3ObjectResource(myBucketName, "*"));

Policy policy = new Policy()
    .withStatements(allowPublicReadStatement, allowRestrictedWriteStatement);

AmazonS3 s3 = AmazonS3ClientBuilder.defaultClient();
s3.setBucketPolicy(myBucketName, policy.toJson());
```

## Amazon SQS Example

One common use of policies is to authorize an Amazon SQS queue to receive messages from an Amazon SNS topic.

```
Policy policy = new Policy().withStatements(
    new Statement(Effect.Allow)
        .withPrincipals(Principal.AllUsers)
        .withActions(SQSActions.SendMessage)
        .withConditions(ConditionFactory.newSourceArnCondition(myTopicArn)));

Map queueAttributes = new HashMap();
queueAttributes.put(QueueAttributeName.Policy.toString(), policy.toJson());

AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
sqs.setQueueAttributes(new SetQueueAttributesRequest(myQueueUrl, queueAttributes));
```

## Amazon SNS Example

Some services offer additional conditions that can be used in policies. Amazon SNS provides conditions for allowing or denying subscriptions to SNS topics based on the protocol (e.g., email, HTTP, HTTPS, Amazon SQS) and endpoint (e.g., email address, URL, Amazon SQS ARN) of the request to subscribe to a topic.

```
Condition endpointCondition =
    SNSConditionFactory.newEndpointCondition("*@mycompany.com");

Policy policy = new Policy().withStatements(
    new Statement(Effect.Allow)
        .withPrincipals(Principal.AllUsers)
        .withActions(SNSActions.Subscribe)
        .withConditions(endpointCondition));

AmazonSNS sns = AmazonSNSClientBuilder.defaultClient();
sns.setTopicAttributes(
    new SetTopicAttributesRequest(myTopicArn, "Policy", policy.toJson()));
```

## Setting the JVM TTL for DNS Name Lookups

The Java virtual machine (JVM) caches DNS name lookups. When the JVM resolves a hostname to an IP address, it caches the IP address for a specified period of time, known as the *time-to-live* (TTL).

Because AWS resources use DNS name entries that occasionally change, we recommend that you configure your JVM with a TTL value of no more than 60 seconds. This ensures that when a resource's IP address changes, your application will be able to receive and use the resource's new IP address by requerying the DNS.

On some Java configurations, the JVM default TTL is set so that it will *never* refresh DNS entries until the JVM is restarted. Thus, if the IP address for an AWS resource changes while your application is still running, it won't be able to use that resource until you *manually restart* the JVM and the cached IP information is refreshed. In this case, it's crucial to set the JVM's TTL so that it will periodically refresh its cached IP information.

### Note

The default TTL can vary according to the version of your JVM and whether a [security manager](#) is installed. Many JVMs provide a default TTL less than 60 seconds. If you're using such a JVM and not using a security manager, you can ignore the remainder of this topic.

## How to Set the JVM TTL

To modify the JVM's TTL, set the `networkaddress.cache.ttl` property value. Use one of the following methods, depending on your needs:

- **globally, for all applications that use the JVM.** Set `networkaddress.cache.ttl` in the `$JAVA_HOME/jre/lib/security/java.security` file:

```
networkaddress.cache.ttl=60
```

- **for your application only,** set `networkaddress.cache.ttl` in your application's initialization code:

```
java.security.Security.setProperty("networkaddress.cache.ttl" , "60");
```

# Programming Examples

This section provides tutorials and examples of using the AWS SDK for Java to program AWS services.

## Tip

See Additional Documentation and Resources for more examples and additional resources available for AWS SDK for Java developers!

## SDK Code Samples

The AWS SDK for Java comes packaged with code samples that demonstrate many of the features of the SDK in buildable, runnable programs. You can study or modify these to implement your own AWS solutions using the AWS SDK for Java.

### How to Get the Samples

The AWS SDK for Java code samples are provided in the `samples` directory of the SDK. If you downloaded and installed the SDK using the information in *Set up the AWS SDK for Java*, you already have the samples on your system.

You can also view the latest samples on the AWS SDK for Java GitHub repository, in the `src/samples` directory.

### Building and Running the Samples Using the Command Line

The samples include [Ant](#) build scripts so that you can easily build and run them from the command line. Each sample also contains a README file in HTML format that contains information specific to each sample.

## Tip

If you're browsing the sample code on GitHub, click the **Raw** button in the source code display when viewing the sample's `README.html` file. In raw mode, the HTML will render as intended in your browser.

## Prerequisites

Before running any of the AWS SDK for Java samples, you need to set your AWS credentials in the environment or with the AWS CLI, as specified in *Set up AWS Credentials and Region for Development*. The samples use the default credential provider chain whenever possible. So by setting your credentials in this way, you can avoid the risky practice of inserting your AWS credentials in files within the source code directory (where they may inadvertently be checked in and shared publicly).

## Running the Samples

## Programming Examples

### To run a sample from the command line

1. Change to the directory containing the sample's code. For example, if you're in the root directory of the AWS SDK download and want to run the AwsConsoleApp sample, you would type:

```
cd samples/AwsConsoleApp
```

2. Build and run the sample with Ant. The default build target performs both actions, so you can just enter:

```
ant
```

The sample prints information to standard output—for example:

```
=====
Welcome to the AWS Java SDK!
=====
You have access to 4 Availability Zones.
You have 0 Amazon EC2 instance(s) running.
You have 13 Amazon SimpleDB domain(s) containing a total of 62 items.
You have 23 Amazon S3 bucket(s), containing 44 objects with a total size of 154767691 bytes
```

### Building and Running the Samples Using the Eclipse IDE

If you use the AWS Toolkit for Eclipse, you can also start a new project in Eclipse based on the AWS SDK for Java or add the SDK to an existing Java project.

#### **Prerequisites**

After installing the AWS Toolkit for Eclipse, we recommend configuring the Toolkit with your security credentials. You can do this anytime by choosing **Preferences** from the **Window** menu in Eclipse, and then choosing the **AWS Toolkit** section.

#### **Running the Samples**

##### To run a sample using the AWS Toolkit for Eclipse

1. Open Eclipse.
2. Create a new AWS Java project. In Eclipse, on the **File** menu, choose **New**, and then click **Project**. The **New Project** wizard opens.
3. Expand the **AWS** category, then choose **AWS Java Project**.
4. Choose **Next**. The project settings page is displayed.
5. Enter a name in the **Project Name** box. The AWS SDK for Java Samples group displays the samples available in the SDK, as described previously.

## Programming Examples

6. Select the samples you want to include in your project by selecting each check box.
7. Enter your AWS credentials. If you've already configured the AWS Toolkit for Eclipse with your credentials, this is automatically filled in.
8. Choose **Finish**. The project is created and added to the **Project Explorer**.

### To run the project

1. Choose the sample `.java` file you want to run. For example, for the Amazon S3 sample, choose `S3Sample.java`.
2. Choose **Run** from the **Run** menu.

### To add the SDK to an existing project

1. Right-click the project in **Project Explorer**, point to **Build Path**, and then choose **Add Libraries**.
2. Choose **AWS Java SDK**, choose **Next**, and then follow the remaining on-screen instructions.

## DynamoDB Examples

This section provides examples of programming **DynamoDB** using the **AWS SDK for Java**.

### Note

The examples include only the code needed to demonstrate each technique. The [complete example code is available on GitHub](#). From there, you can download a single source file or clone the repository locally to get all the examples to build and run.

## Working with Tables in DynamoDB

Tables are the containers for all items in a DynamoDB database. Before you can add or remove data from DynamoDB, you must create a table.

For each table, you must define:

- A table *name* that is unique for your account and region.
- A *primary key* for which every value must be unique; no two items in your table can have the same primary key value.

A primary key can be *simple*, consisting of a single partition (HASH) key, or *composite*, consisting of a partition and a sort (RANGE) key.

Each key value has an associated *data type*, enumerated by the [ScalarAttributeType](#) class. The key value can be binary (B), numeric (N), or a string (S). For more information, see [Naming Rules and Data Types](#) in the *DynamoDB Developer Guide*.

- *Provisioned throughput* values that define the number of reserved read/write capacity units for the table.

## Tip

Amazon DynamoDB pricing is based on the provisioned throughput values that you set on your tables, so reserve only as much capacity as you think you'll need for your table.

Provisioned throughput for a table can be modified at any time, so you can adjust capacity if your needs change.

## Create a Table

Use the [DynamoDB client's](#) `createTable` method to create a new DynamoDB table. You need to construct table attributes and a table schema, both of which are used to identify the primary key of your table. You must also supply initial provisioned throughput values and a table name.

## Note

If a table with the name you chose already exists, an [AmazonServiceException](#) is thrown.

## Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.AttributeDefinition;
import com.amazonaws.services.dynamodbv2.model.CreateTableRequest;
import com.amazonaws.services.dynamodbv2.model.CreateTableResult;
import com.amazonaws.services.dynamodbv2.model.KeySchemaElement;
import com.amazonaws.services.dynamodbv2.model.KeyType;
import com.amazonaws.services.dynamodbv2.model.ProvisionedThroughput;
import com.amazonaws.services.dynamodbv2.model.ScalarAttributeType;
```

## Create a Table with a Simple Primary Key

This code creates a table with a simple primary key ("Name").

## Code

```
CreateTableRequest request = new CreateTableRequest()
    .withAttributeDefinitions(new AttributeDefinition(
        "Name", ScalarAttributeType.S))
    .withKeySchema(new KeySchemaElement("Name", KeyType.HASH))
    .withProvisionedThroughput(new ProvisionedThroughput(
        new Long(10), new Long(10)))
    .withTableName(table_name);
```

## Programming Examples

```
final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    CreateTableResult result = ddb.createTable(request);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete sample](#).

*Create a Table with a Composite Primary Key*

Add another `AttributeDefinition` and `KeySchemaElement` to `CreateTableRequest`.

### Code

```
CreateTableRequest request = new CreateTableRequest()
    .withAttributeDefinitions(
        new AttributeDefinition("Language", ScalarAttributeType.S),
        new AttributeDefinition("Greeting", ScalarAttributeType.S))
    .withKeySchema(
        new KeySchemaElement("Language", KeyType.HASH),
        new KeySchemaElement("Greeting", KeyType.RANGE))
    .withProvisionedThroughput(
        new ProvisionedThroughput(new Long(10), new Long(10)))
    .withTableName(table_name);
```

See the [complete sample](#).

### List Tables

You can list the tables in a particular region by calling the `DynamoDB` client's `listTables` method.

#### Note

If the named table doesn't exist for your account and region, a `ResourceNotFoundException` is thrown.

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.ListTablesResult;
import java.util.List;
```

### Code

## Programming Examples

```
final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

boolean more_tables = true;
while(more_tables) {
    String last_name = null;
    try {
        ListTablesResult table_list = null;
        if (last_name == null) {
            table_list = ddb.listTables();
        } else {
            table_list = ddb.listTables(last_name);
        }

        List<String> table_names = table_list.getTableNames();

        if (table_names.size() > 0) {
            for (String cur_name : table_names) {
                System.out.format("* %s\n", cur_name);
            }
        } else {
            System.out.println("No tables found!");
            System.exit(0);
        }
    }

    last_name = table_list.getLastEvaluatedTableName();
    if (last_name == null) {
        more_tables = false;
    }
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
}
```

By default, up to 100 tables are returned per call—use `getLastEvaluatedTableName` on the returned `ListTablesResult` object to get the last table that was evaluated. You can use this value to start the listing after the last returned value of the previous listing.

See the [complete sample](#).

### ***Describe (Get Information about) a Table***

Call the `DynamoDB` client's `describeTable` method.

#### Note

If the named table doesn't exist for your account and region, a `ResourceNotFoundException` is thrown.

## Programming Examples

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.AttributeDefinition;
import com.amazonaws.services.dynamodbv2.model.ProvisionedThroughputDescription;
import com.amazonaws.services.dynamodbv2.model.TableDescription;
import java.util.List;
```

### Code

```
final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    TableDescription table_info =
        ddb.describeTable(table_name).getTable();

    if (table_info != null) {
        System.out.format("Table name : %s\n",
                          table_info.getTableName());
        System.out.format("Table ARN   : %s\n",
                          table_info.getTableArn());
        System.out.format("Status      : %s\n",
                          table_info.getTableStatus());
        System.out.format("Item count  : %d\n",
                          table_info.getItemCount().longValue());
        System.out.format("Size (bytes): %d\n",
                          table_info.getTableSizeBytes().longValue());

        ProvisionedThroughputDescription throughput_info =
            table_info.getProvisionedThroughput();
        System.out.println("Throughput");
        System.out.format("  Read Capacity : %d\n",
                          throughput_info.getReadCapacityUnits().longValue());
        System.out.format("  Write Capacity: %d\n",
                          throughput_info.getWriteCapacityUnits().longValue());

        List<AttributeDefinition> attributes =
            table_info.getAttributeDefinitions();
        System.out.println("Attributes");
        for (AttributeDefinition a : attributes) {
            System.out.format("  %s (%s)\n",
                              a.getAttributeName(), a.getAttributeType());
        }
    }
} catch (AmazonServiceException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}
```

## Programming Examples

See the [complete sample](#).

### **Modify (Update) a Table**

You can modify your table's provisioned throughput values at any time by calling the DynamoDB client's `updateTable` method.

#### Note

If the named table doesn't exist for your account and region, a `ResourceNotFoundException` is thrown.

#### Imports

```
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.ProvisionedThroughput;
import com.amazonaws.AmazonServiceException;
```

#### Code

```
ProvisionedThroughput table_throughput = new ProvisionedThroughput(
    read_capacity, write_capacity);

final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    ddb.updateTable(table_name, table_throughput);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete sample](#).

### **Delete a Table**

Call the DynamoDB client's `deleteTable` method and pass it the table's name.

#### Note

If the named table doesn't exist for your account and region, a `ResourceNotFoundException` is thrown.

#### Imports

## Programming Examples

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
```

### Code

```
final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    ddb.deleteTable(table_name);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete sample](#).

### More Info

- Guidelines for Working with Tables in the *DynamoDB Developer Guide*
- Working with Tables in DynamoDB in the *DynamoDB Developer Guide*

## Working with Items in DynamoDB

In DynamoDB, an item is a collection of *attributes*, each of which has a *name* and a *value*. An attribute value can be a scalar, set, or document type. For more information, see [Naming Rules and Data Types](#) in the *DynamoDB Developer Guide*.

### Retrieve (Get) an Item from a Table

Call the [AmazonDynamoDB](#)'s `getItem` method and pass it a `GetItemRequest` object with the table name and primary key value of the item you want. It returns a `GetItemResult` object.

You can use the returned `GetItemResult` object's `getItem()` method to retrieve a [Map](#) of key (String) and value ([AttributeValue](#)) pairs that are associated with the item.

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.AttributeValue;
import com.amazonaws.services.dynamodbv2.model.GetItemRequest;
import java.util.HashMap;
import java.util.Map;
import java.util.Set;
```

### Code

## Programming Examples

```
HashMap<String,AttributeValue> key_to_get =
    new HashMap<String,AttributeValue>();

key_to_get.put("Name", new AttributeValue(name));

GetItemRequest request = null;
if (projection_expression != null) {
    request = new GetItemRequest()
        .withKey(key_to_get)
        .withTableName(table_name)
        .withProjectionExpression(projection_expression);
} else {
    request = new GetItemRequest()
        .withKey(key_to_get)
        .withTableName(table_name);
}

final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    Map<String,AttributeValue> returned_item =
        ddb.getItem(request).getItem();
    if (returned_item != null) {
        Set<String> keys = returned_item.keySet();
        for (String key : keys) {
            System.out.format("%s: %s\n",
                key, returned_item.get(key).toString());
        }
    } else {
        System.out.format("No item found with the key %s!\n", name);
    }
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete sample](#).

### Add a New Item to a Table

Create a [Map](#) of key-value pairs that represent the item's attributes. These must include values for the table's primary key fields. If the item identified by the primary key already exists, its fields are *updated* by the request.

#### Note

If the named table doesn't exist for your account and region, a [ResourceNotFoundException](#) is thrown.

## Programming Examples

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.AttributeValue;
import com.amazonaws.services.dynamodbv2.model.ResourceNotFoundException;
import java.util.ArrayList;
import java.util.HashMap;
```

### Code

```
HashMap<String,AttributeValue> item_values =
    new HashMap<String,AttributeValue>();

item_values.put("Name", new AttributeValue(name));

for (String[] field : extra_fields) {
    item_values.put(field[0], new AttributeValue(field[1]));
}

final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    ddb.putItem(table_name, item_values);
} catch (ResourceNotFoundException e) {
    System.err.format("Error: The table \"%s\" can't be found.\n", table_name);
    System.err.println("Be sure that it exists and that you've typed its name correctly!");
    System.exit(1);
} catch (AmazonServiceException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}
```

See the [complete sample](#).

### ***Update an Existing Item in a Table***

You can update an attribute for an item that already exists in a table by using the [AmazonDynamoDB's updateItem method](#), providing a table name, primary key value, and a map of fields to update.

### Note

If the named table doesn't exist for your account and region, or if the item identified by the primary key you passed in doesn't exist, a [ResourceNotFoundException](#) is thrown.

### Imports

## Programming Examples

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.dynamodbv2.AmazonDynamoDBClient;
import com.amazonaws.services.dynamodbv2.model.AttributeAction;
import com.amazonaws.services.dynamodbv2.model.AttributeValue;
import com.amazonaws.services.dynamodbv2.model.AttributeValueUpdate;
import com.amazonaws.services.dynamodbv2.model.ResourceNotFoundException;
import java.util.ArrayList;
import java.util.HashMap;
```

### Code

```
HashMap<String,AttributeValue> item_key =
    new HashMap<String,AttributeValue>();

item_key.put("Name", new AttributeValue(name));

HashMap<String,AttributeValueUpdate> updated_values =
    new HashMap<String,AttributeValueUpdate>();

for (String[] field : extra_fields) {
    updated_values.put(field[0], new AttributeValueUpdate(
        new AttributeValue(field[1]), AttributeAction.PUT));
}

final AmazonDynamoDBClient ddb = new AmazonDynamoDBClient();

try {
    ddb.updateItem(table_name, item_key, updated_values);
} catch (ResourceNotFoundException e) {
    System.err.println(e.getMessage());
    System.exit(1);
} catch (AmazonServiceException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}
```

See the [complete sample](#).

### More Info

- Guidelines for Working with Items in the *DynamoDB Developer Guide*
- Working with Items in DynamoDB in the *DynamoDB Developer Guide*

## Managing Tomcat Session State with DynamoDB

Tomcat applications often store session-state data in memory. However, this approach doesn't scale well because once the application grows beyond a single web server, the session state must be shared among servers. A common solution is to set up a dedicated session-state server with MySQL. However, this

## Programming Examples

approach also has drawbacks: you must administer another server, the session-state server is a single point of failure, and the MySQL server itself can cause performance problems.

DynamoDB, which is a NoSQL database store from AWS, avoids these drawbacks by providing an effective solution for sharing session state across web servers.

### **Download the Session Manager**

You can download the session manager from the [aws/aws-dynamodb-session-tomcat](#) project on GitHub. This project also hosts the session manager source code, so you can contribute to the project by sending us pull requests or opening issues.

### **Configure the Session-State Provider**

To use the DynamoDB session-state provider, you must do the following:

1. Configure the Tomcat server to use the provider.
2. Set the security credentials of the provider so that it can access AWS.

#### *Configure a Tomcat Server to Use DynamoDB as the Session-State Server*

Copy `AmazonDynamoDBSessionManagerForTomcat-1.x.x.jar` to the `lib` directory of your Tomcat installation. `AmazonDynamoDBSessionManagerForTomcat-1.x.x.jar` is a complete, standalone JAR that contains all the code and dependencies to run the DynamoDB Tomcat Session Manager.

Edit your server's `context.xml` file to specify `com.amazonaws.services.dynamodb.sessionmanager.DynamoDBSessionManager` as your session manager.

```
<?xml version="1.0" encoding="UTF-8"?>
<Context>
    <WatchedResource>WEB-INF/web.xml</WatchedResource>
    <Manager className="com.amazonaws.services.dynamodb.sessionmanager.DynamoDBSessionManager"
            createIfNotExist="true" />
</Context>
```

#### *Configure Your AWS Security Credentials*

You can specify AWS security credentials for the session manager in multiple ways. They are loaded in the following order of precedence:

1. The `AwsAccessKey` and `AwsSecretKey` attributes of the `Manager` element explicitly provide credentials.
2. The `AwsCredentialsFile` attribute on the `Manager` element specifies a properties file from which to load credentials.

If you don't specify credentials through the `Manager` element, `DefaultAWSCredentialsProviderChain` continues searching for credentials in the following order:

1. Environment variables – `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY`
2. Java system properties – `aws.accessKeyId` and `aws.secretKey`
3. Instance profile credentials delivered through the Amazon EC2 instance metadata service (IMDS)

#### *Configure with Elastic Beanstalk*

## Programming Examples

If you're using the session manager in Elastic Beanstalk, ensure your project has an `.ebextensions` directory at the top level of your output artifact structure. Put the following files in `.ebextensions` directory:

- The `AmazonDynamoDBSessionManagerForTomcat-1.x.x.jar` file
- A `context.xml` file, described previously, to configure the session manager
- A configuration file that copies the JAR into Tomcat's `lib` directory and applies the overridden `context.xml` file.

For more information about customizing Elastic Beanstalk environments, see [AWS Elastic Beanstalk Environment Configuration](#) in the *Elastic Beanstalk Developer Guide*.

If you deploy to Elastic Beanstalk with the AWS Toolkit for Eclipse, you can have the toolkit set up the session manager for you; use the **New AWS Java Web Project** wizard and choose DynamoDB for session management. The AWS Toolkit for Eclipse configures the required files and puts them in the `.ebextensions` directory in the `WebContent` directory of your project. If you have problems finding this directory, be sure you aren't hiding files that begin with a period.

### **Manage Tomcat Session State with DynamoDB**

If the Tomcat server is running on an Amazon EC2 instance that is configured to use IAM roles for EC2 instances, you don't need to specify any credentials in the `context.xml` file. In this case, the AWS SDK for Java uses IAM roles credentials obtained through the instance metadata service (IMDS).

When your application starts, it looks for a DynamoDB table named, by default, **Tomcat\_SessionState**. The table should have a string hash key named "sessionId" (case-sensitive), no range key, and the desired values for `ReadCapacityUnits` and `WriteCapacityUnits`.

We recommend that you create this table before running your application for the first time. If you don't create the table, however, the extension creates it during initialization. See the `context.xml` options in the next section for a list of attributes that configure how the session-state table is created when it doesn't exist.

#### **Tip**

For information about working with DynamoDB tables and provisioned throughput, see the *DynamoDB Developer Guide*.

After the application is configured and the table is created, you can use sessions with any other session provider.

### **Options Specified in `context.xml`**

You can use the following configuration attributes in the `Manager` element of your `context.xml` file:

- `AwsAccessKey` – Access key ID to use.
- `AwsSecretKey` – Secret key to use.
- `AwsCredentialsFile` – A properties file containing `accessKey` and `secretKey` properties with your AWS security credentials.
- `Table` – Optional string attribute. The name of the table used to store session data. The default is **Tomcat\_SessionState**.

## Programming Examples

- *RegionId* – Optional string attribute. The AWS Region in which to use DynamoDB. For a list of available AWS Regions, see [Regions and Endpoints](#) in the *Amazon Web Services General Reference*.
- *Endpoint* – Optional string attribute that, if present, overrides any value set for the *Region* option. This attribute specifies the regional endpoint of the DynamoDB service to use. For a list of available AWS Regions, see [Regions and Endpoints](#) in *Amazon Web Services General Reference*.
- *ReadCapacityUnits* – Optional int attribute. The read capacity units to use if the session manager creates the table. The default is 10.
- *WriteCapacityUnits* – Optional int attribute. The write capacity units to use if the session manager creates the table. The default is 5.
- *CreateIfNotExist* – Optional Boolean attribute. The *CreateIfNotExist* attribute controls whether the session manager autocreates the table if it doesn't exist. The default is true. If this flag is set to false and the table doesn't exist, an exception is thrown during Tomcat startup.

## Troubleshooting

If you encounter issues with the session manager, the first place to look is in *catalina.out*. If you have access to the Tomcat installation, you can go directly to this log file and look for any error messages from the session manager. If you're using Elastic Beanstalk, you can view the environment logs with the AWS Management Console or the AWS Toolkit for Eclipse.

## Limitations

The session manager doesn't support session locking. Therefore, applications that use many concurrent AJAX calls to manipulate session data may not be appropriate for use with the session manager, due to race conditions on session data writes and saves back to the data store.

## Amazon EC2 Examples

This section provides examples of programming [Amazon EC2](#) with the AWS SDK for Java.

### Tutorial: Starting an EC2 Instance

This tutorial demonstrates how to use the AWS SDK for Java to start an EC2 instance.

### Prerequisites

Before you begin, be sure that you have created an AWS account and that you have set up your AWS credentials. For more information, see [Getting Started](#).

#### Create an Amazon EC2 Security Group

Create a *security group*, which acts as a virtual firewall that controls the network traffic for one or more EC2 instances. By default, Amazon EC2 associates your instances with a security group that allows no inbound traffic. You can create a security group that allows your EC2 instances to accept certain traffic. For example, if you need to connect to a Linux instance, you must configure the security group to allow SSH traffic. You can create a security group using the Amazon EC2 console or the AWS SDK for Java.

You create a security group for use in either EC2-Classic or EC2-VPC. For more information about EC2-Classic and EC2-VPC, see [Supported Platforms](#) in the *Amazon EC2 User Guide for Linux Instances*.

For more information about creating a security group using the Amazon EC2 console, see [Amazon EC2 Security Groups](#) in the *Amazon EC2 User Guide for Linux Instances*.

### To create a security group

1. Create and initialize a `CreateSecurityGroupRequest` instance. Use the `withGroupName` method to set the security group name, and the `withDescription` method to set the security group description, as follows:

```
CreateSecurityGroupRequest csgr = new CreateSecurityGroupRequest();
csgr.withGroupName("JavaSecurityGroup").withDescription("My security group");
```

The security group name must be unique within the AWS region in which you initialize your Amazon EC2 client. You must use US-ASCII characters for the security group name and description.

2. Pass the request object as a parameter to the `createSecurityGroup` method. The method returns a `CreateSecurityGroupResult` object, as follows:

```
CreateSecurityGroupResult createSecurityGroupResult =
    amazonEC2Client.createSecurityGroup(createSecurityGroupRequest);
```

If you attempt to create a security group with the same name as an existing security group, `createSecurityGroup` throws an exception.

By default, a new security group does not allow any inbound traffic to your Amazon EC2 instance. To allow inbound traffic, you must explicitly authorize security group ingress. You can authorize ingress for individual IP addresses, for a range of IP addresses, for a specific protocol, and for TCP/UDP ports.

### To authorize security group ingress

1. Create and initialize an `IpPermission` instance. Use the `withIpv4Ranges` method to set the range of IP addresses to authorize ingress for, and use the `withIpProtocol` method to set the IP protocol. Use the `withFromPort` and `withToPort` methods to specify range of ports to authorize ingress for, as follows:

```
IpPermission ipPermission =
    new IpPermission();

ipPermission.withIpRanges("111.111.111.111/32", "150.150.150.150/32")
    .withIpProtocol("tcp")
    .withFromPort(22)
    .withToPort(22);
```

All the conditions that you specify in the `IpPermission` object must be met in order for ingress to be allowed.

Specify the IP address using CIDR notation. If you specify the protocol as TCP/UDP, you must provide a source port and a destination port. You can authorize ports only if you specify TCP or UDP.

2. Create and initialize an `AuthorizeSecurityGroupIngressRequest` instance. Use the `withGroupName` method to specify the security group name, and pass the `IpPermission` object you initialized earlier to the `withIpPermissions` method, as follows:

## Programming Examples

```
AuthorizeSecurityGroupIngressRequest authorizeSecurityGroupIngressRequest =  
    new AuthorizeSecurityGroupIngressRequest();  
  
authorizeSecurityGroupIngressRequest.withGroupName("JavaSecurityGroup")  
    .withIpPermissions(ipPermission);
```

3. Pass the request object into the `authorizeSecurityGroupIngress` method, as follows:

```
amazonEC2Client.authorizeSecurityGroupIngress(authorizeSecurityGroupIngressRequest);
```

If you call `authorizeSecurityGroupIngress` with IP addresses for which ingress is already authorized, the method throws an exception. Create and initialize a new `IpPermission` object to authorize ingress for different IPs, ports, and protocols before calling `AuthorizeSecurityGroupIngress`.

Whenever you call the `authorizeSecurityGroupIngress` or `authorizeSecurityGroupEgress` methods, a rule is added to your security group.

### Create a Key Pair

You must specify a key pair when you launch an EC2 instance and then specify the private key of the key pair when you connect to the instance. You can create a key pair or use an existing key pair that you've used when launching other instances. For more information, see [Amazon EC2 Key Pairs](#) in the *Amazon EC2 User Guide for Linux Instances*.

### To create a key pair and save the private key

1. Create and initialize a `CreateKeyPairRequest` instance. Use the `withKeyName` method to set the key pair name, as follows:

```
CreateKeyPairRequest createKeyPairRequest = new CreateKeyPairRequest();  
  
createKeyPairRequest.withKeyName(keyName);
```

### Important

Key pair names must be unique. If you attempt to create a key pair with the same key name as an existing key pair, you'll get an exception.

2. Pass the request object to the `createKeyPair` method. The method returns a `CreateKeyPairResult` instance, as follows:

```
CreateKeyPairResult createKeyPairResult =  
    amazonEC2Client.createKeyPair(createKeyPairRequest);
```

## Programming Examples

3. Call the result object's `getKeyPair` method to obtain a `KeyPair` object. Call the `KeyPair` object's `getKeyMaterial` method to obtain the unencrypted PEM-encoded private key, as follows:

```
KeyPair keyPair = new KeyPair();
keyPair = createKeyPairResult.getKeyPair();
String privateKey = keyPair.getKeyMaterial();
```

### Run an Amazon EC2 Instance

Use the following procedure to launch one or more identically configured EC2 instances from the same Amazon Machine Image (AMI). After you create your EC2 instances, you can check their status. After your EC2 instances are running, you can connect to them.

#### To launch an Amazon EC2 instance

1. Create and initialize a `RunInstancesRequest` instance. Make sure that the AMI, key pair, and security group that you specify exist in the region that you specified when you created the client object.

```
RunInstancesRequest runInstancesRequest =
    new RunInstancesRequest();

runInstancesRequest.withImageId("ami-4b814f22")
    .withInstanceType("m1.small")
    .withMinCount(1)
    .withMaxCount(1)
    .withKeyName("my-key-pair")
    .withSecurityGroups("my-security-group");
```

#### `withImageId`

The ID of the AMI. For a list of public AMIs provided by Amazon, see Amazon Machine Images.

#### `withInstanceType`

An instance type that is compatible with the specified AMI. For more information, see [Instance Types in the Amazon EC2 User Guide for Linux Instances](#).

#### `withMinCount`

The minimum number of EC2 instances to launch. If this is more instances than Amazon EC2 can launch in the target Availability Zone, Amazon EC2 launches no instances.

#### `withMaxCount`

The maximum number of EC2 instances to launch. If this is more instances than Amazon EC2 can launch in the target Availability Zone, Amazon EC2 launches the largest possible number of instances above `MinCount`. You can launch between 1 and the maximum number of instances you're allowed for the instance type. For more information, see [How many instances can I run in Amazon EC2 in the Amazon EC2 General FAQ](#).

#### `withKeyName`

The name of the EC2 key pair. If you launch an instance without specifying a key pair, you can't connect to it. For more information, see [Create a Key Pair](#).

### **withSecurityGroups**

One or more security groups. For more information, see *Create an Amazon EC2 Security Group*.

2. Launch the instances by passing the request object to the `runInstances` method. The method returns a `RunInstancesResult` object, as follows:

```
RunInstancesResult result = amazonEC2Client.runInstances(  
    runInstancesRequest);
```

After your instance is running, you can connect to it using your key pair. For more information, see [Connect to Your Linux Instance](#). in the *Amazon EC2 User Guide for Linux Instances*.

## **Using IAM Roles to Grant Access to AWS Resources on Amazon EC2**

All requests to Amazon Web Services (AWS) must be cryptographically signed using credentials issued by AWS. You can use *IAM roles* to conveniently grant secure access to AWS resources from your Amazon EC2 instances.

This topic provides information about how to use IAM roles with Java SDK applications running on Amazon EC2. For more information about IAM instances, see [IAM Roles for Amazon EC2](#) in the *Amazon EC2 User Guide for Linux Instances*.

### ***The default provider chain and EC2 instance profiles***

If your application creates an AWS client using the default constructor, then the client will search for credentials using the *default credentials provider chain*, in the following order:

1. In system environment variables: `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY`.
2. In the Java system properties: `aws.accessKeyId` and `aws.secretKey`.
3. In the default credentials file (the location of this file varies by platform).
4. In the *instance profile credentials*, which exist within the instance metadata associated with the IAM role for the EC2 instance.

The final step in the default provider chain is available only when running your application on an Amazon EC2 instance, but provides the greatest ease of use and best security when working with Amazon EC2 instances. You can also pass an `InstanceProfileCredentialsProvider` instance directly to the client constructor to get instance profile credentials without proceeding through the entire default provider chain.

For example:

```
AmazonS3 s3 = AmazonS3ClientBuilder.standard()  
    .withCredentials(new InstanceProfileCredentialsProvider())  
    .build();
```

When using this approach, the SDK retrieves temporary AWS credentials that have the same permissions as those associated with the IAM role associated with the Amazon EC2 instance in its instance profile. Although these credentials are temporary and would eventually expire, `InstanceProfileCredentialsProvider` periodically refreshes them for you so that the obtained credentials continue to allow access to AWS.

## Important

The automatic credentials refresh happens *only* when you use the default client constructor, which creates its own `InstanceProfileCredentialsProvider` as part of the default provider chain, or when you pass an `InstanceProfileCredentialsProvider` instance directly to the client constructor. If you use another method to obtain or pass instance profile credentials, you are responsible for checking for and refreshing expired credentials.

If the client constructor can't find credentials using the credentials provider chain, it will throw an `AmazonClientException`.

### **Walkthrough: Using IAM roles for EC2 instances**

The following walkthrough shows you how to retrieve an object from Amazon S3 using an IAM role to manage access.

#### *Create an IAM Role*

Create an IAM role that grants read-only access to Amazon S3.

#### **To create the IAM role**

1. Open the [IAM console](#).
2. In the navigation pane, select **Roles**, then **Create New Role**.
3. Enter a name for the role, then select **Next Step**. Remember this name, since you'll need it when you launch your Amazon EC2 instance.
4. On the **Select Role Type** page, under **AWS Service Roles**, select **Amazon EC2**.
5. On the **Set Permissions** page, under **Select Policy Template**, select **Amazon S3 Read Only Access**, then **Next Step**.
6. On the **Review** page, select **Create Role**.

#### *Launch an EC2 Instance and Specify Your IAM Role*

You can launch an Amazon EC2 instance with an IAM role using the Amazon EC2 console or the AWS SDK for Java.

## Programming Examples

- To launch an Amazon EC2 instance using the console, follow the directions in [Getting Started with Amazon EC2 Linux Instances](#) in the *Amazon EC2 User Guide for Linux Instances*.

When you reach the **Review Instance Launch** page, select **Edit instance details**. In **IAM role**, choose the IAM role that you created previously. Complete the procedure as directed.

### Note

You'll need to create or use an existing security group and key pair to connect to the instance.

- To launch an Amazon EC2 instance with an IAM role using the AWS SDK for Java, see [Run an Amazon EC2 Instance](#).

### Create your Application

Let's build the sample application to run on the EC2 instance. First, create a directory that you can use to hold your tutorial files (for example, `GetS3ObjectApp`).

Next, copy the AWS SDK for Java libraries into your newly-created directory. If you downloaded the AWS SDK for Java to your `~/Downloads` directory, you can copy them using the following commands:

```
cp -r ~/Downloads/aws-java-sdk-{1.7.5}/lib .
cp -r ~/Downloads/aws-java-sdk-{1.7.5}/third-party .
```

Open a new file, call it `GetS3Object.java`, and add the following code:

```
import java.io.*;

import com.amazonaws.auth.*;
import com.amazonaws.services.s3.*;
import com.amazonaws.services.s3.model.*;
import com.amazonaws.AmazonClientException;
import com.amazonaws.AmazonServiceException;

public class GetS3Object {
    private static String bucketName = "text-content";
    private static String key = "text-object.txt";

    public static void main(String[] args) throws IOException
    {
        AmazonS3 s3Client = AmazonS3ClientBuilder.defaultClient();

        try {
            System.out.println("Downloading an object");
            S3Object s3object = s3Client.getObject(
                new GetObjectRequest(bucketName, key));
            displayTextInputStream(s3object.getObjectContent());
        }
    }
}
```

## Programming Examples

```
        catch(AmazonServiceException ase) {
            System.err.println("Exception was thrown by the service");
        }
        catch(AmazonClientException ace) {
            System.err.println("Exception was thrown by the client");
        }
    }

    private static void displayTextInputStream(InputStream input) throws IOException
    {
        // Read one text line at a time and display.
        BufferedReader reader = new BufferedReader(new InputStreamReader(input));
        while(true)
        {
            String line = reader.readLine();
            if(line == null) break;
            System.out.println( "      " + line );
        }
        System.out.println();
    }
}
```

Open a new file, call it build.xml, and add the following lines:

```
<project name="Get Amazon S3 Object" default="run" basedir=".">
    <path id="aws.java.sdk.classpath">
        <fileset dir=".lib" includes="**/*.jar"/>
        <fileset dir=".third-party" includes="**/*.jar"/>
        <pathelement location="lib"/>
        <pathelement location="."/>
    </path>

    <target name="build">
        <javac debug="true"
            includeantruntime="false"
            srcdir="."
            destdir="."
            classpathref="aws.java.sdk.classpath"/>
    </target>

    <target name="run" depends="build">
        <java classname="GetS3Object" classpathref="aws.java.sdk.classpath" fork="true"/>
    </target>
</project>
```

Build and run the modified program. Note that there are no credentials are stored in the program. Therefore, unless you have your AWS credentials specified already, the code will throw `AmazonServiceException`. For example:

## Programming Examples

```
$ ant
Buildfile: /path/to/my/GetS3ObjectApp/build.xml

build:
[javac] Compiling 1 source file to /path/to/my/GetS3ObjectApp

run:
[java] Downloading an object
[java] AmazonServiceException

BUILD SUCCESSFUL
```

### Transfer the Compiled Program to Your EC2 Instance

Transfer the program to your Amazon EC2 instance using secure copy (**scp**), along with the AWS SDK for Java libraries. The sequence of commands looks something like the following.

```
scp -p -i {my-key-pair}.pem GetS3Object.class ec2-user@{public_dns}:GetS3Object.class
scp -p -i {my-key-pair}.pem build.xml ec2-user@{public_dns}:build.xml
scp -r -p -i {my-key-pair}.pem lib ec2-user@{public_dns}:lib
scp -r -p -i {my-key-pair}.pem third-party ec2-user@{public_dns}:third-party
```

### Note

Depending on the Linux distribution that you used, the *user name* might be "ec2-user", "root", or "ubuntu". To get the public DNS name of your instance, open the [EC2 console](#) and look for the **Public DNS** value in the **Description** tab (for example, ec2-198-51-100-1.compute-1.amazonaws.com).

In the preceding commands:

- **GetS3Object.class** is your compiled program
- **build.xml** is the ant file used to build and run your program
- the **lib** and **third-party** directories are the corresponding library folders from the AWS SDK for Java.
- The **-r** switch indicates that **scp** should do a recursive copy of all of the contents of the **library** and **third-party** directories in the AWS SDK for Java distribution.
- The **-p** switch indicates that **scp** should preserve the permissions of the source files when it copies them to the destination.

### Tip

The **-p** switch works only on Linux, macOS, or Unix. If you are copying files from Windows, you may need to fix the file permissions on your instance using the following command:

## Programming Examples

```
chmod -R u+rwx GetS3Object.class build.xml lib third-party
```

### Run the Sample Program on the EC2 Instance

To run the program, connect to your Amazon EC2 instance. For more information, see [Connect to Your Linux Instance](#) in the *Amazon EC2 User Guide for Linux Instances*.

If **ant** is not available on your instance, install it using the following command:

```
sudo yum install ant
```

Then, run the program using **ant** as follows:

```
ant run
```

The program will write the contents of your Amazon S3 object to your command window.

## Tutorial: Amazon EC2 Spot Instances

### Overview

Spot Instances allow you to bid on unused Amazon Elastic Compute Cloud (Amazon EC2) capacity and run the acquired instances for as long as your bid exceeds the current *Spot Price*. Amazon EC2 changes the Spot Price periodically based on supply and demand, and customers whose bids meet or exceed it gain access to the available Spot Instances. Like On-Demand Instances and Reserved Instances, Spot Instances provide you another option for obtaining more compute capacity.

Spot Instances can significantly lower your Amazon EC2 costs for batch processing, scientific research, image processing, video encoding, data and web crawling, financial analysis, and testing. Additionally, Spot Instances give you access to large amounts of additional capacity in situations where the need for that capacity is not urgent.

To use Spot Instances, place a Spot Instance request specifying the maximum price you are willing to pay per instance hour; this is your bid. If your bid exceeds the current Spot Price, your request is fulfilled and your instances will run until either you choose to terminate them or the Spot Price increases above your bid (whichever is sooner).

It's important to note:

- You will often pay less per hour than your bid. Amazon EC2 adjusts the Spot Price periodically as requests come in and available supply changes. Everyone pays the same Spot Price for that period regardless of whether their bid was higher. Therefore, you might pay less than your bid, but you will never pay more than your bid.
- If you're running Spot Instances and your bid no longer meets or exceeds the current Spot Price, your instances will be terminated. This means that you will want to make sure that your workloads and applications are flexible enough to take advantage of this opportunistic capacity.

Spot Instances perform exactly like other Amazon EC2 instances while running, and like other Amazon EC2 instances, Spot Instances can be terminated when you no longer need them. If you terminate your instance,

## Programming Examples

you pay for any partial hour used (as you would for On-Demand or Reserved Instances). However, if the Spot Price goes above your bid and your instance is terminated by Amazon EC2, you will not be charged for any partial hour of usage.

This tutorial shows how to use AWS SDK for Java to do the following.

- Submit a Spot Request
- Determine when the Spot Request becomes fulfilled
- Cancel the Spot Request
- Terminate associated instances

### **Prerequisites**

To use this tutorial you must have the AWS SDK for Java installed, as well as having met its basic installation prerequisites. See *Set up the AWS SDK for Java* for more information.

### **Step 1: Setting Up Your Credentials**

To begin using this code sample, you need to add AWS credentials to the `AwsCredentials.properties` file as follows:

1. Open the `AwsCredentials.properties` file.
2. Set your access key / secret key id combination in the `AwsCredentials.properties` file.

#### **Note**

We recommend that you use the credentials of an IAM user to provide these values. For more information, see *Sign Up for AWS and Create an IAM User*.

Now that you have configured your settings, you can get started using the code in the example.

### **Step 2: Setting Up a Security Group**

A *security group* acts as a firewall that controls the traffic allowed in and out of a group of instances. By default, an instance is started without any security group, which means that all incoming IP traffic, on any TCP port will be denied. So, before submitting our Spot Request, we will set up a security group that allows the necessary network traffic. For the purposes of this tutorial, we will create a new security group called "GettingStarted" that allows Secure Shell (SSH) traffic from the IP address where you are running your application from. To set up a new security group, you need to include or run the following code sample that sets up the security group programmatically.

After we create an `AmazonEC2` client object, we create a `CreateSecurityGroupRequest` object with the name, "GettingStarted" and a description for the security group. Then we call the `ec2.createSecurityGroup` API to create the group.

To enable access to the group, we create an `ipPermission` object with the IP address range set to the CIDR representation of the subnet for the local computer; the "/10" suffix on the IP address indicates the subnet for the specified IP address. We also configure the `ipPermission` object with the TCP protocol and port 22

## Programming Examples

(SSH). The final step is to call `ec2.authorizeSecurityGroupIngress` with the name of our security group and the `ipPermission` object.

```
<?dbhtml linenumbering.everyNth="1" ?>
// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Create a new security group.
try {
    CreateSecurityGroupRequest securityGroupRequest = new CreateSecurityGroupRequest("GettingStartedGroup");
    ec2.createSecurityGroup(securityGroupRequest);
} catch (AmazonServiceException ase) {
    // Likely this means that the group is already created, so ignore.
    System.out.println(ase.getMessage());
}

String ipAddr = "0.0.0.0/0";

// Get the IP of the current host, so that we can limit the Security
// Group by default to the ip range associated with your subnet.
try {
    InetAddress addr = InetAddress.getLocalHost();

    // Get IP Address
    ipAddr = addr.getHostAddress() + "/10";
} catch (UnknownHostException e) {
}

// Create a range that you would like to populate.
ArrayList<String> ipRanges = new ArrayList<String>();
ipRanges.add(ipAddr);

// Open up port 22 for TCP traffic to the associated IP
// from above (e.g. ssh traffic).
ArrayList<IpPermission> ipPermissions = new ArrayList<IpPermission>();
IpPermission ipPermission = new IpPermission();
ipPermission.setIpProtocol("tcp");
ipPermission.setFromPort(new Integer(22));
ipPermission.setToPort(new Integer(22));
ipPermission.setIpRanges(ipRanges);
ipPermissions.add(ipPermission);

try {
    // Authorize the ports to the used.
    AuthorizeSecurityGroupIngressRequest ingressRequest =
        new AuthorizeSecurityGroupIngressRequest("GettingStartedGroup", ipPermissions);
    ec2.authorizeSecurityGroupIngress(ingressRequest);
} catch (AmazonServiceException ase) {
    // Ignore because this likely means the zone has
    // already been authorized.
```

## Programming Examples

```
System.out.println(ase.getMessage());  
}
```

You can view this entire code sample in the [CreateSecurityGroupApp.java](#) code sample. Note you only need to run this application once to create a new security group.

You can also create the security group using the AWS Toolkit for Eclipse. See [Managing Security Groups from AWS Explorer](#) for more information.

### Step 3: Submitting Your Spot Request

To submit a Spot request, you first need to determine the instance type, Amazon Machine Image (AMI), and maximum bid price you want to use. You must also include the security group we configured previously, so that you can log into the instance if desired.

There are several instance types to choose from; go to Amazon EC2 Instance Types for a complete list. For this tutorial, we will use t1.micro, the cheapest instance type available. Next, we will determine the type of AMI we would like to use. We'll use ami-8c1fce5, the most up-to-date Amazon Linux AMI available when we wrote this tutorial. The latest AMI may change over time, but you can always determine the latest version AMI by following these steps:

1. Log into the AWS Management Console, click the **EC2** tab, and, from the EC2 Console Dashboard, attempt to launch an instance.



AWS Management Console to launch an instance

2. In the window that displays AMIs, just use the AMI ID as shown in the following screen shot. Alternatively, you can use the [DescribeImages API](#), but leveraging that command is outside the scope of this tutorial.



Identifying the most-recent AMI

There are many ways to approach bidding for Spot instances; to get a broad overview of the various approaches you should view the [Bidding for Spot Instances](#) video. However, to get started, we'll describe three common strategies: bid to ensure cost is less than on-demand pricing; bid based on the value of the resulting computation; bid so as to acquire computing capacity as quickly as possible.

- *Reduce Cost below On-Demand* You have a batch processing job that will take a number of hours or days to run. However, you are flexible with respect to when it starts and when it completes. You want to see if you can complete it for less cost than with On-Demand Instances. You examine the Spot Price history for instance types using either the AWS Management Console or the Amazon EC2 API. For more information, go to [Viewing Spot Price History](#). After you've analyzed the price history for your desired instance type in a given Availability Zone, you have two alternative approaches for your bid:
  - You could bid at the upper end of the range of Spot Prices (which are still below the On-Demand price), anticipating that your one-time Spot request would most likely be fulfilled and run for enough consecutive compute time to complete the job.
  - Or, you could bid at the lower end of the price range, and plan to combine many instances launched over time through a persistent request. The instances would run long enough--in

## Programming Examples

aggregate--to complete the job at an even lower total cost. (We will explain how to automate this task later in this tutorial.)

- *Pay No More than the Value of the Result* You have a data processing job to run. You understand the value of the job's results well enough to know how much they are worth in terms of computing costs. After you've analyzed the Spot Price history for your instance type, you choose a bid price at which the cost of the computing time is no more than the value of the job's results. You create a persistent bid and allow it to run intermittently as the Spot Price fluctuates at or below your bid.
- *Acquire Computing Capacity Quickly* You have an unanticipated, short-term need for additional capacity that is not available through On-Demand Instances. After you've analyzed the Spot Price history for your instance type, you bid above the highest historical price to provide a high likelihood that your request will be fulfilled quickly and continue computing until it completes.

After you choose your bid price, you are ready to request a Spot Instance. For the purposes of this tutorial, we will bid the On-Demand price (\$0.03) to maximize the chances that the bid will be fulfilled. You can determine the types of available instances and the On-Demand prices for instances by going to Amazon EC2 Pricing page. To request a Spot Instance, you simply need to build your request with the parameters you chose earlier. We start by creating a `RequestSpotInstancesRequest` object. The request object requires the number of instances you want to start and the bid price. Additionally, you need to set the `LaunchSpecification` for the request, which includes the instance type, AMI ID, and security group you want to use. Once the request is populated, you call the `requestSpotInstances` method on the `AmazonEC2Client` object. The following example shows how to request a Spot Instance.

```
// Retrieves the credentials from a AWS Credentials.properties file.
AWSCredentials credentials = null;
try {
    credentials = new PropertiesCredentials(
        GettingStartedApp.class.getResourceAsStream("AwsCredentials.properties"));
} catch (IOException e1) {
    System.out.println("Credentials were not properly entered into AwsCredentials.properties");
    System.out.println(e1.getMessage());
    System.exit(-1);
}

// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 1 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(1));

// Setup the specifications of the launch. This includes the
// instance type (e.g. t1.micro) and the latest Amazon Linux
// AMI id available. Note, you should always use the latest
// Amazon Linux AMI id or another of your choosing.
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fce5");
```

## Programming Examples

```
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Add the launch specifications to the request.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstance API.
RequestSpotInstancesResult requestResult = ec2.requestSpotInstances(requestRequest);
```

Running this code will launch a new Spot Instance Request. There are other options you can use to configure your Spot Requests. To learn more, please visit [Tutorial: Advanced Amazon EC2 Spot Request Management](#) or the [RequestSpotInstances](#) class in the [AWS SDK for Java Reference](#).

### Note

You will be charged for any Spot Instances that are actually launched, so make sure that you cancel any requests and terminate any instances you launch to reduce any associated fees.

### **Step 4: Determining the State of Your Spot Request**

Next, we want to create code to wait until the Spot request reaches the "active" state before proceeding to the last step. To determine the state of our Spot request, we poll the [describeSpotInstanceRequests](#) method for the state of the Spot request ID we want to monitor.

The request ID created in Step 2 is embedded in the response to our [requestSpotInstances](#) request. The following example code shows how to gather request IDs from the [requestSpotInstances](#) response and use them to populate an [ArrayList](#).

```
// Call the RequestSpotInstance API.
RequestSpotInstancesResult requestResult = ec2.requestSpotInstances(requestRequest);
List<SpotInstanceRequest> requestResponses = requestResult.getSpotInstanceRequests();

// Setup an arraylist to collect all of the request ids we want to
// watch hit the running state.
ArrayList<String> spotInstanceRequestIds = new ArrayList<String>();

// Add all of the request ids to the hashset, so we can determine when they hit the
// active state.
for (SpotInstanceRequest requestResponse : requestResponses) {
    System.out.println("Created Spot Request: "+requestResponse.getSpotInstanceRequestId());
    spotInstanceRequestIds.add(requestResponse.getSpotInstanceRequestId());
}
```

## Programming Examples

To monitor your request ID, call the `describeSpotInstanceRequests` method to determine the state of the request. Then loop until the request is not in the "open" state. Note that we monitor for a state of not "open", rather a state of, say, "active", because the request can go straight to "closed" if there is a problem with your request arguments. The following code example provides the details of how to accomplish this task.

```
// Create a variable that will track whether there are any
// requests still in the open state.
boolean anyOpen;

do {
    // Create the describeRequest object with all of the request ids
    // to monitor (e.g. that we started).
    DescribeSpotInstanceRequestsRequest describeRequest = new DescribeSpotInstanceRequestsRequest();
    describeRequest.setSpotInstanceRequestIds(spotInstanceRequestIds);

    // Initialize the anyOpen variable to false - which assumes there
    // are no requests open unless we find one that is still open.
    anyOpen=false;

    try {
        // Retrieve all of the requests we want to monitor.
        DescribeSpotInstanceRequestsResult describeResult = ec2.describeSpotInstanceRequests();
        List<SpotInstanceRequest> describeResponses = describeResult.getSpotInstanceRequestList();

        // Look through each request and determine if they are all in
        // the active state.
        for (SpotInstanceRequest describeResponse : describeResponses) {
            // If the state is open, it hasn't changed since we attempted
            // to request it. There is the potential for it to transition
            // almost immediately to closed or cancelled so we compare
            // against open instead of active.
            if (describeResponse.getState().equals("open")) {
                anyOpen = true;
                break;
            }
        }
    } catch (AmazonServiceException e) {
        // If we have an exception, ensure we don't break out of
        // the loop. This prevents the scenario where there was
        // a blip on the wire.
        anyOpen = true;
    }

    try {
        // Sleep for 60 seconds.
        Thread.sleep(60*1000);
    } catch (Exception e) {
        // Do nothing because it woke up early.
    }
} while (anyOpen);
```

## Programming Examples

After running this code, your Spot Instance Request will have completed or will have failed with an error that will be output to the screen. In either case, we can proceed to the next step to clean up any active requests and terminate any running instances.

### **Step 5: Cleaning Up Your Spot Requests and Instances**

Lastly, we need to clean up our requests and instances. It is important to both cancel any outstanding requests and terminate any instances. Just canceling your requests will not terminate your instances, which means that you will continue to pay for them. If you terminate your instances, your Spot requests may be canceled, but there are some scenarios—such as if you use persistent bids|mash|where terminating your instances is not sufficient to stop your request from being re-fulfilled. Therefore, it is a best practice to both cancel any active bids and terminate any running instances.

The following code demonstrates how to cancel your requests.

```
try {
    // Cancel requests.
    CancelSpotInstanceRequestsRequest cancelRequest =
        new CancelSpotInstanceRequestsRequest(spotInstanceRequestIds);
    ec2.cancelSpotInstanceRequests(cancelRequest);
} catch (AmazonServiceException e) {
    // Write out any exceptions that may have occurred.
    System.out.println("Error cancelling instances");
    System.out.println("Caught Exception: " + e.getMessage());
    System.out.println("Reponse Status Code: " + e.getStatusCode());
    System.out.println("Error Code: " + e.getErrorCode());
    System.out.println("Request ID: " + e.getRequestId());
}
```

To terminate any outstanding instances, you will need the instance ID associated with the request that started them. The following code example takes our original code for monitoring the instances and adds an `ArrayList` in which we store the instance ID associated with the `describeInstance` response.

```
// Create a variable that will track whether there are any requests
// still in the open state.
boolean anyOpen;
// Initialize variables.
ArrayList<String> instanceIds = new ArrayList<String>();

do {
    // Create the describeRequest with all of the request ids to
    // monitor (e.g. that we started).
    DescribeSpotInstanceRequestsRequest describeRequest = new DescribeSpotInstanceRequestsRequest();
    describeRequest.setSpotInstanceRequestIds(spotInstanceRequestIds);

    // Initialize the anyOpen variable to false, which assumes there
    // are no requests open unless we find one that is still open.
    anyOpen = false;

    try {
```

## Programming Examples

```
// Retrieve all of the requests we want to monitor.
DescribeSpotInstanceRequestsResult describeResult =
    ec2.describeSpotInstanceRequests(describeRequest);

List<SpotInstanceRequest> describeResponses =
    describeResult.getSpotInstanceRequests();

// Look through each request and determine if they are all
// in the active state.
for (SpotInstanceRequest describeResponse : describeResponses) {
    // If the state is open, it hasn't changed since we
    // attempted to request it. There is the potential for
    // it to transition almost immediately to closed or
    // cancelled so we compare against open instead of active.
    if (describeResponse.getState().equals("open")) {
        anyOpen = true; break;
    }
    // Add the instance id to the list we will
    // eventually terminate.
    instanceIds.add(describeResponse.getInstanceId());
}
} catch (AmazonServiceException e) {
    // If we have an exception, ensure we don't break out
    // of the loop. This prevents the scenario where there
    // was a blip on the wire.
    anyOpen = true;
}

try {
    // Sleep for 60 seconds.
    Thread.sleep(60*1000);
} catch (Exception e) {
    // Do nothing because it woke up early.
}
} while (anyOpen);
```

Using the instance IDs, stored in the `ArrayList`, terminate any running instances using the following code snippet.

```
try {
    // Terminate instances.
    TerminateInstancesRequest terminateRequest = new TerminateInstancesRequest(instanceIds);
    ec2.terminateInstances(terminateRequest);
} catch (AmazonServiceException e) {
    // Write out any exceptions that may have occurred.
    System.out.println("Error terminating instances");
    System.out.println("Caught Exception: " + e.getMessage());
    System.out.println("Response Status Code: " + e.getStatusCode());
    System.out.println("Error Code: " + e.getErrorCode());
```

## Programming Examples

```
        System.out.println("Request ID: " + e.getRequestId());  
    }
```

### ***Bringing It All Together***

To bring this all together, we provide a more object-oriented approach that combines the preceding steps we showed: initializing the EC2 Client, submitting the Spot Request, determining when the Spot Requests are no longer in the open state, and cleaning up any lingering Spot request and associated instances. We create a class called Requests that performs these actions.

We also create a GettingStartedApp class, which has a main method where we perform the high level function calls. Specifically, we initialize the Requests object described previously. We submit the Spot Instance request. Then we wait for the Spot request to reach the "Active" state. Finally, we clean up the requests and instances.

The complete source code for this example can be viewed or downloaded at [GitHub](#).

Congratulations! You have just completed the getting started tutorial for developing Spot Instance software with the AWS SDK for Java.

### ***Next Steps***

Proceed with [Tutorial: Advanced Amazon EC2 Spot Request Management](#).

## [Tutorial: Advanced Amazon EC2 Spot Request Management](#)

Amazon EC2 spot instances allow you to bid on unused Amazon EC2 capacity and run those instances for as long as your bid exceeds the current *spot price*. Amazon EC2 changes the spot price periodically based on supply and demand. For more information about spot instances, see [Spot Instances](#) in the *Amazon EC2 User Guide for Linux Instances*.

### ***Prerequisites***

To use this tutorial you must have the AWS SDK for Java installed, as well as having met its basic installation prerequisites. See [Set up the AWS SDK for Java](#) for more information.

### ***Setting up your credentials***

To begin using this code sample, you need to add AWS credentials to the `AwsCredentials.properties` file as follows:

1. Open the `AwsCredentials.properties` file.
2. Set your access key / secret key id combination in the `AwsCredentials.properties` file.

### **Note**

We recommend that you use the credentials of an IAM user to provide these values. For more information, see [Sign Up for AWS and Create an IAM User](#).

Now that you have configured your settings, you can get started using the code in the example.

### **Setting up a security group**

A security group acts as a firewall that controls the traffic allowed in and out of a group of instances. By default, an instance is started without any security group, which means that all incoming IP traffic, on any TCP port will be denied. So, before submitting our Spot Request, we will set up a security group that allows the necessary network traffic. For the purposes of this tutorial, we will create a new security group called "GettingStarted" that allows Secure Shell (SSH) traffic from the IP address where you are running your application from. To set up a new security group, you need to include or run the following code sample that sets up the security group programmatically.

After we create an `AmazonEC2` client object, we create a `CreateSecurityGroupRequest` object with the name, "GettingStarted" and a description for the security group. Then we call the `ec2.createSecurityGroup` API to create the group.

To enable access to the group, we create an `ipPermission` object with the IP address range set to the CIDR representation of the subnet for the local computer; the "/10" suffix on the IP address indicates the subnet for the specified IP address. We also configure the `ipPermission` object with the TCP protocol and port 22 (SSH). The final step is to call `ec2.authorizeSecurityGroupIngress` with the name of our security group and the `ipPermission` object.

(The following code is the same as what we used in the first tutorial.)

```
// Retrieves the credentials from the shared credentials file
AWSCredentialsProvider credentials = new ProfileCredentialsProvider("my-profile");

// Create the AmazonEC2Client object so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.standard()
    .withCredentials(credentials)
    .build();

// Create a new security group.
try {
    CreateSecurityGroupRequest securityGroupRequest =
        new CreateSecurityGroupRequest("GettingStartedGroup",
            "Getting Started Security Group");
    ec2.createSecurityGroup(securityGroupRequest);
} catch (AmazonServiceException ase) {
    // Likely this means that the group is already created, so ignore.
    System.out.println(ase.getMessage());
}

String ipAddr = "0.0.0.0/0";

// Get the IP of the current host, so that we can limit the Security Group
// by default to the ip range associated with your subnet.
try {
    // Get IP Address
    InetAddress addr = InetAddress.getLocalHost();
    ipAddr = addr.getHostAddress() + "/10";
}
catch (UnknownHostException e) {
    // Fail here...
}
```

## Programming Examples

```
}

// Create a range that you would like to populate.
ArrayList<String> ipRanges = new ArrayList<String>();
ipRanges.add(ipAddr);

// Open up port 22 for TCP traffic to the associated IP from
// above (e.g. ssh traffic).
ArrayList<IpPermission> ipPermissions = new ArrayList<IpPermission> ();
IpPermission ipPermission = new IpPermission();
ipPermission.setIpProtocol("tcp");
ipPermission.setFromPort(new Integer(22));
ipPermission.setToPort(new Integer(22));
ipPermission.setIpRanges(ipRanges);
ipPermissions.add(ipPermission);

try {
    // Authorize the ports to the used.
    AuthorizeSecurityGroupIngressRequest ingressRequest =
        new AuthorizeSecurityGroupIngressRequest(
            "GettingStartedGroup", ipPermissions);
    ec2.authorizeSecurityGroupIngress(ingressRequest);
}
catch (AmazonServiceException ase) {
    // Ignore because this likely means the zone has already
    // been authorized.
    System.out.println(ase.getMessage());
}
```

You can view this entire code sample in the `advanced.CreateSecurityGroupApp.java` code sample. Note you only need to run this application once to create a new security group.

### Note

You can also create the security group using the AWS Toolkit for Eclipse. See [Managing Security Groups from AWS Explorer](#) in the *AWS Toolkit for Eclipse User Guide* for more information.

### Detailed spot instance request creation options

As we explained in *Tutorial: Amazon EC2 Spot Instances*, you need to build your request with an instance type, an Amazon Machine Image (AMI), and maximum bid price.

Let's start by creating a `RequestSpotInstanceRequest` object. The request object requires the number of instances you want and the bid price. Additionally, we need to set the `LaunchSpecification` for the request, which includes the instance type, AMI ID, and security group you want to use. After the request is populated, we call the `requestSpotInstances` method on the `AmazonEC2Client` object. An example of how to request a Spot instance follows.

## Programming Examples

(The following code is the same as what we used in the first tutorial.)

```
// Create the AmazonEC2 client so we can call various APIs.  
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();  
  
// Initializes a Spot Instance Request  
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();  
  
// Request 1 x t1.micro instance with a bid price of $0.03.  
requestRequest.setSpotPrice("0.03");  
requestRequest.setInstanceCount(Integer.valueOf(1));  
  
// Set up the specifications of the launch. This includes the  
// instance type (e.g. t1.micro) and the latest Amazon Linux  
// AMI id available. Note, you should always use the latest  
// Amazon Linux AMI id or another of your choosing.  
LaunchSpecification launchSpecification = new LaunchSpecification();  
launchSpecification.setImageId("ami-8c1fece5");  
launchSpecification.setInstanceType("t1.micro");  
  
// Add the security group to the request.  
ArrayList<String> securityGroups = new ArrayList<String>();  
securityGroups.add("GettingStartedGroup");  
launchSpecification.setSecurityGroups(securityGroups);  
  
// Add the launch specification.  
requestRequest.setLaunchSpecification(launchSpecification);  
  
// Call the RequestSpotInstance API.  
RequestSpotInstancesResult requestResult =  
    ec2.requestSpotInstances(requestRequest);
```

### Persistent vs. one-time requests

When building a Spot request, you can specify several optional parameters. The first is whether your request is one-time only or persistent. By default, it is a one-time request. A one-time request can be fulfilled only once, and after the requested instances are terminated, the request will be closed. A persistent request is considered for fulfillment whenever there is no Spot Instance running for the same request. To specify the type of request, you simply need to set the Type on the Spot request. This can be done with the following code.

```
// Retrieves the credentials from an AWSCredentials.properties file.  
AWSCredentials credentials = null;  
try {  
    credentials = new PropertiesCredentials(  
        GettingStartedApp.class.getResourceAsStream("AwsCredentials.properties"));  
}  
catch (IOException e1) {  
    System.out.println(  
        "Credentials were not properly entered into AwsCredentials.properties.");
```

## Programming Examples

```
System.out.println(e1.getMessage());
System.exit(-1);
}

// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest =
    new RequestSpotInstancesRequest();

// Request 1 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(1));

// Set the type of the bid to persistent.
requestRequest.setType("persistent");

// Set up the specifications of the launch. This includes the
// instance type (e.g. t1.micro) and the latest Amazon Linux
// AMI id available. Note, you should always use the latest
// Amazon Linux AMI id or another of your choosing.
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstance API.
RequestSpotInstancesResult requestResult =
    ec2.requestSpotInstances(requestRequest);
```

### ***Limits the duration of a request***

You can also optionally specify the length of time that your request will remain valid. You can specify both a starting and ending time for this period. By default, a Spot request will be considered for fulfillment from the moment it is created until it is either fulfilled or canceled by you. However you can constrain the validity period if you need to. An example of how to specify this period is shown in the following code.

```
// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
```

## Programming Examples

```
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 1 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(1));

// Set the valid start time to be two minutes from now.
Calendar cal = Calendar.getInstance();
cal.add(Calendar.MINUTE, 2);
requestRequest.setValidFrom(cal.getTime());

// Set the valid end time to be two minutes and two hours from now.
cal.add(Calendar.HOUR, 2);
requestRequest.setValidUntil(cal.getTime());

// Set up the specifications of the launch. This includes
// the instance type (e.g. t1.micro)

// and the latest Amazon Linux AMI id available.
// Note, you should always use the latest Amazon
// Linux AMI id or another of your choosing.
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstances API.
RequestSpotInstancesResult requestResult = ec2.requestSpotInstances(requestRequest);
```

### **Grouping your Amazon EC2 spot instance requests**

You have the option of grouping your Spot instance requests in several different ways. We'll look at the benefits of using launch groups, Availability Zone groups, and placement groups.

If you want to ensure your Spot instances are all launched and terminated together, then you have the option to leverage a launch group. A launch group is a label that groups a set of bids together. All instances in a launch group are started and terminated together. Note, if instances in a launch group have already been fulfilled, there is no guarantee that new instances launched with the same launch group will also be fulfilled. An example of how to set a Launch Group is shown in the following code example.

```
// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();
```

## Programming Examples

```
// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 5 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(5));

// Set the launch group.
requestRequest.setLaunchGroup("ADVANCED-DEMO-LAUNCH-GROUP");

// Set up the specifications of the launch. This includes
// the instance type (e.g. t1.micro) and the latest Amazon Linux
// AMI id available. Note, you should always use the latest
// Amazon Linux AMI id or another of your choosing.
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstance API.
RequestSpotInstancesResult requestResult =
    ec2.requestSpotInstances(requestRequest);
```

If you want to ensure that all instances within a request are launched in the same Availability Zone, and you don't care which one, you can leverage Availability Zone groups. An Availability Zone group is a label that groups a set of instances together in the same Availability Zone. All instances that share an Availability Zone group and are fulfilled at the same time will start in the same Availability Zone. An example of how to set an Availability Zone group follows.

```
// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 5 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(5));

// Set the availability zone group.
requestRequest.setAvailabilityZoneGroup("ADVANCED-DEMO-AZ-GROUP");
```

## Programming Examples

```
// Set up the specifications of the launch. This includes the instance
// type (e.g. t1.micro) and the latest Amazon Linux AMI id available.
// Note, you should always use the latest Amazon Linux AMI id or another
// of your choosing.
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstances API.
RequestSpotInstancesResult requestResult =
    ec2.requestSpotInstances(requestRequest);
```

You can specify an Availability Zone that you want for your Spot Instances. The following code example shows you how to set an Availability Zone.

```
// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 1 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(1));

// Set up the specifications of the launch. This includes the instance
// type (e.g. t1.micro) and the latest Amazon Linux AMI id available.
// Note, you should always use the latest Amazon Linux AMI id or another
// of your choosing.
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Set up the availability zone to use. Note we could retrieve the
// availability zones using the ec2.describeAvailabilityZones() API. For
// this demo we will just use us-east-1a.
```

## Programming Examples

```
SpotPlacement placement = new SpotPlacement("us-east-1b");
launchSpecification.setPlacement(placement);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstance API.
RequestSpotInstancesResult requestResult =
    ec2.requestSpotInstances(requestRequest);
```

Lastly, you can specify a *placement group* if you are using High Performance Computing (HPC) Spot instances, such as cluster compute instances or cluster GPU instances. Placement groups provide you with lower latency and high-bandwidth connectivity between the instances. An example of how to set a placement group follows.

```
// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 1 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(1));

// Set up the specifications of the launch. This includes the instance
// type (e.g. t1.micro) and the latest Amazon Linux AMI id available.
// Note, you should always use the latest Amazon Linux AMI id or another
// of your choosing.

LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Set up the placement group to use with whatever name you desire.
// For this demo we will just use "ADVANCED-DEMO-PLACEMENT-GROUP".
SpotPlacement placement = new SpotPlacement();
placement.setGroupName("ADVANCED-DEMO-PLACEMENT-GROUP");
launchSpecification.setPlacement(placement);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstance API.
```

## Programming Examples

```
RequestSpotInstancesResult requestResult =
    ec2.requestSpotInstances(requestRequest);
```

All of the parameters shown in this section are optional. It is also important to realize that most of these parameters—with the exception of whether your bid is one-time or persistent—can reduce the likelihood of bid fulfillment. So, it is important to leverage these options only if you need them. All of the preceding code examples are combined into one long code sample, which can be found in the `com.amazonaws.codesamples.advanced.InlineGettingStartedCodeSampleApp.java` class.

### **How to persist a root partition after interruption or termination**

One of the easiest ways to manage interruption of your Spot instances is to ensure that your data is checkpointed to an Amazon Elastic Block Store (Amazon EBS) volume on a regular cadence. By checkpointing periodically, if there is an interruption you will lose only the data created since the last checkpoint (assuming no other non-idempotent actions are performed in between). To make this process easier, you can configure your Spot Request to ensure that your root partition will not be deleted on interruption or termination. We've inserted new code in the following example that shows how to enable this scenario.

In the added code, we create a `BlockDeviceMapping` object and set its associated Elastic Block Storage (EBS) to an `EBS` object that we've configured to not be deleted if the Spot Instance is terminated. We then add this `BlockDeviceMapping` to the `ArrayList` of mappings that we include in the launch specification.

```
// Retrieves the credentials from an AWS Credentials.properties file.
AWSCredentials credentials = null;
try {
    credentials = new PropertiesCredentials(
        GettingStartedApp.class.getResourceAsStream("AwsCredentials.properties"));
}
catch (IOException e1) {
    System.out.println(
        "Credentials were not properly entered into AwsCredentials.properties.");
    System.out.println(e1.getMessage());
    System.exit(-1);
}

// Create the AmazonEC2 client so we can call various APIs.
AmazonEC2 ec2 = AmazonEC2ClientBuilder.defaultClient();

// Initializes a Spot Instance Request
RequestSpotInstancesRequest requestRequest = new RequestSpotInstancesRequest();

// Request 1 x t1.micro instance with a bid price of $0.03.
requestRequest.setSpotPrice("0.03");
requestRequest.setInstanceCount(Integer.valueOf(1));

// Set up the specifications of the launch. This includes the instance
// type (e.g. t1.micro) and the latest Amazon Linux AMI id available.
// Note, you should always use the latest Amazon Linux AMI id or another
// of your choosing.
```

## Programming Examples

```
LaunchSpecification launchSpecification = new LaunchSpecification();
launchSpecification.setImageId("ami-8c1fece5");
launchSpecification.setInstanceType("t1.micro");

// Add the security group to the request.
ArrayList<String> securityGroups = new ArrayList<String>();
securityGroups.add("GettingStartedGroup");
launchSpecification.setSecurityGroups(securityGroups);

// Create the block device mapping to describe the root partition.
BlockDeviceMapping blockDeviceMapping = new BlockDeviceMapping();
blockDeviceMapping.setDeviceName("/dev/sda1");

// Set the delete on termination flag to false.
EbsBlockDevice ebs = new EbsBlockDevice();
ebs.setDeleteOnTermination(Boolean.FALSE);
blockDeviceMapping.setEbs(ebs);

// Add the block device mapping to the block list.
ArrayList<BlockDeviceMapping> blockList = new ArrayList<BlockDeviceMapping>();
blockList.add(blockDeviceMapping);

// Set the block device mapping configuration in the launch specifications.
launchSpecification.setBlockDeviceMappings(blockList);

// Add the launch specification.
requestRequest.setLaunchSpecification(launchSpecification);

// Call the RequestSpotInstances API.
RequestSpotInstancesResult requestResult =
    ec2.requestSpotInstances(requestRequest);
```

Assuming you wanted to re-attach this volume to your instance on startup, you can also use the block device mapping settings. Alternatively, if you attached a non-root partition, you can specify the Amazon EBS volumes you want to attach to your Spot instance after it resumes. You do this simply by specifying a snapshot ID in your `EbsBlockDevice` and alternative device name in your `BlockDeviceMapping` objects. By leveraging block device mappings, it can be easier to bootstrap your instance.

Using the root partition to checkpoint your critical data is a great way to manage the potential for interruption of your instances. For more methods on managing the potential of interruption, please visit the [Managing Interruption](#) video.

### ***How to tag your spot requests and instances***

Adding tags to EC2 resources can simplify the administration of your cloud infrastructure. A form of metadata, tags can be used to create user-friendly names, enhance searchability, and improve coordination between multiple users. You can also use tags to automate scripts and portions of your processes. To read more about tagging Amazon EC2 resources, go to [Using Tags](#) in the *Amazon EC2 User Guide for Linux Instances*.

#### *Tagging requests*

## Programming Examples

To add tags to your spot requests, you need to tag them *after* they have been requested. The return value from `requestSpotInstances()` provides you with a `RequestSpotInstancesResult` object that you can use to get the spot request IDs for tagging:

```
// Call the RequestSpotInstance API.  
RequestSpotInstancesResult requestResult = ec2.requestSpotInstances(requestRequest);  
List<SpotInstanceRequest> requestResponses = requestResult.getSpotInstanceRequests();  
  
// A list of request IDs to tag  
ArrayList<String> spotInstanceRequestIds = new ArrayList<String>();  
  
// Add the request ids to the hashset, so we can determine when they hit the  
// active state.  
for (SpotInstanceRequest requestResponse : requestResponses) {  
    System.out.println("Created Spot Request: " + requestResponse.getSpotInstanceRequestId());  
    spotInstanceRequestIds.add(requestResponse.getSpotInstanceRequestId());  
}
```

Once you have the IDs, you can tag the requests by adding their IDs to a `CreateTagsRequest` and calling the EC2 client's `createTags()` method:

```
// The list of tags to create  
ArrayList<Tag> requestTags = new ArrayList<Tag>();  
requestTags.add(new Tag("keyname1", "value1"));  
  
// Create the tag request  
CreateTagsRequest createTagsRequest_requests = new CreateTagsRequest();  
createTagsRequest_requests.setResources(spotInstanceRequestIds);  
createTagsRequest_requests.setTags(requestTags);  
  
// Tag the spot request  
try {  
    ec2.createTags(createTagsRequest_requests);  
}  
catch (AmazonServiceException e) {  
    System.out.println("Error terminating instances");  
    System.out.println("Caught Exception: " + e.getMessage());  
    System.out.println("Reponse Status Code: " + e.getStatusCode());  
    System.out.println("Error Code: " + e.getErrorCode());  
    System.out.println("Request ID: " + e.getRequestId());  
}
```

### Tagging instances

Similarly to spot requests themselves, you can only tag an instance once it has been created, which will happen once the spot request has been met (it is no longer in the *open* state).

You can check the status of your requests by calling the EC2 client's `describeSpotInstanceRequests()` method with a `DescribeSpotInstanceRequestsRequest` object. The returned `DescribeSpotInstanceRequestsResult` object contains a list of `SpotInstanceRequest` objects that you can use

## Programming Examples

to query the status of your spot requests and obtain their instance IDs once they are no longer in the *open* state.

Once the spot request is no longer open, you can retrieve its instance ID from the `SpotInstanceRequest` object by calling its `getInstanceId()` method.

```
boolean anyOpen; // tracks whether any requests are still open

// a list of instances to tag.
ArrayList<String> instanceIds = new ArrayList<String>();

do {
    DescribeSpotInstanceRequestsRequest describeRequest =
        new DescribeSpotInstanceRequestsRequest();
    describeRequest.setSpotInstanceRequestIds(spotInstanceRequestIds);

    anyOpen=false; // assume no requests are still open

    try {
        // Get the requests to monitor
        DescribeSpotInstanceRequestsResult describeResult =
            ec2.describeSpotInstanceRequests(describeRequest);

        List<SpotInstanceRequest> describeResponses =
            describeResult.getSpotInstanceRequests();

        // are any requests open?
        for (SpotInstanceRequest describeResponse : describeResponses) {
            if (describeResponse.getState().equals("open")) {
                anyOpen = true;
                break;
            }
            // get the corresponding instance ID of the spot request
            instanceIds.add(describeResponse.getInstanceId());
        }
    }
    catch (AmazonServiceException e) {
        // Don't break the loop due to an exception (it may be a temporary issue)
        anyOpen = true;
    }

    try {
        Thread.sleep(60*1000); // sleep 60s.
    }
    catch (Exception e) {
        // Do nothing if the thread woke up early.
    }
} while (anyOpen);
```

Now you can tag the instances that are returned:

## Programming Examples

```
// Create a list of tags to create
ArrayList<Tag> instanceTags = new ArrayList<Tag>();
instanceTags.add(new Tag("keyname1", "value1"));

// Create the tag request
CreateTagsRequest createTagsRequest_instances = new CreateTagsRequest();
createTagsRequest_instances.setResources(instanceIds);
createTagsRequest_instances.setTags(instanceTags);

// Tag the instance
try {
    ec2.createTags(createTagsRequest_instances);
}
catch (AmazonServiceException e) {
    // Write out any exceptions that may have occurred.
    System.out.println("Error terminating instances");
    System.out.println("Caught Exception: " + e.getMessage());
    System.out.println("Reponse Status Code: " + e.getStatusCode());
    System.out.println("Error Code: " + e.getErrorCode());
    System.out.println("Request ID: " + e.getRequestId());
}
```

### **Cancelling spot requests and terminating instances**

#### *Cancelling a spot request*

To cancel a spot instance request, call `cancelSpotInstanceRequests` on the EC2 client with a `CancelSpotInstanceRequestsRequest` object.

```
try {
    CancelSpotInstanceRequestsRequest cancelRequest = new CancelSpotInstanceRequestsRequest();
    ec2.cancelSpotInstanceRequests(cancelRequest);
} catch (AmazonServiceException e) {
    System.out.println("Error cancelling instances");
    System.out.println("Caught Exception: " + e.getMessage());
    System.out.println("Reponse Status Code: " + e.getStatusCode());
    System.out.println("Error Code: " + e.getErrorCode());
    System.out.println("Request ID: " + e.getRequestId());
}
```

#### *Terminating spot instances*

You can terminate any spot instances that are running by passing their IDs to the EC2 client's `terminateInstances()` method.

```
try {
    TerminateInstancesRequest terminateRequest = new TerminateInstancesRequest(instanceIds);
    ec2.terminateInstances(terminateRequest);
} catch (AmazonServiceException e) {
```

## Programming Examples

```
System.out.println("Error terminating instances");
System.out.println("Caught Exception: " + e.getMessage());
System.out.println("Reponse Status Code: " + e.getStatusCode());
System.out.println("Error Code: " + e.getErrorCode());
System.out.println("Request ID: " + e.getRequestId());
}
```

### ***Bringing it all together***

To bring this all together, we provide a more object-oriented approach that combines the steps we showed in this tutorial into one easy to use class. We instantiate a class called `Requests` that performs these actions. We also create a `GettingStartedApp` class, which has a main method where we perform the high level function calls.

The complete source code for this example can be viewed or downloaded at [GitHub](#).

Congratulations! You've completed the Advanced Request Features tutorial for developing Spot Instance software with the AWS SDK for Java.

## Amazon S3 Examples

This section provides examples of programming [Amazon S3](#) using the [AWS SDK for Java](#).

### Note

The examples include only the code needed to demonstrate each technique. The [complete example code is available on GitHub](#). From there, you can download a single source file or clone the repository locally to get all the examples to build and run.

## Creating, Listing, and Deleting Amazon S3 Buckets

Every object (file) in Amazon S3 must reside within a *bucket*, which represents a collection (container) of objects. Each bucket is known by a *key* (name), which must be unique. For detailed information about buckets and their configuration, see [Working with Amazon S3 Buckets](#) in the *Amazon S3 Developer Guide*.

### Best Practice

We recommend that you enable the `AbortIncompleteMultipartUpload` lifecycle rule on your Amazon S3 buckets.

This rule directs Amazon S3 to abort multipart uploads that don't complete within a specified number of days after being initiated. When the set time limit is exceeded, Amazon S3 aborts the upload and then deletes the incomplete upload data.

For more information, see [Lifecycle Configuration for a Bucket with Versioning](#) in the *Amazon S3 User Guide*.

## Note

These code snippets assume that you understand the material in *Using the AWS SDK for Java* and have configured default AWS credentials using the information in *Set up AWS Credentials and Region for Development*.

### Create a Bucket

Use the [AmazonS3](#) client's `createBucket` method. The new `Bucket` is returned.

#### Imports

```
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.model.Bucket;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.AmazonServiceException;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    Bucket b = s3.createBucket(bucket_name);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
System.out.println("Done!");
```

See the [complete example](#).

### List Buckets

Use the [AmazonS3](#) client's `listBucket` method. If successful, a list of `Bucket` is returned.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.services.s3.model.Bucket;
import java.util.List;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
List<Bucket> buckets = s3.listBuckets();
System.out.println("Your Amazon S3 buckets:");
```

## Programming Examples

```
for (Bucket b : buckets) {  
    System.out.println("★ " + b.getName());  
}
```

See the [complete example](#).

### Delete a Bucket

Before you can delete an Amazon S3 bucket, you must ensure that the bucket is empty or an error will result. If you have a [versioned bucket](#), you must also delete any versioned objects associated with the bucket.

#### Note

The [complete example](#) includes each of these steps in order, providing a complete solution for deleting an Amazon S3 bucket and its contents.

#### *Remove Objects from an Unversioned Bucket Before Deleting It*

Use the [AmazonS3](#) client's `listObjects` method to retrieve the list of objects and `deleteObject` to delete each one.

#### Imports

```
import com.amazonaws.AmazonServiceException;  
import com.amazonaws.services.s3.AmazonS3;  
import com.amazonaws.services.s3.AmazonS3Client;  
import com.amazonaws.services.s3.model.ObjectListing;  
import com.amazonaws.services.s3.model.S3ObjectSummary;  
import java.util.Iterator;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();  
try {  
    System.out.println(" - removing objects from bucket");  
    ObjectListing object_listing = s3.listObjects(bucket_name);  
    while (true) {  
        for (Iterator<?> iterator =  
            object_listing.getObjectSummaries().iterator();  
            iterator.hasNext();) {  
            S3ObjectSummary summary = (S3ObjectSummary) iterator.next();  
            s3.deleteObject(bucket_name, summary.getKey());  
        }  
  
        // more object_listing to retrieve?  
        if (object_listing.isTruncated()) {  
            object_listing = s3.listNextBatchOfObjects(object_listing);  
        }  
    }  
}
```

## Programming Examples

```
        } else {
            break;
        }
    };
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete example](#).

### *Remove Objects from a Versioned Bucket Before Deleting It*

If you're using a [versioned bucket](#), you also need to remove any stored versions of the objects in the bucket before the bucket can be deleted.

Using a pattern similar to the one used when removing objects within a bucket, remove versioned objects by using the [AmazonS3](#) client's [listVersions](#) method to list any versioned objects, and then [deleteVersion](#) to delete each one.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.services.s3.model.ListVersionsRequest;
import com.amazonaws.services.s3.model.ObjectListing;
import com.amazonaws.services.s3.model.S3ObjectSummary;
import com.amazonaws.services.s3.model.S3VersionSummary;
import com.amazonaws.services.s3.model.VersionListing;
import java.util.Iterator;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    System.out.println(" - removing objects from bucket");
    ObjectListing object_listing = s3.listObjects(bucket_name);
    while (true) {
        for (Iterator<?> iterator =
                object_listing.getObjectSummaries().iterator();
                iterator.hasNext();) {
            S3ObjectSummary summary = (S3ObjectSummary) iterator.next();
            s3.deleteObject(bucket_name, summary.getKey());
        }

        // more object_listing to retrieve?
        if (object_listing.isTruncated()) {
            object_listing = s3.listNextBatchOfObjects(object_listing);
        } else {
    }
```

## Programming Examples

```
        break;
    }
};

System.out.println(" - removing versions from bucket");
VersionListing version_listing = s3.listVersions(
    new ListVersionsRequest().withBucketName(bucket_name));
while (true) {
    for (Iterator<?> iterator =
        version_listing.getVersionSummaries().iterator();
        iterator.hasNext();) {
        S3VersionSummary vs = (S3VersionSummary) iterator.next();
        s3.deleteVersion(
            bucket_name, vs.getKey(), vs.getVersionId());
    }

    if (version_listing.isTruncated()) {
        version_listing = s3.listNextBatchOfVersions(
            version_listing);
    } else {
        break;
    }
}
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete example](#).

### Delete an Empty Bucket

Once you remove the objects from a bucket (including any versioned objects), you can delete the bucket itself by using the [AmazonS3](#) client's `deleteBucket` method.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    s3.deleteBucket(bucket_name);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

## Programming Examples

See the [complete example](#).

### Performing Operations on Amazon S3 Objects

An Amazon S3 object represents a *file* or collection of data. Every object must reside within a *bucket*.

#### Note

These code snippets assume that you understand the material in *Using the AWS SDK for Java* and have configured default AWS credentials using the information in *Set up AWS Credentials and Region for Development*.

<b>Upload an Object</b>	<b>85</b>
<b>List Objects</b>	<b>85</b>
<b>Download an Object</b>	<b>86</b>
<b>Copy, Move, or Rename Objects</b>	<b>87</b>
<b>Delete an Object</b>	<b>88</b>
<b>Delete Multiple Objects at Once</b>	<b>88</b>

#### ***Upload an Object***

Use the [AmazonS3](#) client's `putObject` method, supplying a bucket name, key name, and file to upload. *The bucket must exist, or an error will result.*

#### Imports

```
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.AmazonServiceException;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    s3.putObject(bucket_name, key_name, file_path);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete example](#).

#### ***List Objects***

To get a list of objects within a bucket, use the [AmazonS3](#) client's `listObjects` method, supplying the name of a bucket.

## Programming Examples

The `listObjects` method returns an `ObjectListing` object that provides information about the objects in the bucket. To list the object names (keys), use the `getObjectSummaries` method to get a `List` of `S3ObjectSummary` objects, each of which represents a single object in the bucket. Then call its `getKey` method to retrieve the object's name.

### Imports

```
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.services.s3.model.ObjectListing;
import com.amazonaws.services.s3.model.S3ObjectSummary;
import java.util.List;
```

### Code

```
final AmazonS3 s3 = new AmazonS3Client();
ObjectListing ol = s3.listObjects(bucket_name);
List<S3ObjectSummary> objects = ol.getObjectSummaries();
for (S3ObjectSummary os: objects) {
    System.out.println("* " + os.getKey());
}
```

See the [complete example](#).

## Download an Object

Use the `AmazonS3` client's `getObject` method, passing it the name of a bucket and object to download. If successful, the method returns an `S3Object`. *The specified bucket and object key must exist, or an error will result.*

You can get the object's contents by calling `getObjectContent` on the `S3Object`. This returns an `S3ObjectInputStream` that behaves as a standard Java `InputStream` object.

The following example downloads an object from S3 and saves its contents to a file (using the same name as the object's key).

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.services.s3.model.S3Object;
import com.amazonaws.services.s3.model.S3ObjectInputStream;
import java.io.File;
import java.io.FileNotFoundException;
import java.io.FileOutputStream;
import java.io.IOException;
```

### Code

## Programming Examples

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    S3Object o = s3.getObject(bucket_name, key_name);
    S3ObjectInputStream s3is = o.getObjectContent();
    FileOutputStream fos = new FileOutputStream(new File(key_name));
    byte[] read_buf = new byte[1024];
    int read_len = 0;
    while ((read_len = s3is.read(read_buf)) > 0) {
        fos.write(read_buf, 0, read_len);
    }
    s3is.close();
    fos.close();
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
} catch (FileNotFoundException e) {
    System.err.println(e.getMessage());
    System.exit(1);
} catch (IOException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}
```

See the [complete example](#).

### ***Copy, Move, or Rename Objects***

You can copy an object from one bucket to another by using the `AmazonS3` client's `copyObject` method. It takes the name of the bucket to copy from, the object to copy, and the destination bucket and name.

#### **Imports**

```
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.AmazonServiceException;
```

#### **Code**

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    s3.copyObject(from_bucket, object_key, to_bucket, object_key);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete example](#).

## Note

You can use `copyObject` with `deleteObject` to **move** or **rename** an object, by first copying the object to a new name (you can use the same bucket as both the source and destination) and then deleting the object from its old location.

## Delete an Object

Use the [AmazonS3](#) client's `deleteObject` method, passing it the name of a bucket and object to delete. *The specified bucket and object key must exist, or an error will result.*

### Imports

```
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.AmazonServiceException;
```

### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    s3.deleteObject(bucket_name, object_key);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the [complete example](#).

## Delete Multiple Objects at Once

Using the [AmazonS3](#) client's `deleteObjects` method, you can delete multiple objects from the same bucket by passing their names to the [DeleteObjectRequest withKeys](#) method.

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.services.s3.model.DeleteObjectsRequest;
```

### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    DeleteObjectsRequest dor = new DeleteObjectsRequest(bucket_name)
        .withKeys(object_keys);
```

## Programming Examples

```
s3.deleteObjects(dor);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

See the complete example.

## Managing Access to Amazon S3 Buckets Using Bucket Policies

You can set, get, or delete a *bucket policy* to manage access to your Amazon S3 buckets.

### **Set a Bucket Policy**

You can set the bucket policy for a particular S3 bucket by:

- Calling the [AmazonS3](#) client's `setBucketPolicy` and providing it with a `SetBucketPolicyRequest`
- Setting the policy directly by using the `setBucketPolicy` overload that takes a bucket name and policy text (in JSON format)

#### Imports

```
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.AmazonServiceException;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    s3.setBucketPolicy(bucket_name, policy_text);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

#### Use the Policy Class to Generate or Validate a Policy

When providing a bucket policy to `setBucketPolicy`, you can do the following:

- Specify the policy directly as a string of JSON-formatted text
- Build the policy using the [Policy](#) class

By using the [Policy](#) class, you don't have to be concerned about correctly formatting your text string, as shown in the following example.

#### Imports

```
import com.amazonaws.auth.policy.Action;
import com.amazonaws.auth.policy.Policy;
import com.amazonaws.auth.policy.Principal;
```

## Programming Examples

```
import com.amazonaws.auth.policy.Resource;
import com.amazonaws.auth.policy.Statement;
import com.amazonaws.auth.policy.actions.S3Actions;
```

### Code

```
Policy bucket_policy = new Policy().withStatements(
    new Statement(Statement.Effect.Allow)
        .withPrincipals(Principal.AllUsers)
        .withActions(S3Actions.GetObject)
        .withResources(new Resource(
            "arn:aws:s3:::" + bucket_name + "/*")));
```

To get the JSON policy text from the Policy class, use its `toJson` method.

```
return bucket_policy.toJson();
```

The `Policy` class also provides a `fromJson` method that can attempt to build a policy using a passed-in JSON string. The method validates it to ensure that the text can be transformed into a valid policy structure, and will fail with an `IllegalArgumentException` if the policy text is invalid.

```
try {
    bucket_policy = Policy.fromJson(file_text.toString());
} catch (IllegalArgumentException e) {
    System.out.format("Invalid policy text in file: \"%s\"", policy_file);
    System.out.println(e.getMessage());
}
```

You can use this technique to prevalidate a policy that you read in from a file or other means.

See the [complete example](#).

## **Get a Bucket Policy**

To retrieve the policy for an Amazon S3 bucket, call the `AmazonS3` client's `getBucketPolicy` method, passing it the name of the bucket to get the policy from.

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
import com.amazonaws.services.s3.model.BucketPolicy;
```

### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    BucketPolicy bucket_policy = s3.getBucketPolicy(bucket_name);
```

## Programming Examples

```
    policy_text = bucket_policy.getPolicyText();
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

If the named bucket doesn't exist, if you don't have access to it, or if it has no bucket policy, an `AmazonServiceException` is thrown.

See the [complete example](#).

### **Delete a Bucket Policy**

To delete a bucket policy, call the `AmazonS3` client's `deleteBucketPolicy`, providing it with the bucket name.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.AmazonS3;
import com.amazonaws.services.s3.AmazonS3Client;
```

#### Code

```
final AmazonS3 s3 = new AmazonS3Client();
try {
    s3.deleteBucketPolicy(bucket_name);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
```

This method succeeds even if the bucket doesn't already have a policy. If you specify a bucket name that doesn't exist or if you don't have access to the bucket, an `AmazonServiceException` is thrown.

See the [complete example](#).

#### **More Info**

- Access Policy Language Overview in the *Amazon S3 Developer Guide*
- Bucket Policy Examples in the *Amazon S3 Developer Guide*

## Using TransferManager for Amazon S3 Operations

You can use the AWS SDK for Java `TransferManager` class to reliably transfer files from the local environment to Amazon S3 and to copy objects from one S3 location to another. `TransferManager` can get the progress of a transfer and pause or resume uploads and downloads.

### Best Practice

We recommend that you enable the [AbortIncompleteMultipartUpload](#) lifecycle rule on your Amazon S3 buckets.

This rule directs Amazon S3 to abort multipart uploads that don't complete within a specified number of days after being initiated. When the set time limit is exceeded, Amazon S3 aborts the upload and then deletes the incomplete upload data.

For more information, see [Lifecycle Configuration for a Bucket with Versioning](#) in the *Amazon S3 User Guide*.

### Note

These code snippets assume that you understand the material in *Using the AWS SDK for Java* and have configured default AWS credentials using the information in *Set up AWS Credentials and Region for Development*.

## Upload Files and Directories

[TransferManager](#) can upload files, file lists, and directories to any Amazon S3 buckets that you've previously created.

[Upload a Single File](#) 92

[Upload a List of Files](#) 93

[Upload a Directory](#) 94

### Upload a Single File

Call the [TransferManager](#) `upload` method, providing an Amazon S3 bucket name, a key (object) name, and a standard Java `File` object that represents the file to upload.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.transfer.TransferManager;
import com.amazonaws.services.s3.transfer.Upload;
import java.io.File;
```

#### Code

```
File f = new File(file_path);
TransferManager xfer_mgr = new TransferManager();
try {
    Upload xfer = xfer_mgr.upload(bucket_name, key_name, f);
```

## Programming Examples

```
// loop with Transfer.isDone()
// or block with Transfer.waitForCompletion()
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

The `upload` method returns *immediately*, providing an `Upload` object to use to check the transfer state or to wait for it to complete.

See Wait for a Transfer to Complete for information about using `waitForCompletion` to successfully complete a transfer before calling `TransferManager`'s `shutdownNow` method. While waiting for the transfer to complete, you can poll or listen for updates about its status and progress. See Get Transfer Status and Progress for more information.

See the [complete example](#).

### Upload a List of Files

To upload multiple files in one operation, call the `TransferManager` `uploadFileList` method, providing the following:

- An Amazon S3 bucket name
- A *key prefix* to prepend to the names of the created objects (the path within the bucket in which to place the objects)
- A `File` object that represents the relative directory from which to create file paths
- A `List` object containing a set of `File` objects to upload

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.transfer.TransferManager;
import com.amazonaws.services.s3.transfer.MultipleFileUpload;
import java.io.File;
import java.util.ArrayList;
```

### Code

```
ArrayList<File> files = new ArrayList<File>();
for (String path : file_paths) {
    files.add(new File(path));
}

TransferManager xfer_mgr = new TransferManager();
try {
    MultipleFileUpload xfer = xfer_mgr.uploadFileList(bucket_name,
        key_prefix, new File("."), files);
    // loop with Transfer.isDone()
    // or block with Transfer.waitForCompletion()
```

## Programming Examples

```
} catch (AmazonServiceException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

See [Wait for a Transfer to Complete](#) for information about using `waitForCompletion` to successfully complete a transfer before calling `TransferManager`'s `shutdownNow` method. While waiting for the transfer to complete, you can poll or listen for updates about its status and progress. See [Get Transfer Status and Progress](#) for more information.

The `MultipleFileUpload` object returned by `uploadFileList` can be used to query the transfer state or progress. See [Poll the Current Progress of a Transfer](#) and [Get Transfer Progress with a ProgressListener](#) for more information.

You can also use `MultipleFileUpload`'s `getSubTransfers` method to get the individual `Upload` objects for each file being transferred. For more information, see [Get the Progress of Subtransfers](#).

See the [complete example](#).

### *Upload a Directory*

You can use `TransferManager`'s `uploadDirectory` method to upload an entire directory of files, with the option to copy files in subdirectories recursively. You provide an Amazon S3 bucket name, an S3 key prefix, a `File` object representing the local directory to copy, and a boolean value indicating whether you want to copy subdirectories recursively (`true` or `false`).

### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.transfer.TransferManager;
import com.amazonaws.services.s3.transfer.Upload;
import com.amazonaws.services.s3.transfer.MultipleFileUpload;
import java.io.File;
```

### Code

```
TransferManager xfer_mgr = new TransferManager();
try {
    MultipleFileUpload xfer = xfer_mgr.uploadDirectory(bucket_name,
        key_prefix, new File(dir_path), recursive);
    // loop with Transfer.isDone()
    // or block with Transfer.waitForCompletion()
} catch (AmazonServiceException e) {
    System.err.println(e.getMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

See [Wait for a Transfer to Complete](#) for information about using `waitForCompletion` to successfully complete a transfer before calling `TransferManager`'s `shutdownNow` method. While waiting for the transfer to

## Programming Examples

complete, you can poll or listen for updates about its status and progress. See Get Transfer Status and Progress for more information.

The [MultipleFileUpload](#) object returned by [uploadFileList](#) can be used to query the transfer state or progress. See Poll the Current Progress of a Transfer and Get Transfer Progress with a ProgressListener for more information.

You can also use [MultipleFileUpload](#)'s [getSubTransfers](#) method to get the individual [Upload](#) objects for each file being transferred. For more information, see Get the Progress of Subtransfers.

See the [complete example](#).

### Download Files or Directories

Use the [TransferManager](#) class to download either a single file (Amazon S3 object) or a directory (an Amazon S3 bucket name followed by an object prefix) from Amazon S3.

[Download a Single File](#)

95

[Download a Directory](#)

96

#### Download a Single File

Use the [TransferManager](#)'s [download](#) method, providing the Amazon S3 bucket name containing the object you want to download, the key (object) name, and a [File](#) object that represents the file to create on your local system.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.transfer.TransferManager;
import com.amazonaws.services.s3.transfer.Download;
import java.io.File;
```

#### Code

```
File f = new File(file_path);
TransferManager xfer_mgr = new TransferManager();
try {
    Download xfer = xfer_mgr.download(bucket_name, key_name, f);
    // loop with Transfer.isDone()
    // or block with Transfer.waitForCompletion()
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

See Wait for a Transfer to Complete for information about using [waitForCompletion](#) to successfully complete a transfer before calling [TransferManager](#)'s [shutdownNow](#) method. While waiting for the transfer to complete, you can poll or listen for updates about its status and progress. See Get Transfer Status and Progress for more information.

See the [complete example](#).

## Programming Examples

### *Download a Directory*

To download a set of files that share a common key prefix (analogous to a directory on a file system) from Amazon S3, use the [TransferManager](#) `downloadDirectory` method. The method takes the Amazon S3 bucket name containing the objects you want to download, the object prefix shared by all of the objects, and a [File](#) object that represents the directory to download the files into on your local system. If the named directory doesn't exist yet, it will be created.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.transfer.TransferManager;
import com.amazonaws.services.s3.transfer.MultipleFileDownload;
import java.io.File;
```

#### Code

```
TransferManager xfer_mgr = new TransferManager();
try {
    MultipleFileDownload xfer = xfer_mgr.downloadDirectory(
        bucket_name, key_prefix, new File(dir_path));
    // loop with Transfer.isDone()
    // or block with Transfer.waitForCompletion()
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

See [Wait for a Transfer to Complete](#) for information about using `waitForCompletion` to successfully complete a transfer before calling [TransferManager](#)'s `shutdownNow` method. While waiting for the transfer to complete, you can poll or listen for updates about its status and progress. See [Get Transfer Status and Progress](#) for more information.

See the [complete example](#).

### ***Copy Objects***

To copy an object from one S3 bucket to another, use the [TransferManager](#) `copy` method.

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.services.s3.transfer.Copy;
import com.amazonaws.services.s3.transfer.TransferManager;
```

#### Code

```
TransferManager xfer_mgr = new TransferManager();
try {
    Copy xfer = xfer_mgr.copy(from_bucket, from_key, to_bucket, to_key);
```

## Programming Examples

```
// loop with Transfer.isDone()
// or block with Transfer.waitForCompletion()
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

See the [complete example](#).

### **Wait for a Transfer to Complete**

If your application (or thread) can block until the transfer completes, you can use the `Transfer` interface's `waitForCompletion` method to block until the transfer is complete or an exception occurs.

```
try {
    xfer.waitForCompletion();
} catch (AmazonServiceException e) {
    System.err.println("Amazon service error: " + e.getMessage());
    System.exit(1);
} catch (AmazonClientException e) {
    System.err.println("Amazon client error: " + e.getMessage());
    System.exit(1);
} catch (InterruptedException e) {
    System.err.println("Transfer interrupted: " + e.getMessage());
    System.exit(1);
}
```

You get progress of transfers if you poll for events *before* calling `waitForCompletion`, implement a polling mechanism on a separate thread, or receive progress updates asynchronously using a `ProgressListener`.

See the [complete example](#).

### **Get Transfer Status and Progress**

Each of the classes returned by the `TransferManager` `upload*`, `download*`, and `copy` methods returns an instance of one of the following classes, depending on whether it's a single-file or multiple-file operation.

Class	Returned by
Copy	copy
Download	download
MultipleFileDownload	downloadDirectory
Upload	upload
MultipleFileUpload	uploadFileList, uploadDirectory

All of these classes implement the `Transfer` interface. `Transfer` provides useful methods to get the progress of a transfer, pause or resume the transfer, and get the transfer's current or final status.

[Poll the Current Progress of a Transfer](#)

98

## Programming Examples

<b>Get Transfer Progress with a ProgressListener</b>	<b>98</b>
<b>Get the Progress of Subtransfers</b>	<b>99</b>

### *Poll the Current Progress of a Transfer*

This loop prints the progress of a transfer, examines its current progress while running and, when complete, prints its final state.

#### Imports

```
import com.amazonaws.services.s3.transfer.TransferProgress;
```

#### Code

```
do {
    try {
        Thread.sleep(100);
    } catch (InterruptedException e) {
        return;
    }
    TransferProgress progress = xfer.getProgress();
    long so_far = progress.getBytesTransferred();
    long total = progress.getTotalBytesToTransfer();
    double pct = progress.getPercentTransferred();
} while (xfer.isDone() == false);
```

See the [complete example](#).

### *Get Transfer Progress with a ProgressListener*

You can attach a [ProgressListener](#) to any transfer by using the [Transfer](#) interface's [addProgressListener](#) method.

A [ProgressListener](#) requires only one method, [progressChanged](#), which takes a [ProgressEvent](#) object. You can use the object to get the total bytes of the operation by calling its [getBytes](#) method, and the number of bytes transferred so far by calling [getBytesTransferred](#).

#### Imports

```
import com.amazonaws.AmazonServiceException;
import com.amazonaws.event.ProgressEvent;
import com.amazonaws.event.ProgressListener;
import com.amazonaws.services.s3.transfer.Upload;
import java.io.File;
```

#### Code

```
File f = new File(file_path);
TransferManager xfer_mngr = new TransferManager();
try {
```

## Programming Examples

```
Upload u = xfer_mgr.upload(bucket_name, key_name, f);
u.addProgressListener(new ProgressListener() {
    public void progressChanged(ProgressEvent e) {
        double pct = e.getBytesTransferred() * 100.0 / e.getBytes();
    }
});
// block with Transfer.waitForCompletion()
TransferState xfer_state = u.getState();
System.out.println(": " + xfer_state);
} catch (AmazonServiceException e) {
    System.err.println(e.getErrorMessage());
    System.exit(1);
}
xfer_mgr.shutdownNow();
```

See the [complete example](#).

### Get the Progress of Subtransfers

The `MultipleFileUpload` class can return information about its subtransfers by calling its `getSubTransfers` method. It returns an unmodifiable `Collection` of `Upload` objects that provide the individual transfer status and progress of each subtransfer.

#### Imports

```
import com.amazonaws.services.s3.transfer.Upload;
import com.amazonaws.services.s3.transfer.MultipleFileUpload;
import java.util.ArrayList;
import java.util.Collection;
```

#### Code

```
Collection<? extends Upload> sub_xfers = new ArrayList<Upload>();
sub_xfers = multi_upload.getSubTransfers();
```

See the [complete example](#).

#### More Info

- [Object Keys](#) in the *Amazon S3 Developer Guide*

## Amazon SQS Examples

This section provides examples of programming [Amazon SQS](#) using the [AWS SDK for Java](#).

## Note

The examples include only the code needed to demonstrate each technique. The [complete example code is available on GitHub](#). From there, you can download a single source file or clone the repository locally to get all the examples to build and run.

## Working with Amazon SQS Message Queues

A *message queue* is the logical container used for sending messages reliably in Amazon SQS. There are two types of queues: *standard* and *first-in, first-out* (FIFO). To learn more about queues and the differences between these types, see the *Amazon SQS Developer Guide*.

This topic describes how to create, list, delete, and get the URL of an Amazon SQS queue by using the AWS SDK for Java.

### Create a Queue

Use the `AmazonSQS` client's `createQueue` method, providing a `CreateQueueRequest` object that describes the queue parameters.

#### Imports

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
import com.amazonaws.services.sqs.model.AmazonSQSEException;
import com.amazonaws.services.sqs.model.CreateQueueRequest;
```

#### Code

```
AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
CreateQueueRequest cq_request = new CreateQueueRequest(QUEUE_NAME)
    .addAttributesEntry("DelaySeconds", "60")
    .addAttributesEntry("MessageRetentionPeriod", "86400");

try {
    sqs.createQueue(cq_request);
} catch (AmazonSQSEException e) {
    if (!e.getErrorCode().equals("QueueAlreadyExists")) {
        throw e;
    }
}
```

You can use the simplified form of `createQueue`, which needs only a queue name, to create a standard queue.

```
sqs.createQueue("MyQueue" + new Date().getTime());
```

See the [complete sample](#).

## Programming Examples

### ***Listing Queues***

To list the Amazon SQS queues for your account, call the [AmazonSQS](#) client's `listQueues` method.

#### **Imports**

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
import com.amazonaws.services.sqs.model.ListQueuesResult;
```

#### **Code**

```
AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
ListQueuesResult lq_result = sqs.listQueues();
System.out.println("Your SQS Queue URLs:");
for (String url : lq_result.getQueueUrls()) {
    System.out.println(url);
}
```

Using the `listQueues` overload without any parameters returns *all queues*. You can filter the returned results by passing it a `ListQueuesRequest` object.

#### **Imports**

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
import com.amazonaws.services.sqs.model.ListQueuesRequest;
```

#### **Code**

```
AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
String name_prefix = "Queue";
lq_result = sqs.listQueues(new ListQueuesRequest(name_prefix));
System.out.println("Queue URLs with prefix: " + name_prefix);
for (String url : lq_result.getQueueUrls()) {
    System.out.println(url);
}
```

See the [complete sample](#).

### ***Get the URL for a Queue***

Call the [AmazonSQS](#) client's `getQueueUrl` method.

#### **Imports**

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
```

#### **Code**

## Programming Examples

```
AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
String queue_url = sqs.getQueueUrl(QUEUE_NAME).getQueueUrl();
```

See the [complete sample](#).

### Delete a Queue

Provide the queue's URL to the [AmazonSQS](#) client's `deleteQueue` method.

#### Imports

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
```

#### Code

```
AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
sqs.deleteQueue(queue_url);
```

See the [complete sample](#).

### More Info

- How Amazon SQS Queues Work in the [Amazon SQS Developer Guide](#)
- [CreateQueue](#) in the [Amazon SQS API Reference](#)
- [GetQueueUrl](#) in the [Amazon SQS API Reference](#)
- [ListQueues](#) in the [Amazon SQS API Reference](#)
- [DeleteQueues](#) in the [Amazon SQS API Reference](#)

## Sending, Receiving, and Deleting Amazon SQS Messages

This topic describes how to send, receive and delete Amazon SQS messages. Messages are always delivered using an [SQS Queue](#).

### Send a Message

Add a single message to an Amazon SQS queue by calling the [AmazonSQS](#) client's `sendMessage` method. Provide a [SendMessageRequest](#) object that contains the queue's URL, message body, and optional delay value (in seconds).

#### Imports

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
import com.amazonaws.services.sqs.model.SendMessageRequest;
```

#### Code

## Programming Examples

```
final AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();

String queueUrl = sqs.getQueueUrl(QUEUE_NAME).getQueueUrl();

SendMessageRequest cm_request = new SendMessageRequest()
    .withQueueUrl(queueUrl)
    .withMessageBody("hello world")
    .withDelaySeconds(5);
sqs.sendMessage(cm_request);
```

### *Send Multiple Messages at Once*

You can send more than one message in a single request. To send multiple messages, use the [AmazonSQS](#) client's `sendMessageBatch` method, which takes a `SendMessageBatchRequest` containing the queue URL and a list of messages (each one a `SendMessageBatchRequestEntry`) to send. You can also set an optional delay value per message.

#### Imports

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
import com.amazonaws.services.sqs.model.SendMessageBatchRequest;
import com.amazonaws.services.sqs.model.SendMessageBatchRequestEntry;
```

#### Code

```
final AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();

String queueUrl = sqs.getQueueUrl(QUEUE_NAME).getQueueUrl();

SendMessageBatchRequest smb_request = new SendMessageBatchRequest()
    .withQueueUrl(queueUrl)
    .withEntries(
        new SendMessageBatchRequestEntry(
            "msg_1", "Hello from message 1"),
        new SendMessageBatchRequestEntry(
            "msg_2", "Hello from message 2")
            .withDelaySeconds(10));
sqs.sendMessageBatch(smb_request);
```

See the [complete sample](#).

### **Receive Messages**

Retrieve any messages that are currently in the queue by calling the [AmazonSQS](#) client's `receiveMessage` method, passing it the queue's URL. Messages are returned as a list of `Message` objects.

#### Imports

## Programming Examples

```
import com.amazonaws.services.sqs.AmazonSQS;
import com.amazonaws.services.sqs.AmazonSQSClientBuilder;
import com.amazonaws.services.sqs.model.Message;
```

### Code

```
final AmazonSQS sqs = AmazonSQSClientBuilder.defaultClient();
String queueUrl = sqs.getQueueUrl(QUEUE_NAME).getQueueUrl();
List<Message> messages = sqs.receiveMessage(queueUrl).getMessages();
```

### **Delete Messages after Receipt**

After receiving a message and processing its contents, delete the message from the queue by sending the message's receipt handle and queue URL to the [AmazonSQS](#) client's `deleteMessage` method.

### Code

```
for (Message m : messages) {
    sqs.deleteMessage(queueUrl, m.getReceiptHandle());
}
```

See the [complete sample](#).

### **More Info**

- [How Amazon SQS Queues Work](#) in the *Amazon SQS Developer Guide*
- [SendMessage](#) in the *Amazon SQS API Reference*
- [SendMessageBatch](#) in the *Amazon SQS API Reference*
- [ReceiveMessage](#) in the *Amazon SQS API Reference*
- [DeleteMessage](#) in the *Amazon SQS API Reference*

## Getting Temporary Credentials with AWS STS

You can use AWS Security Token Service ([AWS STS](#)) to get temporary, limited-privilege credentials that can be used to access AWS services.

There are three steps involved in using AWS STS:

1. Activate a region (optional).
2. Retrieve temporary security credentials from AWS STS.
3. Use the credentials to access AWS resources.

## Note

Activating a region is *optional*; by default, temporary security credentials are obtained from the global endpoint `sts.amazonaws.com`. However, to reduce latency and to enable you to build redundancy into your requests by using additional endpoints if an AWS STS request to the first endpoint fails, you can activate regions that are geographically closer to your services or applications that use the credentials.

## (Optional) Activate and use an AWS STS region

To activate a region for use with AWS STS, use the AWS Management Console to select and activate the region.

### To activate additional STS regions

1. Sign in as an IAM user with permissions to perform IAM administration tasks "`iam:*`" for the account for which you want to activate AWS STS in a new region.
2. Open the IAM console and in the navigation pane click **Account Settings**.
3. Expand the **STS Regions** list, find the region that you want to use, and then click **Activate**.

After this, you can direct calls to the STS endpoint that is associated with that region.

## Note

For more information about activating STS regions and for a list of the available AWS STS endpoints, see [Activating and Deactivating AWS STS in an AWS Region](#) in the *IAM User Guide*.

## Retrieve temporary security credentials from AWS STS

### To retrieve temporary security credentials using the AWS SDK for Java

1. Create an `AWSecurityTokenServiceClient` object:

```
AWSecurityTokenServiceClient sts_client = new AWSecurityTokenServiceClient();
```

When creating the client with no arguments, the default credential provider chain is used to retrieve credentials. You can provide a specific credential provider if you want. For more information, see [Providing AWS Credentials](#) in the AWS SDK for Java.

2. *Optional*; requires that you have activated the region) Set the endpoint for the STS client:

```
sts_client.setEndpoint("sts-endpoint.amazonaws.com");
```

where `sts-endpoint` represents the STS endpoint for your region.

### Important

Do not use the `setRegion` method to set a regional endpoint—for backwards compatibility, that method continues to use the single global endpoint of `sts.amazonaws.com`.

3. Create a `GetSessionTokenRequest` object, and optionally set the duration in seconds for which the temporary credentials are valid:

```
GetSessionTokenRequest session_token_request = new GetSessionTokenRequest();
session_token_request.setDurationSeconds(7200); // optional.
```

The duration of temporary credentials can range from 900 seconds (15 minutes) to 129600 seconds (36 hours) for IAM users. If a duration isn't specified, then 43200 seconds (12 hours) is used by default.

For a root AWS account, the valid range of temporary credentials is from 900 to 3600 seconds (1 hour), with a default value of 3600 seconds if no duration is specified.

### Important

It is *strongly recommended*, from a security standpoint, that you *use IAM users instead of the root account for AWS access*. For more information, see *IAM Best Practices* in the *IAM User Guide*.

4. Call `getSessionToken` on the STS client to get a session token, using the `GetSessionTokenRequest` object:

```
GetSessionTokenResult session_token_result =
    sts_client.getSessionToken(session_token_request);
```

5. Get session credentials using the result of the call to `getSessionToken`:

```
Credentials session_creds = session_token_result.getCredentials();
```

The `session_creds` provide access only for the duration that was specified by the `GetSessionTokenRequest` object. Once the credentials expire, you will need to call `getSessionToken` again to obtain a new session token for continued access to AWS.

## Use the temporary credentials to access AWS resources

Once you have temporary security credentials, you can use them to initialize an AWS service client to use its resources, using the technique described in [Explicitly Specifying Credentials](#).

For example, to create an S3 client using temporary service credentials:

## Programming Examples

```
BasicSessionCredentials sessionCredentials = new BasicSessionCredentials(  
    session_creds.getAccessKeyId(),  
    session_creds.getSecretAccessKey(),  
    session_creds.getSessionToken());  
  
AmazonS3 s3 = AmazonS3ClientBuilder.standard()  
    .withCredentials(new AWSStaticCredentialsProvider(sessionCredentials))  
    .build();
```

You can now use the `AmazonS3` object to make Amazon S3 requests.

### For more information

For more information about how to use temporary security credentials to access AWS resources, visit the following sections in the *IAM User Guide*:

- [Requesting Temporary Security Credentials](#)
- [Controlling Permissions for Temporary Security Credentials](#)
- [Using Temporary Security Credentials to Request Access to AWS Resources](#)
- [Activating and Deactivating AWS STS in an AWS Region](#)

## Amazon SWF

[Amazon SWF](#) is a workflow-management service that helps developers build and scale distributed workflows that can have parallel or sequential steps consisting of activities, child workflows or even [Lambda](#) tasks.

There are two ways to work with Amazon SWF using the AWS SDK for Java, by using the SWF *client* object, or by using the AWS Flow Framework for Java. The AWS Flow Framework for Java is more difficult to configure initially, since it makes heavy use of annotations and relies on additional libraries such as AspectJ and the Spring Framework. However, for large or complex projects, you will save coding time by using the AWS Flow Framework for Java. For more information, see the *AWS Flow Framework for Java Developer Guide*.

This section provides examples of programming Amazon SWF by using the AWS SDK for Java client directly.

### Amazon SWF Basics

These are general patterns for working with Amazon SWF using the AWS SDK for Java. It is meant primarily for reference. For a more complete introductory tutorial, see *Building a Simple Amazon SWF Application*.

### Dependencies

Basic Amazon SWF applications will require the following dependencies, which are included with the AWS SDK for Java:

- `aws-java-sdk-1.11.*.jar`
- `commons-logging-1.1.*.jar`
- `httpclient-4.3.*.jar`
- `httpcore-4.3.*.jar`
- `jackson-annotations-2.5.*.jar`

## Programming Examples

- jackson-core-2.5.\*.jar
- jackson-databind-2.5.\*.jar
- joda-time-2.8.\*.jar

### Note

the version numbers of these packages will differ depending on the version of the SDK that you have, but the versions that are supplied with the SDK have been tested for compatibility, and are the ones you should use.

AWS Flow Framework for Java applications require additional setup, *and* additional dependencies. See the *AWS Flow Framework for Java Developer Guide* for more information about using the framework.

### Imports

In general, you can use the following imports for code development:

```
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflow;
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflowClientBuilder;
```

It's a good practice to import only the classes you require, however. You will likely end up specifying particular classes in the `com.amazonaws.services.simpleworkflow.model` workspace:

```
import com.amazonaws.services.simpleworkflow.model.PollForActivityTaskRequest;
import com.amazonaws.services.simpleworkflow.model.RespondActivityTaskCompletedRequest;
import com.amazonaws.services.simpleworkflow.model.RespondActivityTaskFailedRequest;
import com.amazonaws.services.simpleworkflow.model.TaskList;
```

If you are using the AWS Flow Framework for Java, you will import classes from the `com.amazonaws.services.simpleworkflow.flow` workspace. For example:

```
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflow;
import com.amazonaws.services.simpleworkflow.flow.ActivityWorker;
```

### Note

The AWS Flow Framework for Java has additional requirements beyond those of the base AWS SDK for Java. For more information, see the *AWS Flow Framework for Java Developer Guide*.

### Using the SWF client class

## Programming Examples

Your basic interface to Amazon SWF is through either the [AmazonSimpleWorkflowClient](#) or [AmazonSimpleWorkflowAsyncClient](#) classes. The main difference between these is that the `*AsyncClient` class return `Future` objects for concurrent (asynchronous) programming.

```
AmazonSimpleWorkflowClient swf = AmazonSimpleWorkflowClientBuilder.defaultClient();
```

## Building a Simple Amazon SWF Application

This topic will introduce you to programming Amazon SWF applications with the AWS SDK for Java, while presenting a few important concepts along the way.

### About the example

The example project will create a workflow with a single activity that accepts workflow data passed through the AWS cloud (In the tradition of HelloWorld, it'll be the name of someone to greet) and then prints a greeting in response.

While this seems very simple on the surface, Amazon SWF applications consist of a number of parts working together:

- A **domain**, used as a logical container for your workflow execution data.
- One or more **workflows** which represent code components that define logical order of execution of your workflow's activities and child workflows.
- A **workflow worker**, also known as a *decider*, that polls for decision tasks and schedules activities or child workflows in response.
- One or more **activities**, each of which represents a unit of work in the workflow.
- An **activity worker** that polls for activity tasks and runs activity methods in response.
- One or more **task lists**, which are queues maintained by Amazon SWF used to issue requests to the workflow and activity workers. Tasks on a task list meant for workflow workers are called *decision tasks*. Those meant for activity workers are called *activity tasks*.
- A **workflow starter** that begins your workflow execution.

Behind the scenes, Amazon SWF orchestrates the operation of these components, coordinating their flow from the AWS cloud, passing data between them, handling timeouts and heartbeat notifications, and logging workflow execution history.

## Prerequisites

### Development environment

The development environment used in this tutorial consists of:

- The [AWS SDK for Java](#).
- [Apache Maven](#) (3.3.1).
- JDK 1.7 or later. This tutorial was developed and tested using JDK 1.8.0.
- A good Java text editor (your choice).

## Note

If you use a different build system than Maven, you can still create a project using the appropriate steps for your environment and use the the concepts provided here to follow along. More information about configuring and using the AWS SDK for Java with various build systems is provided in *Getting Started*.

Likewise, but with more effort, the steps shown here can be implemented using any of the AWS SDKs with support for Amazon SWF.

All of the necessary external dependencies are included with the AWS SDK for Java, so there's nothing additional to download.

### AWS access

To access Amazon Web Services (AWS), you must have an active AWS account. For information about signing up for AWS and creating an IAM user (recommended over using root account credentials), see *Sign Up for AWS and Create an IAM User*.

This tutorial uses the terminal (command-line) to run the example code, and expects that you have your AWS credentials and configuration accessible to the SDK. The easiest way to do this is to use the environment variables `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY`. You should also set the `AWS_REGION` to the region you want to use.

For example, on Linux, macOS, or Unix, set the variables this way:

```
export AWS_ACCESS_KEY_ID=your_access_key_id
export AWS_SECRET_ACCESS_KEY=your_secret_access_key
export AWS_REGION=us-east-1
```

To set these variables on Windows, use these commands:

```
set AWS_ACCESS_KEY_ID=your_access_key_id
set AWS_SECRET_ACCESS_KEY=your_secret_access_key
set AWS_REGION=us-east-1
```

## Important

Substitute your own access key, secret access key and region information for the example values shown here.

For more information about configuring your credentials for the SDK, see *Set up AWS Credentials and Region for Development*.

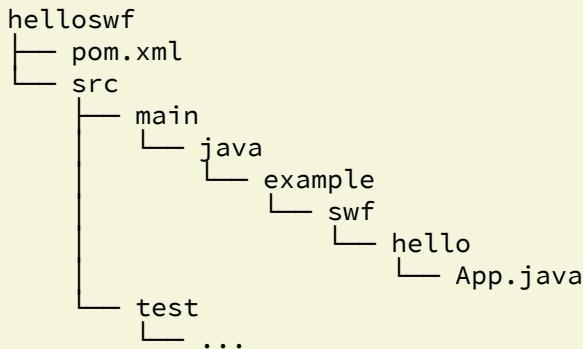
## Create a SWF project

## Programming Examples

1. Start a new project with Maven:

```
mvn archetype:generate -DartifactId=helloswf \
-DgroupId=example.swf.hello -DinteractiveMode=false
```

This will create a new project with a standard maven project structure:



You can ignore or delete the `test` directory and all it contains, we won't be using it for this tutorial. You can also delete `App.java`, since we'll be replacing it with new classes.

2. Edit the project's `pom.xml` file and add the `aws-java-sdk-simpleworkflow` module by adding a dependency for it within the `<dependencies>` block.

```
<dependencies>
<dependency>
    <groupId>com.amazonaws</groupId>
    <artifactId>aws-java-sdk-simpleworkflow</artifactId>
    <version>1.11.78</version>
</dependency>
</dependencies>
```

3. Make sure that Maven builds your project with JDK 1.7+ support. Add the following to your project (either before or after the `<dependencies>` block) in `pom.xml`:

```
<build>
    <plugins>
        <plugin>
            <artifactId>maven-compiler-plugin</artifactId>
            <configuration>
                <source>1.7</source>
                <target>1.7</target>
            </configuration>
        </plugin>
    </plugins>
</build>
```

## Code the project

The example project will consist of four separate applications, which we'll visit one by one:

- **HelloTypes.java**—contains the project's domain, activity and workflow type data, shared with the other components. It also handles registering these types with SWF.
- **ActivityWorker.java**—contains the activity worker, which polls for activity tasks and runs activities in response.
- **WorkflowWorker.java**—contains the workflow worker (decider), which polls for decision tasks and schedules new activities.
- **WorkflowStarter.java**—contains the workflow starter, which starts a new workflow execution, which will cause SWF to start generating decision and workflow tasks for your workers to consume.

### Common steps for all source files

All of the files that you create to house your Java classes will have a few things in common. In the interest of time, these steps *will be implied every time you add a new file to the project*:

1. Create the file in the in the project's `src/main/java/example/swf/hello/` directory.
2. Add a `package` declaration to the beginning of each file to declare its namespace. The example project uses:

```
package aws.example.helloswf;
```

3. Add `import` declarations for the `AmazonSimpleWorkflowClient` class and for multiple classes in the `com.amazonaws.services.simpleworkflow.model` namespace. To simplify things, we'll use:

```
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflow;
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflowClientBuilder;
import com.amazonaws.services.simpleworkflow.model.*;
```

### Register a domain, workflow and activity types

We'll begin by creating a new executeable class, `HelloTypes.java`. This file will contain shared data that different parts of your workflow will need to know about, such as the name and version of your activity and workflow types, the domain name and the task list name.

1. Open your text editor and create the file `HelloTypes.java`, adding a package declaration and imports according to the common steps.
2. Declare the `HelloTypes` class and provide it with values to use for your registered activity and workflow types:

```
public class HelloTypes {
    public static final String DOMAIN = "HelloDomain";
    public static final String TASKLIST = "HelloTasklist";
    public static final String WORKFLOW = "HelloWorkflow";
    public static final String WORKFLOW_VERSION = "1.0";
    public static final String ACTIVITY = "HelloActivity";
```

## Programming Examples

```
    public static final String ACTIVITY_VERSION = "1.0";
}
```

These values will be used throughout the code.

3. After the String declarations, create an instance of the [AmazonSimpleWorkflowClient](#) class. This is the basic interface to the Amazon SWF methods provided by the AWS SDK for Java.

```
private static final AmazonSimpleWorkflow swf =
    AmazonSimpleWorkflowClientBuilder.defaultClient();
```

4. Add a new function to register a SWF domain. A *domain* is a logical container for a number of related SWF activity and workflow types. SWF components can only communicate with each other if they exist within the same domain.

```
public static void registerDomain() {
    try {
        System.out.println("** Registering the domain '" + DOMAIN + "'.");
        swf.registerDomain(new RegisterDomainRequest()
            .withName(DOMAIN)
            .withWorkflowExecutionRetentionPeriodInDays("1"));
    } catch (DomainAlreadyExistsException e) {
        System.out.println("** Domain already exists!");
    }
}
```

When you register a domain, you provide it with a *name* (any set of 1 – 256 characters excluding :, /, |, control characters or the literal string 'arn') and a *retention period*, which is the number of days that Amazon SWF will keep your workflow's execution history data after a workflow execution has completed. The maximum workflow execution retention period is 90 days. See [RegisterDomainRequest](#) for more information.

If a domain with that name already exists, a [DomainAlreadyExistsException](#) is raised. Because we're unconcerned if the domain has already been created, we can ignore the exception.

### Tip

This code demonstrates a common pattern when working with AWS SDK for Java methods, data for the method is supplied by a class in the `simpleworkflow.model` namespace, which you instantiate and populate using the chainable `.with*` methods.

5. Add a function to register a new activity type. An *activity* represents a unit of work in your workflow.

```
public static void registerActivityType() {
    try {
        System.out.println("** Registering the activity type '" + ACTIVITY + "
```

```

        " -" + ACTIVITY_VERSION + ".");
swf.registerActivityType(new RegisterActivityTypeRequest()
    .withDomain(DOMAIN)
    .withName(ACTIVITY)
    .withVersion(ACTIVITY_VERSION)
    .withDefaultTaskList(new TaskList().withName(TASKLIST))
    .withDefaultTaskScheduleToStartTimeout("30")
    .withDefaultTaskStartToCloseTimeout("600")
    .withDefaultTaskScheduleToCloseTimeout("630")
    .withDefaultTaskHeartbeatTimeout("10"));
} catch (TypeAlreadyExistsException e) {
    System.out.println("** Activity type already exists!");
}
}

```

An activity type is identified by a *name* and a *version*, which are used to uniquely identify the activity from any others in the domain that it's registered in. Activities also contain a number of optional parameters, such as the default task-list used to receive tasks and data from SWF and a number of different timeouts that you can use to place constraints upon how long different parts of the activity execution can take. See [RegisterActivityTypeRequest](#) for more information.

### Tip

All timeout values are specified in *seconds*. See [Amazon SWF Timeout Types](#) for a full description of how timeouts affect your workflow executions.

If the activity type that you're trying to register already exists, an [TypeAlreadyExistsException](#) is raised.

6. Add a function to register a new workflow type. A *workflow*, also known as a *decider* represents the logic of your workflow's execution.

```

public static void registerWorkflowType() {
    try {
        System.out.println("** Registering the workflow type '" + WORKFLOW +
            " -" + WORKFLOW_VERSION + ".");
        swf.registerWorkflowType(new RegisterWorkflowTypeRequest()
            .withDomain(DOMAIN)
            .withName(WORKFLOW)
            .withVersion(WORKFLOW_VERSION)
            .withDefaultChildPolicy(ChildPolicy.TERMINATE)
            .withDefaultTaskList(new TaskList().withName(TASKLIST))
            .withDefaultTaskStartToCloseTimeout("30"));
    } catch (TypeAlreadyExistsException e) {
        System.out.println("** Workflow type already exists!");
    }
}

```

## Programming Examples

Similar to activity types, workflow types are identified by a *name* and a *version* and also have configurable timeouts. See [RegisterWorkflowTypeRequest](#) for more information.

If the workflow type that you're trying to register already exists, an [TypeAlreadyExistsException](#) is raised.

7. Finally, make the class executable by providing it a `main` method, which will register the domain, the activity type, and the workflow type in turn:

```
public static void main(String[] args) {  
    registerDomain();  
    registerWorkflowType();  
    registerActivityType();  
}
```

You can build and run the application now to run the registration script, or continue with coding the activity and workflow workers. Once the domain, workflow and activity have been registered, you won't need to run this again—these types persist until you deprecate them yourself.

### *Implement the activity worker*

An *activity* is the basic unit of work in a workflow. A workflow provides the logic, scheduling activities to be run (or other actions to be taken) in response to decision tasks. A typical workflow usually consists of a number of activities that can run synchronously, asynchronously, or a combination of both.

The *activity worker* is the bit of code that polls for activity tasks that are generated by Amazon SWF in response to workflow decisions. When it receives an activity task, it runs the corresponding activity and returns a success/failure response back to the workflow.

We'll implement a simple activity worker that drives a single activity.

1. Open your text editor and create the file `ActivityWorker.java`, adding a package declaration and imports according to the common steps.
2. Add the `ActivityWorker` class to the file, and give it a data member to hold a SWF client that we'll use to interact with Amazon SWF:

```
public class ActivityWorker {  
    private static final AmazonSimpleWorkflow swf =  
        AmazonSimpleWorkflowClientBuilder.defaultClient();  
}
```

3. Add the method that we'll use as an activity:

```
private static String sayHello(String input) throws Throwable {  
    return "Hello, " + input + "!";  
}
```

The activity simply takes a string, combines it into a greeting and returns the result. Although there is little chance that this activity will raise an exception, it's a good idea to design activities that can raise an error if something goes wrong.

4. Add a `main` method that we'll use as the activity task polling method. We'll start it by adding some code to poll the task list for activity tasks:

## Programming Examples

```
public static void main(String[] args) {
    while (true) {
        System.out.println("Polling for an activity task from the tasklist ''"
            + HelloTypes.TASKLIST + " in the domain '" +
            HelloTypes.DOMAIN + "'");

        ActivityTask task = swf.pollForActivityTask(
            new PollForActivityTaskRequest()
                .withDomain(HelloTypes.DOMAIN)
                .withTaskList(
                    new TaskList().withName(HelloTypes.TASKLIST)));
    }

    String task_token = task.getTaskToken();
}
```

The activity receives tasks from Amazon SWF by calling the SWF client's `pollForActivityTask` method, specifying the domain and task list to use in the passed-in `PollForActivityTaskRequest`.

Once a task is received, we retrieve a unique identifier for it by calling the task's `getTaskToken` method.

5. Next, write some code to process the tasks that come in. Add the following to your `main` method, right after the code that polls for the task and retrieves its task token.

```
if (task_token != null) {
    String result = null;
    Throwable error = null;

    try {
        System.out.println("Executing the activity task with input '" +
            task.getInput() + "'.");
        result = sayHello(task.getInput());
    } catch (Throwable th) {
        error = th;
    }

    if (error == null) {
        System.out.println("The activity task succeeded with result ''"
            + result + "'");
        swf.respondActivityTaskCompleted(
            new RespondActivityTaskCompletedRequest()
                .withTaskToken(task_token)
                .withResult(result));
    } else {
        System.out.println("The activity task failed with the error ''"
            + error.getClass().getSimpleName() + "'");
        swf.respondActivityTaskFailed(
            new RespondActivityTaskFailedRequest()
                .withTaskToken(task_token)
                .withReason(error.getClass().getSimpleName())
                .withDetails(error.getMessage()));
    }
}
```

```
    }  
}
```

If the task token is not `null`, then we can start running the activity method (`sayHello`), providing it with the input data that was sent with the task.

If the task *succeeded* (no error was generated), then the worker responds to SWF by calling the SWF client's `respondActivityTaskCompleted` method with a `RespondActivityTaskCompletedRequest` object containing the task token and the activity's result data.

On the other hand, if the task *failed*, then we respond by calling the `respondActivityTaskFailed` method with a `RespondActivityTaskFailedRequest` object, passing it the task token and information about the error.

### Tip

This activity will not shut down gracefully if killed. Although it is beyond the scope of this tutorial, an alternative implementation of this activity worker is provided in the accompanying topic, *Shutting Down Activity and Workflow Workers Gracefully*.

### Implement the workflow worker

Your workflow logic resides in a piece of code known as a **workflow worker**. The workflow worker polls for decision tasks that are sent by Amazon SWF in the domain, and on the default tasklist, that the workflow type was registered with.

When the workflow worker receives a task, it makes some sort of decision (usually whether to schedule a new activity or not) and takes an appropriate action (such as scheduling the activity).

1. Open your text editor and create the file `WorkflowWorker.java`, adding a package declaration and imports according to the common steps.
2. Add a few additional imports to the file:

```
import java.util.ArrayList;  
import java.util.List;  
import java.util.UUID;
```

3. Declare the `WorkflowWorker` class, and create an instance of the `AmazonSimpleWorkflowClient` class used to access SWF methods.

```
public class WorkflowWorker {  
    private static final AmazonSimpleWorkflow swf =  
        AmazonSimpleWorkflowClientBuilder.defaultClient();  
}
```

4. Add the `main` method. The method loops continuously, polling for decision tasks using the SWF client's `pollForDecisionTask` method. The `PollForDecisionTaskRequest` provides the details.

## Programming Examples

```
public static void main(String[] args) {
    PollForDecisionTaskRequest task_request =
        new PollForDecisionTaskRequest()
            .withDomain(HelloTypes.DOMAIN)
            .withTaskList(new TaskList().withName(HelloTypes.TASKLIST));

    while (true) {
        System.out.println(
            "Polling for a decision task from the tasklist '" +
            HelloTypes.TASKLIST + "' in the domain '" +
            HelloTypes.DOMAIN + "'.");
        DecisionTask task = swf.pollForDecisionTask(task_request);

        String taskToken = task.getTaskToken();
        if (taskToken != null) {
            try {
                executeDecisionTask(taskToken, task.getEvents());
            } catch (Throwable th) {
                th.printStackTrace();
            }
        }
    }
}
```

Once a task is received, we call its `getTaskToken` method, which returns a string that can be used to identify the task. If the returned token is not `null`, then we process it further in the `executeDecisionTask` method, passing it the task token and the list of `HistoryEvent` objects sent with the task.

5. Add the `executeDecisionTask` method, taking the task token (a `String`) and the `HistoryEvent` list.

```
private static void executeDecisionTask(String taskToken, List<HistoryEvent> events)
    throws Throwable {
    List<Decision> decisions = new ArrayList<Decision>();
    String workflow_input = null;
    int scheduled_activities = 0;
    int open_activities = 0;
    boolean activity_completed = false;
    String result = null;
}
```

We also set up some data members to keep track of things such as:

- A list of `Decision` objects used to report the results of processing the task.
- A String to hold workflow input provided by the "WorkflowExecutionStarted" event

## Programming Examples

- a count of the scheduled and open (running) activities to avoid scheduling the same activity when it has already been scheduled or is currently running.
- a boolean to indicate that the activity has completed.
- A String to hold the activity results, for returning it as our workflow result.

6. Next, add some code to `executeDecisionTask` to process the `HistoryEvent` objects that were sent with the task, based on the event type reported by the `getEventType` method.

```
System.out.println("Executing the decision task for the history events: []");
for (HistoryEvent event : events) {
    System.out.println(" " + event);
    switch(event.getEventType()) {
        case "WorkflowExecutionStarted":
            workflow_input =
                event.getWorkflowExecutionStartedEventAttributes()
                    .getInput();
            break;
        case "ActivityTaskScheduled":
            scheduled_activities++;
            break;
        case "ScheduleActivityTaskFailed":
            scheduled_activities--;
            break;
        case "ActivityTaskStarted":
            scheduled_activities--;
            open_activities++;
            break;
        case "ActivityTaskCompleted":
            open_activities--;
            activity_completed = true;
            result = event.getActivityTaskCompletedEventAttributes()
                .getResult();
            break;
        case "ActivityTaskFailed":
            open_activities--;
            break;
        case "ActivityTaskTimedOut":
            open_activities--;
            break;
    }
}
System.out.println("]");
```

For the purposes of our workflow, we are most interested in:

- the "WorkflowExecutionStarted" event, which indicates that the workflow execution has started (typically meaning that you should run the first activity in the workflow), and that provides the initial input provided to the workflow. In this case, it's the name portion of our greeting, so it's saved in a String for use when scheduling the activity to run.

## Programming Examples

- the "ActivityTaskCompleted" event, which is sent once the scheduled activity is complete. The event data also includes the return value of the completed activity. Since we have only one activity, we'll use that value as the result of the entire workflow.

The other event types can be used if your workflow requires them. See the [HistoryEvent](#) class description for information about each event type.

### Note

Strings in switch statements were introduced in Java 7. If you're using an earlier version of Java, you can make use of the [EventType](#) class to convert the String returned by `history_event.getType()` to an enum value and then back to a String if necessary:

```
EventType et = EventType.fromValue(event.getEventType());
```

7. After the `switch` statement, add more code to respond with an appropriate *decision* based on the task that was received.

```
if (activity_completed) {
    decisions.add(
        new Decision()
            .withDecisionType(DecisionType.CompleteWorkflowExecution)
            .withCompleteWorkflowExecutionDecisionAttributes(
                new CompleteWorkflowExecutionDecisionAttributes()
                    .withResult(result)));
} else {
    if (open_activities == 0 && scheduled_activities == 0) {

        ScheduleActivityTaskDecisionAttributes attrs =
            new ScheduleActivityTaskDecisionAttributes()
                .withActivityType(new ActivityType()
                    .withName(HelloTypes.ACTIVITY)
                    .withVersion(HelloTypes.ACTIVITY_VERSION))
                .withActivityId(UUID.randomUUID().toString())
                .withInput(workflow_input);

        decisions.add(
            new Decision()
                .withDecisionType(DecisionType.ScheduleActivityTask)
                .withScheduleActivityTaskDecisionAttributes(attrs));
    } else {
        // an instance of HelloActivity is already scheduled or running. Do nothing, another
        // task will be scheduled once the activity completes, fails or times out
    }
}

System.out.println("Exiting the decision task with the decisions " + decisions);
```

## Programming Examples

- If the activity hasn't been scheduled yet, we respond with a `ScheduleActivityTask` decision, which provides information in a `ScheduleActivityTaskDecisionAttributes` structure about the activity that Amazon SWF should schedule next, also including any data that Amazon SWF should send to the activity.
- If the activity was completed, then we consider the entire workflow completed and respond with a `CompletedWorkflowExecution` decision, filling in a `CompleteWorkflowExecutionDecisionAttributes` structure to provide details about the completed workflow. In this case, we return the result of the activity.

In either case, the decision information is added to the `Decision` list that was declared at the top of the method.

8. Complete the decision task by returning the list of `Decision` objects collected while processing the task. Add this code at the end of the `executeDecisionTask` method that we've been writing:

```
swf.respondDecisionTaskCompleted(  
    new RespondDecisionTaskCompletedRequest()  
        .withTaskToken(taskToken)  
        .withDecisions(decisions));
```

The SWF client's `respondDecisionTaskCompleted` method takes the task token that identifies the task as well as the list of `Decision` objects.

### *Implement the workflow starter*

Finally, we'll write some code to start the workflow execution.

1. Open your text editor and create the file `WorkflowStarter.java`, adding a package declaration and imports according to the common steps.
2. Add the `WorkflowStarter` class:

```
public class WorkflowStarter {  
    private static final AmazonSimpleWorkflow swf =  
        AmazonSimpleWorkflowClientBuilder.defaultClient();  
    public static final String WORKFLOW_EXECUTION = "HelloWorldWorkflowExecution";  
  
    public static void main(String[] args) {  
        String workflow_input = "Amazon SWF";  
        if (args.length > 0) {  
            workflow_input = args[0];  
        }  
  
        System.out.println("Starting the workflow execution '" + WORKFLOW_EXECUTION +  
            "' with input '" + workflow_input + ".");  
  
        WorkflowType wf_type = new WorkflowType()  
            .withName(HelloTypes.WORKFLOW)  
            .withVersion(HelloTypes.WORKFLOW_VERSION);  
  
        Run run = swf.startWorkflowExecution(new StartWorkflowExecutionRequest()  
            .withDomain(HelloTypes.DOMAIN)
```

## Programming Examples

```
.withWorkflowType(wf_type)
.withWorkflowId(WORKFLOW_EXECUTION)
.withInput(workflow_input)
.withExecutionStartToCloseTimeout("90"));

System.out.println("Workflow execution started with the run id '" +
    run.getRunId() + "'.");

}
```

The `WorkflowStarter` class consists of a single method, `main`, which takes an optional argument passed on the command-line as input data for the workflow.

The SWF client method, `startWorkflowExecution`, takes a `StartWorkflowExecutionRequest` object as input. Here, in addition to specifying the domain and workflow type to run, we provide it with:

- a human-readable workflow execution name
- workflow input data (provided on the command-line in our example)
- a timeout value that represents how long, in seconds, that the entire workflow should take to run.

The `Run` object that `startWorkflowExecution` returns provides a *run ID*, a value that can be used to identify this particular workflow execution in Amazon SWF's history of your workflow executions.

### Note

The run ID is generated by Amazon SWF, and is *not* the same as the workflow execution name that you pass in when starting the workflow execution.

### **Build the example**

To build the example project with Maven, go to the `helloswf` directory and type:

```
mvn package
```

The resulting `helloswf-1.0.jar` will be generated in the `target` directory.

### **Run the example**

The example consists of four separate executable classes, which are run independently of each other.

### Note

If you are using a Linux, macOS, or Unix system, you can run all of them, one after another, in a single terminal window. If you are running Windows, you should open two additional command-line instances and navigate to the `helloswf` directory in each.

## Programming Examples

### *Setting the Java classpath*

Although Maven has handled the dependencies for you, to run the example, you'll need to provide the AWS SDK library and its dependencies on your Java classpath. You can either set the CLASSPATH environment variable to the location of your AWS SDK libraries and the third-party/lib directory in the SDK, which includes necessary dependencies:

```
export CLASSPATH='target/helloswf-1.0.jar:/path/to/sdk/lib/*:/path/to/sdk/third-party/lib/*'  
java example.swf.hello.HelloTypes
```

or use the **java** command's -cp option to set the classpath while running each applications.

```
java -cp target/helloswf-1.0.jar:/path/to/sdk/lib/*:/path/to/sdk/third-party/lib/* \  
example.swf.hello.HelloTypes
```

The style that you use is up to you. If you had no trouble building the code, but then try to run the examples and get a series of "NoClassDefFound" errors, it is likely because the classpath is set incorrectly.

### *Register the domain, workflow and activity types*

Before running your workers and the workflow starter, you'll need to register the domain and your workflow and activity types. The code to do this was implemented in Register a domain, workflow and activity types.

After building, and if you've set the CLASSPATH, you can run the registration code by executing the command:

```
java aws.example.helloswf.HelloTypes
```

### *Start the activity and workflow workers*

Now that the types have been registered, you can start the activity and workflow workers. These will continue to run and poll for tasks until they are killed, so you should either run them in separate terminal windows, or, if you're running on Linux, macOS, or Unix you can use the & operator to cause each of them to spawn a separate process when run.

```
java aws.example.helloswf.ActivityWorker &  
java aws.example.helloswf.WorkflowWorker &
```

If you're running these commands in separate windows, omit the final & operator from each line.

### *Start the workflow execution*

Now that your activity and workflow workers are polling, you can start the workflow execution. This process will run until the workflow returns a completed status. You should run it in a new terminal window (unless you ran your workers as new spawned processes by using the & operator).

```
java aws.example.helloswf.WorkflowStarter
```

## Note

If you want to provide your own input data, which will be passed first to the workflow and then to the activity, add it to the command-line. For example:

```
java aws.example.helloswf.WorkflowStarter "Thelonious"
```

Once you begin the workflow execution, you should start seeing output delivered by both workers and by the workflow execution itself. When the workflow finally completes, its output will be printed to the screen.

### ***Complete source for this example***

You can browse the [complete source](#) for this example on Github in the [aws-java-developer-guide](#) repository.

### ***For more information***

- The workers presented here can result in lost tasks if they are shutdown while a workflow poll is still going on. To find out how to shut down workers gracefully, see *Shutting Down Activity and Workflow Workers Gracefully*.
- To learn more about Amazon SWF, visit the [Amazon SWF](#) home page or view the [Amazon SWF Developer Guide](#).
- You can use the AWS Flow Framework for Java to write more complex workflows in an elegant Java style using annotations. To learn more, see the [AWS Flow Framework for Java Developer Guide](#).

## Lambda Tasks

As an alternative to, or in conjunction with, Amazon SWF activities, you can use [Lambda](#) functions to represent units of work in your workflows, and schedule them similarly to activities.

This topic focuses on how to implement Amazon SWF Lambda tasks using the AWS SDK for Java. For more information about Lambda tasks in general, see [AWS Lambda Tasks](#) in the [Amazon SWF Developer Guide](#).

### ***Set up a cross-service IAM role to run your Lambda function***

Before Amazon SWF can run your Lambda function, you need to set up an IAM role to give Amazon SWF permission to run Lambda functions on your behalf. For complete information about how to do this, see [AWS Lambda Tasks](#).

You will need the Amazon Resource Name (ARN) of this IAM role when you register a workflow that will use Lambda tasks.

### ***Create a Lambda function***

You can write Lambda functions in a number of different languages, including Java. For complete information about how to author, deploy and use Lambda functions, see the [Lambda Developer Guide](#).

## Note

It doesn't matter what language you use to write your Lambda function, it can be scheduled and run by any Amazon SWF workflow, regardless of the language that your workflow code is written in. Amazon SWF handles the details of running the function and passing data to and from it.

Here's a simple Lambda function that could be used in place of the activity in *Building a Simple Amazon SWF Application*.

- This version is written in JavaScript, which can be entered directly using the AWS Management Console:

```
exports.handler = function(event, context) {
    context.succeed("Hello, " + event.who + "!");
};
```

- Here is the same function written in Java, which you could also deploy and run on Lambda:

```
package example.swf.hellolambda;

import com.amazonaws.services.lambda.runtime.Context;
import com.amazonaws.services.lambda.runtime.RequestHandler;
import com.amazonaws.util.json.JSONException;
import com.amazonaws.util.json.JSONObject;

public class SwfHelloLambdaFunction implements RequestHandler<Object, Object> {
    @Override
    public Object handleRequest(Object input, Context context) {
        String who = "Amazon SWF";
        if (input != null) {
            JSONObject js0 = null;
            try {
                js0 = new JSONObject(input.toString());
                who = js0.getString("who");
            } catch (JSONException e) {
                e.printStackTrace();
            }
        }
        return ("Hello, " + who + "!");
    }
}
```

## Tip

To learn more about deploying Java functions to Lambda, see [Creating a Deployment Package \(Java\)](#) in the *Lambda Developer Guide*. You will also want to look at the section titled [Programming Model for Authoring Lambda Functions in Java](#).

## Programming Examples

Lambda functions take an *event* or *input* object as the first parameter, and a *context* object as the second, which provides information about the request to run the Lambda function. This particular function expects input to be in JSON, with a `who` field set to the name used to create the greeting.

### **Register a workflow for use with Lambda**

For a workflow to schedule a Lambda function, you must provide the name of the IAM role that provides Amazon SWF with permission to invoke Lambda functions. You can set this during workflow registration by using the `withDefaultLambdaRole` or `setDefaultLambdaRole` methods of `RegisterWorkflowTypeRequest`.

```
System.out.println("** Registering the workflow type '" + WORKFLOW + "-" + WORKFLOW_VERSION
    + "'.");
try {
    swf.registerWorkflowType(new RegisterWorkflowTypeRequest()
        .withDomain(DOMAIN)
        .withName(WORKFLOW)
        .withDefaultLambdaRole(lambda_role_arn)
        .withVersion(WORKFLOW_VERSION)
        .withDefaultChildPolicy(ChildPolicy.TERMINATE)
        .withDefaultTaskList(new TaskList().withName(TASKLIST))
        .withDefaultTaskStartToCloseTimeout("30"));
}
catch (TypeAlreadyExistsException e) {
```

### **Schedule a Lambda task**

Schedule a Lambda task is similar to scheduling an activity. You provide a `Decision` with a `ScheduleLambdaFunction` `DecisionType` and with `ScheduleLambdaFunctionDecisionAttributes`.

```
running_functions == 0 && scheduled_functions == 0) {
AWSLambda lam = AWSLambdaClientBuilder.defaultClient();
GetFunctionConfigurationResult function_config =
    lam.getFunctionConfiguration(
        new GetFunctionConfigurationRequest()
            .withFunctionName("HelloFunction"));
String function_arn = function_config.getFunctionArn();

ScheduleLambdaFunctionDecisionAttributes attrs =
    new ScheduleLambdaFunctionDecisionAttributes()
        .withId("HelloFunction (Lambda task example)")
        .withName(function_arn)
        .withInput(workflow_input);

decisions.add(
```

In the `ScheduleLambdaFunctionDecisionAttributes`, you must supply a *name*, which is the ARN of the Lambda function to call, and an *id*, which is the name that Amazon SWF will use to identify the Lambda function in history logs.

## Programming Examples

You can also provide optional *input* for the Lambda function and set its *start to close timeout* value, which is the number of seconds that the Lambda function is allowed to run before generating a LambdaFunctionTimedOut event.

### Tip

This code uses the [AWSLambdaClient](#) to retrieve the ARN of the Lambda function, given the function name. You can use this technique to avoid hard-coding the full ARN (which includes your AWS account ID) in your code.

### **Handle Lambda function events in your decider**

Lambda tasks will generate a number of events that you can take action on when polling for decision tasks in your workflow worker, corresponding to the lifecycle of your Lambda task, with [EventType](#) values such as LambdaFunctionScheduled, LambdaFunctionStarted, and LambdaFunctionCompleted. If the Lambda function fails, or takes longer to run than its set timeout value, you will receive either a LambdaFunctionFailed or LambdaFunctionTimedOut event type, respectively.

```
boolean function_completed = false;
String result = null;

System.out.println("Executing the decision task for the history events: []");
for (HistoryEvent event : events) {
    System.out.println(" " + event);
    EventType event_type = EventType.fromValue(event.getEventType());
    switch(event_type) {
        case WorkflowExecutionStarted:
            workflow_input =
                event.getWorkflowExecutionStartedEventAttributes()
                    .getInput();
            break;
        case LambdaFunctionScheduled:
            scheduled_functions++;
            break;
        case ScheduleLambdaFunctionFailed:
            scheduled_functions--;
            break;
        case LambdaFunctionStarted:
            scheduled_functions--;
            running_functions++;
            break;
        case LambdaFunctionCompleted:
            running_functions--;
            function_completed = true;
            result = event.getLambdaFunctionCompletedEventAttributes()
                .getResult();
            break;
    }
}
```

## Programming Examples

```
case LambdaFunctionFailed:  
    running_functions--;  
    break;  
case LambdaFunctionTimedOut:  
    running_functions--;  
    break;
```

### **Receive output from your Lambda function**

When you receive a LambdaFunctionCompleted [EventType](#), you can retrieve your Lambda function's return value by first calling [getLambdaFunctionCompletedEventAttributes](#) on the [HistoryEvent](#) to get a [LambdaFunctionCompletedEventAttributes](#) object, and then calling its [getResult](#) method to retrieve the output of the Lambda function:

```
LambdaFunctionCompleted:  
running_functions--;
```

### **Complete source for this example**

You can browse the complete source code for this example on Github in the [aws-java-developer-guide](https://github.com/awsdocs/aws-java-developer-guide/tree/master/doc_source/snippets/helloswf_lambda/) repository.

## [Shutting Down Activity and Workflow Workers Gracefully](#)

The [Building a Simple Amazon SWF Application](#) topic provided a complete implementation of a simple workflow application consisting of a registration application, an activity and workflow worker, and a workflow starter.

Worker classes are designed to run continuously, polling for tasks sent by Amazon SWF in order to run activities or return decisions. Once a poll request is made, Amazon SWF records the poller and will attempt to assign a task to it.

If the workflow worker is terminated during a long poll, Amazon SWF may still try to send a task to the terminated worker, resulting in a lost task (until the task times out).

One way to handle this situation is to wait for all long poll requests to return before the worker terminates.

In this topic, we'll rewrite the activity worker from helloswf, using Java's shutdown hooks to attempt a graceful shutdown of the activity worker.

Here is the complete code:

```
package aws.example.helloswf;  
  
import java.util.concurrent.CountDownLatch;  
import java.util.concurrent.TimeUnit;  
  
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflow;  
import com.amazonaws.services.simpleworkflow.AmazonSimpleWorkflowClientBuilder;  
import com.amazonaws.services.simpleworkflow.model.ActivityTask;  
import com.amazonaws.services.simpleworkflow.model.PollForActivityTaskRequest;  
import com.amazonaws.services.simpleworkflow.model.RespondActivityTaskCompletedRequest;
```

## Programming Examples

```
import com.amazonaws.services.simpleworkflow.model.RespondActivityTaskFailedRequest;
import com.amazonaws.services.simpleworkflow.model.TaskList;

public class ActivityWorkerWithGracefulShutdown {

    private static final AmazonSimpleWorkflow swf =
        AmazonSimpleWorkflowClientBuilder.defaultClient();
    private static CountDownLatch waitForTermination = new CountDownLatch(1);
    private static volatile boolean terminate = false;

    private static String executeActivityTask(String input) throws Throwable {
        return "Hello, " + input + "!";
    }

    public static void main(String[] args) {
        Runtime.getRuntime().addShutdownHook(new Thread() {
            @Override
            public void run() {
                try {
                    terminate = true;
                    System.out.println("Waiting for the current poll request" +
                        " to return before shutting down.");
                    waitForTermination.await(60, TimeUnit.SECONDS);
                } catch (InterruptedException e) {
                    // ignore
                }
            }
        });
        try {
            pollAndExecute();
        }
        finally {
            waitForTermination.countDown();
        }
    }

    public static void pollAndExecute() {
        while (!terminate) {
            System.out.println("Polling for an activity task from the tasklist ''" +
                + HelloTypes.TASKLIST + "' in the domain '" +
                HelloTypes.DOMAIN + "'.");

            ActivityTask task = swf.pollForActivityTask(new PollForActivityTaskRequest()
                .withDomain(HelloTypes.DOMAIN)
                .withTaskList(new TaskList().withName(HelloTypes.TASKLIST)));

            String taskToken = task.getTaskToken();

            if (taskToken != null) {

```

## Programming Examples

```
String result = null;
Throwable error = null;

try {
    System.out.println("Executing the activity task with input '" +
                       + task.getInput() + ".");
    result = executeActivityTask(task.getInput());
}
catch (Throwable th) {
    error = th;
}

if (error == null) {
    System.out.println("The activity task succeeded with result '" +
                       + result + ".");
    swf.respondActivityTaskCompleted(
        new RespondActivityTaskCompletedRequest()
            .withTaskToken(taskToken)
            .withResult(result));
}
else {
    System.out.println("The activity task failed with the error '" +
                       + error.getClass().getSimpleName() + ".");
    swf.respondActivityTaskFailed(
        new RespondActivityTaskFailedRequest()
            .withTaskToken(taskToken)
            .withReason(error.getClass().getSimpleName())
            .withDetails(error.getMessage()));
}
}
```

In this version, the polling code that was in the main function in the original version has been moved into its own method, `pollAndExecute`.

The main function now uses a [CountDownLatch](#) in conjunction with a [shutdown hook](#) to cause the thread to wait for up to 60 seconds after its termination is requested before letting the thread shut down.

## Registering Domains

Every workflow and activity in [Amazon SWF](#) needs a *domain* to run in.

## To register an Amazon SWF domain

1. Create a new [RegisterDomainRequest](#) object, providing it with at least the domain name and workflow execution retention period (these parameters are both required).
  2. Call the [AmazonSimpleWorkflowClient.registerDomain](#) method with the *RegisterDomainRequest* object.

## Programming Examples

3. Catch the [DomainAlreadyExistsException](#) if the domain you're requesting already exists (in which case, no action is usually required).

The following code demonstrates this procedure:

```
public void register_swf_domain(AmazonSimpleWorkflowClient swf, String name)
{
    RegisterDomainRequest request = new RegisterDomainRequest().withName(name);
    request.setWorkflowExecutionRetentionPeriodInDays("10");
    try
    {
        swf.registerDomain(request);
    }
    catch (DomainAlreadyExistsException e)
    {
        System.out.println("Domain already exists!");
    }
}
```

## Listing Domains

You can list the [Amazon SWF](#) domains associated with your account and AWS region by registration type.

### To list Amazon SWF domains

1. Create a [ListDomainsRequest](#) object, and specify the registration status of the domains that you're interested in—this is required.
2. Call [AmazonSimpleWorkflowClient.listDomains](#) with the *ListDomainRequest* object. Results are provided in a [DomainInfos](#) object.
3. Call [getDomainInfos](#) on the returned object to get a list of [DomainInfo](#) objects.
4. Call [getName](#) on each *DomainInfo* object to get its name.

The following code demonstrates this procedure:

```
public void list_swf_domains(AmazonSimpleWorkflowClient swf)
{
    ListDomainsRequest request = new ListDomainsRequest();
    request.setRegistrationStatus("REGISTERED");
    DomainInfos domains = swf.listDomains(request);
    System.out.println("Current Domains:");
    for (DomainInfo di : domains.getDomainInfos())
    {
        System.out.println(" * " + di.getName());
    }
}
```

# Document History

This topic describes important changes to the *AWS Java Developer Guide* over the course of its history.

**Last documentation update:** Mar 09, 2017

## Jan 26, 2017

Added a new Amazon S3 topic, *Using TransferManager for Amazon S3 Operations*, and a new *Best Practices for AWS Development with the AWS SDK for Java* topic in the *Using the AWS SDK for Java* section.

## Jan 16, 2017

Added a new Amazon S3 topic, *Managing Access to Amazon S3 Buckets Using Bucket Policies*, and two new Amazon SQS topics, *Working with Amazon SQS Message Queues* and *Sending, Receiving, and Deleting Amazon SQS Messages*.

## Dec 16, 2016

Added new example topics for DynamoDB: *Working with Tables in DynamoDB* and *Working with Items in DynamoDB*.

## Sep 26, 2016

The topics in the **Advanced** section have been moved into *Using the AWS SDK for Java*, since they really are central to using the SDK.

## Aug 25, 2016

A new topic, *Creating Service Clients*, has been added to *Using the AWS SDK for Java*, which demonstrates how to use *client builders* to simplify the creation of AWS service clients.

The *Programming Examples* section has been updated with *new examples for S3* which are backed by a [repository on GitHub](#) that contains the complete example code.

## May 02, 2016

A new topic, *Asynchronous Programming*, has been added to the *Using the AWS SDK for Java* section, describing how to work with asynchronous client methods that return *Future* objects or that take an *AsyncHandler*.

## Apr 26, 2016

The *SSL Certificate Requirements* topic has been removed, since it is no longer relevant. Support for SHA-1 signed certificates was deprecated in 2015 and the site that housed the test scripts has been removed.

## Mar 14, 2016

Added a new topic to the Amazon SWF section: *Lambda Tasks*, which describes how to implement a Amazon SWF workflow that calls Lambda functions as tasks as an alternative to using traditional Amazon SWF activities.

## Mar 04, 2016

The *Amazon SWF* section has been updated with new content:

- *Amazon SWF Basics* – Provides basic information about how to include SWF in your projects.
- *Building a Simple Amazon SWF Application* – A new tutorial that provides step-by-step guidance for Java developers new to Amazon SWF.
- *Shutting Down Activity and Workflow Workers Gracefully* – Describes how you can gracefully shut down Amazon SWF worker classes using Java's concurrency classes.

## Feb 23, 2016

## Document History

The source for the *AWS Java Developer Guide* has been moved to [aws-java-developer-guide](#).

### Dec 28, 2015

*Setting the JVM TTL for DNS Name Lookups* has been moved from **Advanced** into *Using the AWS SDK for Java*, and has been rewritten for clarity.

*Using the SDK with Apache Maven* has been updated with information about how to include the SDK's bill of materials (BOM) in your project.

### Aug 04, 2015

*SSL Certificate Requirements* is a new topic in the *Getting Started* section that describes AWS' move to SHA256-signed certificates for SSL connections, and how to fix early 1.6 and previous Java environments to use these certificates, which are *required* for AWS access after September 30, 2015.

#### Note

Java 1.7+ is already capable of working with SHA256-signed certificates.

### May 14, 2014

The *introduction* and *getting started* material has been heavily revised to support the new guide structure and now includes guidance about how to *Set up AWS Credentials and Region for Development*.

The discussion of *code samples* has been moved into its own topic in the Additional Documentation and Resources section.

Information about how to view the SDK revision history has been moved into the introduction.

### May 9, 2014

The overall structure of the AWS SDK for Java documentation has been simplified, and the *Getting Started* and Additional Documentation and Resources topics have been updated.

New topics have been added:

- *Working with AWS Credentials* – discusses the various ways that you can specify credentials for use with the AWS SDK for Java.
- *Using IAM Roles to Grant Access to AWS Resources on Amazon EC2* – provides information about how to securely specify credentials for applications running on EC2 instances.

### Sep 9, 2013

This topic, *Document History*, tracks changes to the *AWS Java Developer Guide*. It is intended as a companion to the release notes history.

## About Amazon Web Services

*Amazon Web Services (AWS)* is a collection of digital infrastructure services that developers can leverage when developing their applications. The services include computing, storage, database, and application synchronization (messaging and queuing). AWS uses a pay-as-you-go service model: you are charged only for the services that you—or your applications—use. For new AWS users, a free usage tier is available. On this tier, services are free below a certain level of usage. For more information about AWS costs and the Free Tier, see [Use the AWS Free Tier](#). To obtain an AWS account, visit the [AWS home page](#) and click **Create a Free Account**.